

Maximum Nullity of Cyclic Covers of Graphs

A monodromy ceiling, and the zero forcing
number of $P(n, k)$

Research Binder

Aryan Padarthy · Mathematics

Compiled 27 September 2026

Contents

- I. Official Abstract
- II. Research Plan
- III. Research Paper
- IV. Research Log, dated, including every dead end
- V. Verification Output, 69 checks
- VI. A Certification Transcript

Every claim in the paper is labelled by the standard it meets: proved, proved by an exact certificate, proved by interval certification, or numerical. Claims that have been refuted or withdrawn, including the author's own, are listed explicitly in the paper's Status section and dated in the log.

I. Official Abstract

Zero forcing is a colouring process: a coloured vertex with exactly one uncoloured neighbour forces that neighbour to colour. The zero forcing number $Z(G)$ is the smallest starting set that colours everything, and it bounds the maximum nullity $M(G)$ over symmetric matrices with the graph's pattern. For the generalized Petersen graphs $P(n,k)$, a 2020 paper claimed $Z(P(n,3))=8$ for n at least 12. A July 2026 note showed this is false and conjectured $Z(P(n,3))=8$ for n at least 13, identifying the missing ingredient as a lower-bound proof valid for all large n .

I prove that eliminating the inner coordinates of any matrix carrying the $P(n,k)$ pattern leaves a single linear recurrence of order exactly $2k+2$. Its nullity is therefore at most $2k+2$ the same number as the best known upper bound on Z . The matrix method thus either settles Z completely or cannot settle it at all.

Reaching that ceiling with a rotation-invariant matrix provably confines n to one divisibility class containing at most one prime. Assembling the matrix from tiles that each contribute the identity to the monodromy removes the restriction: tiles of every length in $[L,2L)$ settle every n at least L . I certified 55 such tiles by interval arithmetic.

That is the missing ingredient. Consequently $Z(P(n,2))$ and $Z(P(n,3))$ are determined for every n , the latter for the first time; $Z(P(n,3))=8$ for every n at least 13, proving the conjecture; and $Z(P(n,4))=10$ for every n at least 30.

II. Research Plan

Project. Maximum Nullity of Cyclic Covers of Graphs: A Monodromy Ceiling and the Zero Forcing Number of Generalized Petersen Graphs **Category.** Mathematics **Finalist.** Aryan Padarthy

A. Rationale

Zero forcing began in two unrelated places: controllability of quantum spin systems, and placement of phasor measurement units in electrical grids. It was then connected to linear algebra by the inequality $M(G) \geq Z(G)$, where $M(G)$ is the largest nullity attainable by a symmetric matrix whose off-diagonal support is the edge set of G . The inequality makes Z a combinatorial ceiling on a purely algebraic quantity, and makes lower bounds on Z obtainable by exhibiting a matrix.

The generalized Petersen graphs $P(n,k)$ are a standard cubic test family. Rashidi, Shajareh Poursalavati and Tavakkoli (2020) proved $Z(P(n,k)) \geq 2k+2$ and determined $Z(P(n,2))$. Their Theorem 3.6, asserting $Z(P(n,3)) = 8$ for $n \geq 12$, is false: Krishnan's July 2026 correction note (arXiv:2607.19412) exhibits $Z(P(12,3)) = 7$, the error being a case analysis that never excludes 7-vertex sets. That note leaves open the stabilization threshold and any rigorous lower bound for $k \geq 4$.

This project addresses those two open problems.

B. Research question

1. Is there an intrinsic ceiling on what a matrix certificate can prove about $Z(P(n,k))$, and where does it sit relative to the known upper bound?
2. Can that ceiling be attained, and by an explicitly checkable matrix?
3. Can $Z(P(n,k))$ be determined for **all** n rather than for sparse families of n ?

C. Procedures

This is a theoretical mathematics project. There are no human participants, no vertebrate animals, no potentially hazardous biological agents, and no hazardous chemicals, activities or devices. All work is proof and computation performed by the finalist on a personal computer.

C1 Structure. Eliminate the inner coordinates from the kernel equations of an arbitrary matrix carrying the $P(n,k)$ pattern and identify the resulting scalar recurrence, its order, and the leading coefficients. Establish when the nullity equals the recurrence order.

C2 Construction. For rotation-invariant matrices, reduce singularity to a polynomial condition on a discrete grid and construct matrices meeting the ceiling. Determine exactly when such a matrix is admissible.

C3 Obstruction. Determine whether the divisibility conditions arising in C2 are intrinsic, and prove the answer rather than observing it.

C4 General n. Construct matrices from repeated blocks whose contributions compose, so that one finite family of blocks settles an infinite set of n .

C5 Certification. Verify every constructed matrix at the highest standard available: exact integer or rational arithmetic where possible, and otherwise an interval (Krawczyk) contraction proof that a true real solution exists. Record which standard each claim meets.

C6 Independent checking. Compute Z directly by exhaustive search over vertex subsets for all sizes reachable, validate that solver against every published value, and reproduce each theorem numerically before relying on it.

D. Data analysis

Results are theorems and explicit matrices, not measurements, so analysis is verification rather than statistics. Every computational claim is re-derived by a single script that reports pass or fail per check. Numerical searches are always wrapped in a bound they cannot legally exceed (the ceiling of C1), so that an impossible answer is detected rather than believed. Where a search fails, the residual and the computational budget are both recorded, because a failure at a small budget is not evidence of impossibility.

E. Risk and safety

None. No laboratory work, no fieldwork, no subjects of any kind.

F. Bibliography

1. S. Rashidi, N. Shajareh Poursalavati, M. Tavakkoli. Computing the zero forcing number for generalized Petersen graphs. *J. Algebra Combin. Discrete Struct. Appl.* 7(2) (2020) 183193.
2. A. Krishnan. A correction to the zero forcing number of the generalized Petersen graphs $P(n,3)$. arXiv:2607.19412 (2026).
3. AIM Minimum Rank Special Graphs Work Group. Zero forcing sets and the minimum rank of graphs. *Linear Algebra Appl.* 428 (2008) 16281648.
4. J. H. Conway, A. J. Jones. Trigonometric diophantine equations. *Acta Arith.* 30 (1976) 229240.
5. I. Niven. *Irrational Numbers*. Carus Mathematical Monographs 11, MAA, 1956.
6. R. E. Moore, R. B. Kearfott, M. J. Cloud. *Introduction to Interval Analysis*. SIAM, 2009. (Krawczyk existence test.)

III. Research Paper

The full paper follows this page. It states, for every claim, which of four standards it meets: proved; proved by an exact certificate; proved by interval certification; or numerical. Its Status section lists every claim of the author's own that has been refuted or withdrawn.

Maximum Nullity of Cyclic Covers of Graphs, and Exact Zero Forcing Numbers of Generalized Petersen Graphs

Aryan Padarthi

September 23, 2026

Abstract

For a graph G let $Z(G)$ be its zero forcing number and $M(G)$ the maximum nullity over matrices whose off-diagonal support is exactly $E(G)$, so that $M(G) \leq Z(G)$. We study both for *cyclic covers*: given a base graph B with edge voltages in $\mathbb{Z}_n$, the derived graph B^n on $V(B) \times \mathbb{Z}_n$. Generalized Petersen graphs are the case of a two-vertex base with voltages 1, k and 0.

Our main structural result is that for such a cover the nullity is governed by a single scalar object determined by the voltages. For $P(n, k)$ we prove it for *every* matrix carrying the pattern, symmetric or not: eliminating the inner coordinates leaves one nine-term linear recurrence of order $2k + 2$, so

$$\text{null } A = \dim \ker(T - I) \leq 2k + 2,$$

T the monodromy, with equality iff $T = I$. The order $2k + 2$ is exactly the size of the forcing set produced by a rotation-bootstrap argument that gives $Z(P(n, k)) \leq 2k + 2$, so *the ceiling of the matrix-certificate method coincides with the upper bound it is trying to match*: either the method settles Z completely or nothing does. For a general cover the same reduction holds for equivariant matrices, with the ceiling given by the degree span of $\det M(\zeta)$.

We then reach the ceiling. Under the rotation the singularity condition is a monic degree- $(k + 1)$ polynomial in $s = 2 \cos(2\pi m/n)$ carrying exactly three free parameters for every k ; three roots may be prescribed and no more, but a symbol with rational coefficients drags its whole Galois orbit onto the grid $\{2 \cos(2\pi m/n)\}$. Choosing the symbol to be a product of minimal polynomials of $2 \cos(2\pi/d)$ yields explicit *integer* certificates and hence

$$Z(P(n, 4)) = M(P(n, 4)) = 10, \quad Z(P(n, 5)) = M(P(n, 5)) = 12,$$

for every n divisible by 60, 70 or 90, respectively by 24 — the first exact values of $Z(P(n, k))$ for $k \geq 4$, a range a 2026 correction note records as open. In the second case the certificate is the adjacency matrix itself, and we classify its nullity completely: with $P_k(z) = z^{2k+2} + z^{2k} - z^{k+1} + z^2 + 1$, $\text{null}(\text{Adj } P(n, k)) = \sum \varphi(d)$ over those d with $\Phi_d \mid P_k$ and $d \mid n$, a Conway–Jones phenomenon; in particular the adjacency matrix is never singular for $k = 3, 6, 9$.

Against this we prove a matching *obstruction*. Realisability of a certificate is the positivity (or nonvanishing) of $c^2 = a\alpha - \beta$, and because $L_k(-s) = (-1)^k L_k(s)$, for *every even* k a symbol with two roots $y, -y$ has $c^2 = 0$ identically, while for odd k no such obstruction exists. For $k = 2$, where the symbol has three roots and no residual conditions, this is the whole story in closed form:

$$c^2 = -(s_1 + s_2)(s_1 + s_3)(s_2 + s_3),$$

so the construction fails precisely when two prescribed roots are antipodal. Counting antipodal classes on the grid then settles $k = 2$ for all n : $Z(P(n, 2)) = M(P(n, 2)) = 6$ for every $n \geq 9$ with $n \neq 10$, by an explicit certificate at the three most negative grid values; for $7 \mid n$ no symmetric certificate exists and an integer non-symmetric one of nullity 6 does, the one place where the larger matrix class is needed rather than merely available. The two exceptions are of

opposite kinds — $Z(P(8, 2)) = 5$, so there the method is right to fail, and $Z(P(10, 2)) = 6$, so there it is the method that falls short.

The same analysis is what settles the remaining cases, because only $c^2 = 0$ is fatal: a symbol with $c^2 < 0$ admits no symmetric certificate but is realisable once the two spoke weights are allowed to differ. Redoing our exhaustive enumeration of rational symbols under the weaker condition turns two cases we had recorded as impossible into exact values. Both come from the *signed* adjacency matrix, the adjacency matrix of $P(n, k)$ with one set of spokes negated, whose nullity we classify by the polynomial $Q_k(z) = z^{2k+2} + z^{2k} + z^{k+1} + z^2 + 1$ — the unsigned P_k with the middle sign reversed. Q_k splits into cyclotomics only for $k = 2, 3, 7$ in $2 \leq k \leq 60$, and at $k = 3, 7$ it gives the ceiling:

$$Z(P(n, 3)) = 8 \quad (10 \mid n), \quad Z(P(n, 7)) = 16 \quad (120 \mid n),$$

by matrices of zeros and ± 1 , the second replacing a floating-point claim by an exact one. Conversely $k = 6, 8, 9, 10$ are *impossible* over $\mathbb{Q}$ in both classes, and now for an identified reason: every candidate clearing the family condition is one of at most 25 degenerate symbols, the bound coming from Niven’s theorem.

Breaking the symmetry to period d , the map from weights to the coefficients of the symbol has rank $\min(2d + 1, k + 1)$ in every case computed, so every symbol becomes reachable exactly when $d \geq \lceil k/2 \rceil$. This yields $Z(P(n, 6)) = 14$ at $n = 60, 72, 84, 90, 96, 108, 120$, $Z(P(96, 7)) = 16$ and $Z(P(96, 8)) = 18$, though these certificates are verified numerically rather than in exact arithmetic. Every exact value above holds for n in a divisibility class, and we prove that this is forced: a period-one symbol is a fixed polynomial, its roots are fixed algebraic numbers, and a fixed algebraic number lies on Γ_n only for n in one divisibility class — which contains at most one prime. For $k \geq 3$ an n -dependent choice of roots is constrained by the identity $e_2 = -k$, a Conway–Jones relation on the grid; for $k = 2$ that condition is vacuous, which is exactly why $k = 2$ is general in n and no other k is.

Escaping this needs the non-invariant matrices, and the ceiling corollary supplies the criterion: null $A = 2k + 2$ if and only if the monodromy is the identity, a condition with no arithmetic in it. We certify that condition rigorously, not numerically, by recasting it as the bilinear system $A(w)K(X) = 0$ with K in graph form and applying a Krawczyk contraction test; the square subsystem must be chosen structurally, since $K^\top AK$ is symmetric and its antisymmetric part accounts for exactly $\binom{2k+2}{2}$ dependencies. This gives $Z(P(n, 3)) = M(P(n, 3)) = 8$ at $n = 17, 19, 21, 23, 25, 27$, the primes 17, 19, 23 among them — values that no period-one certificate can reach, and the first in this paper not lying on a congruence class. Finally we obtain a statement for *all* large n . Tiles carrying a fixed docking pattern at both ends have transfer products depending only on their interiors, so identity tiles of every length in $[L, 2L]$ give $Z(P(n, k)) = 2k + 2$ for every $n \geq L$ — a finite list of finite problems. Carrying this out and certifying all 55 tiles involved,

$$Z(P(n, 3)) = M(P(n, 3)) = 8 \quad (n \geq 17), \quad Z(P(n, 4)) = M(P(n, 4)) = 10 \quad (n \geq 30),$$

with no congruence condition. Together with exhaustive search on the finitely many smaller n , this determines $Z(P(n, k))$ for *every* n when $k = 2$ and when $k = 3$ — the latter for the first time, since the previously published value for $k = 3$ is false. In particular $Z(P(n, 3)) = 8$ for every $n \geq 13$ and 13 is optimal. This proves the conjecture of [2], whose stated missing ingredient — “a lower-bound proof valid for all large n ” — is precisely what the tiling construction supplies. With four exhaustive searches at $n = 13, 14, 15, 16$ this closes the $k = 3$ case outright: $Z(P(n, 3)) = 8$ for every $n \geq 13$, and 13 is optimal since $Z(P(11, 3)) = Z(P(12, 3)) = 7$, so the stabilization threshold that the 2026 correction note leaves open is exactly 13. The tile condition is a rational equation, but certifying it needs no interval arithmetic on rational maps: by the ceiling corollary the monodromy of the one-tile cyclic matrix *is* the tile product, so the same bilinear system certifies it.

We state throughout which claims are proved, which rest on exact certificates, which on interval certification, which are numerical, and which earlier claims of ours are corrected or withdrawn.

1 Background

Given a set S of *filled* vertices of a graph G , the colour change rule says: if a filled vertex has exactly one unfilled neighbour, that neighbour becomes filled. Write $\text{cl}(S)$ for the closure of S under repeated application. S is a *zero forcing set* if $\text{cl}(S) = V$, and $Z(G)$ is the least size of one. $Z(G)$ bounds the maximum nullity over the symmetric matrices with G 's off-diagonal pattern, and arose independently in the control of quantum systems and in sensor placement for power networks.

The families we treat all have $2n$ or $4n$ vertices, are cubic, and carry the rotation $\rho : i \mapsto i + 1$ acting on index subscripts.

- $P(n, k)$: outer $u_i u_{i+1}$, inner $v_i v_{i+k}$, spokes $u_i v_i$.
- $I(n, j, k)$: outer $u_i u_{i+j}$, inner $v_i v_{i+k}$, spokes $u_i v_i$; thus $P(n, k) = I(n, 1, k)$.
- $DP(n, k)$: two outer cycles $u_i u_{i+1}$ and $x_i x_{i+1}$, inner edges $w_i y_{i+k}$ and $y_i w_{i+k}$, spokes $u_i w_i$ and $x_i y_i$.

Rashidi, Shajareh Poursalavati and Tavakkoli [1] proved $Z(P(n, k)) \leq 2k+2$ and $Z(P(n, 2)) = 6$ for $n \geq 10$. Their Theorem 3.6, claiming $Z(P(n, 3)) = 8$ for $n \geq 12$, is false: Krishnan [2] exhibits $Z(P(12, 3)) = 7$, the flaw being a case analysis that never excludes 7-vertex sets. That note closes by recording that determining $Z(P(n, k))$ for general k — the stabilization threshold and a rigorous lower bound for $k \geq 4$ — remains open, and conjectures that $Z(P(n, 3)) = 8$ for $n \geq 13$, identifying the missing ingredient as “a lower-bound proof valid for all large n ”. We settle all three: Theorem 55 proves the conjecture and so fixes the threshold at 13, and Theorems 35 and 53 give rigorous lower bounds, indeed exact values, for $k = 4$.

2 The rotation bootstrap

The following two lemmas are elementary; the point is that together they reduce an “all n ” statement to one finite check.

Lemma 1 (Bootstrap). *Let ρ be an automorphism of a finite graph G and let $S \subseteq V(G)$ satisfy $\rho(S) \subseteq \text{cl}(S)$. Then $\rho(\text{cl}(S)) = \text{cl}(S)$.*

Proof. Closure commutes with automorphisms, and is monotone and idempotent, so $\rho(\text{cl}(S)) = \text{cl}(\rho(S)) \subseteq \text{cl}(\text{cl}(S)) = \text{cl}(S)$. As ρ permutes the finite set $V(G)$, the two sets have equal cardinality, hence are equal. $\square$

Lemma 2 (Closing). *Let G be cubic with vertex set $O \sqcup I$ such that every vertex of I has exactly one neighbour in O , every vertex of O has exactly one neighbour in I , and the remaining neighbours of each vertex of O lie in O . If $O \subseteq \text{cl}(S)$ then $\text{cl}(S) = V(G)$.*

Proof. Suppose some $w \in I$ is unfilled and let $o \in O$ be its unique neighbour in O . The other two neighbours of o lie in O , hence are filled, so w is the unique unfilled neighbour of the filled vertex o and is forced. $\square$

To apply these to a family it suffices to exhibit one set S , described uniformly in n , and verify $\rho(S) \subseteq \text{cl}(S)$ by an explicit finite forcing chain. Since ρ 's orbit of any outer vertex is all of O , Lemmas 1 and 2 then give $\text{cl}(S) = V$ and so $Z(G) \leq |S|$, for every n in range.

3 I-graphs

Theorem 3. *For all $j, k \geq 1$ and $n \geq 2(j + k) + 1$, $Z(I(n, j, k)) \leq 2(j + k)$.*

Proof. Put $m = 2(j + k)$ and $S = \{u_0, \dots, u_{m-1}\}$, so $|S| = m$. Since $n \geq m + 1$ the indices $0, \dots, m$ are pairwise distinct mod n .

Step 1. The outer neighbours of u_i are $u_{i \pm j}$, both in S precisely when $j \leq i \leq m - 1 - j$. For such i the spoke neighbour v_i is therefore the unique unfilled neighbour of u_i and is forced. This fills $v_j, \dots, v_{m-1-j}$, that is $v_j, \dots, v_{j+2k-1}$.

Step 2. Consider v_{j+k} . Its neighbours are $u_{j+k} \in S$ (as $j + k \leq m - 1$), v_j , filled in Step 1, and v_{j+2k} , which is not: Step 1 filled indices at most $j + 2k - 1$. So v_{j+k} forces v_{j+2k} .

Step 3. Consider $u_{j+2k} = u_{m-j}$. Its neighbours are $u_{2k} \in S$ (as $2k \leq m - 1$, using $j \geq 1$), the vertex v_{j+2k} filled in Step 2, and u_m , which is unfilled. So u_{j+2k} forces u_m .

Hence $\rho(S) = \{u_1, \dots, u_m\} \subseteq \text{cl}(S)$. By Lemma 1, $\text{cl}(S)$ is ρ -invariant; it contains u_0 , whose ρ -orbit is all of $O = \{u_0, \dots, u_{n-1}\}$, so $O \subseteq \text{cl}(S)$, and Lemma 2 gives $\text{cl}(S) = V$. $\square$

Corollary 4. *For $k \geq 1$ and $n \geq 2k + 3$, $Z(P(n, k)) \leq 2k + 2$.*

Proof. $P(n, k) = I(n, 1, k)$; take $j = 1$ in Theorem 3. $\square$

Remark 5. The corollary is [1, Thm. 2.2]. We note it because that paper is known to contain a false theorem, and because the corrected argument given in [2] is written only for $k = 3$; the proof above is uniform in j and k . Every step of the chain has been verified by computer for $1 \leq j \leq 4$, $1 \leq k \leq 5$ and six values of n per pair (106 cases).

Remark 6 (Scope). Theorem 3 should not be read as $2(j + k)$ new bounds. If $\gcd(n, j) = 1$ then multiplying indices by $j^{-1} \bmod n$ is an isomorphism $I(n, j, k) \cong I(n, 1, j^{-1}k) = P(n, j^{-1}k)$, so those graphs are generalized Petersen graphs and the corollary already applies. We verified this in five cases, e.g. $I(15, 2, 4) \cong P(15, 2)$ with $Z = 6$ on both sides and $I(13, 2, 5) \cong P(13, 4)$ with $Z = 8$. The genuinely new content of Theorem 3 is therefore confined to $\gcd(n, j) > 1$ and $\gcd(n, k) > 1$.

In that regime the bound is usually far from sharp: among the cases we computed it is attained only at $I(18, 2, 3)$ ($Z = 10 = 2(2 + 3)$), while e.g. $I(20, 5, 8)$ has $Z = 10$ against a bound of 26. The reason is structural. When $\gcd(n, j) = d > 1$ the outer vertices form d cycles of length n/d rather than one cycle, and the set $\{u_0, \dots, u_{m-1}\}$ used in the proof spreads across all of them; a smaller, better-adapted set suffices. We record the bound as proved and its sharpness in this regime as open.

4 Double generalized Petersen graphs

Theorem 7. *For all $k \geq 1$ and $n \geq 2k + 3$, $Z(DP(n, k)) \leq 4k + 4$.*

Proof. Put $S = \{u_0, \dots, u_{2k+1}\} \cup \{x_0, \dots, x_{2k+1}\}$, so $|S| = 4k + 4$.

Step 1. For $1 \leq i \leq 2k$ the outer neighbours $u_{i \pm 1}$ of u_i lie in S , so u_i forces w_i ; symmetrically x_i forces y_i . This fills $w_1, \dots, w_{2k}$ and $y_1, \dots, y_{2k}$.

Step 2. The neighbours of w_{k+1} are $u_{k+1} \in S$, y_1 (filled) and y_{2k+1} (not), so w_{k+1} forces y_{2k+1} . Symmetrically y_{k+1} forces w_{2k+1} .

Step 3. The neighbours of u_{2k+1} are $u_{2k} \in S$, w_{2k+1} (filled in Step 2) and u_{2k+2} , so u_{2k+1} forces u_{2k+2} ; symmetrically x_{2k+1} forces x_{2k+2} .

Hence $\rho(S) \subseteq \text{cl}(S)$. The outer set here is $O = \{u_i\} \cup \{x_i\}$, on which ρ acts with two orbits, both met by $\text{cl}(S)$; Lemmas 1 and 2 finish as before. $\square$

No zero forcing results for $DP(n, k)$ were found in the literature; the family is studied for Hamiltonicity, super-connectivity and its automorphism group. Every step above was verified by computer for $1 \leq k \leq 5$ and eight values of n per k .

5 Computed values

All values are exact, by exhaustive search over vertex subsets with a 128-bit mask closure, reduced by ρ : a minimum forcing set may be assumed to contain a fixed orbit representative. The implementation reproduced every published value before use — Krishnan’s table of $Z(P(n, 3))$ for $7 \leq n \leq 20$ (including the corrected $Z(P(12, 3)) = 7$ and the non-monotone values at $n = 10, 11$), $Z(P(n, 2)) = 6$ for $10 \leq n \leq 16$, and $Z(P(2k + 1, k)) = 6$ for $k = 5, 6, 7$ — and was independently cross-checked against a brute-force implementation and two exact solver formulations.




Figure 1: Exact $Z(P(n, k))$ for $k = 4, 5, 6$ against the bound $2k + 2$. Points are separated by whether $\gcd(n, k) = 1$. The sequences are non-monotone: $k = 4$ drops at $n = 12, 16$ and $k = 6$ at $n = 23, 25$.

k	bound $2k + 2$	least n attaining it	least n from which it holds always
2	6	7	9 (it fails at $n = 8$)
3	8	10	13 (it fails at $n = 11, 12$)
4	10	18	not determined
5	12	23	not determined
6	14	28	not determined; $Z(P(29, 6)) = 12$

Table 1: The bound is attained well before it is attained *always*: at $k = 3$ it first holds at $n = 10$, fails again at $n = 11$ and 12 , and holds for every $n \geq 13$ (Theorem 55). The two columns are different quantities and conflating them is easy; at $k = 6$ the bound is attained at $n = 28$ and lost at $n = 29$, so no stabilization threshold is known there. All values by exhaustive search.

Conjecture 8. For each fixed k and each $d \mid k$, the values $Z(P(n, k))$ over n with $\gcd(n, k) = d$ are eventually constant and equal to $2k + 2$.

An earlier form of this conjecture, asserting a single threshold $n_0(k) = 5k - 2$ beyond which $Z(P(n, k)) = 2k + 2$ for *all* n , is false. It fits $k = 3, 4, 5$, and predicted that $n = 28$ would be the first n with $Z(P(n, 6)) = 14$; that prediction, recorded before the computation finished, was borne

out. But $Z(P(29, 6)) = 12$, and $29 > 28$, so 28 is not a stabilization point and the single-threshold form fails at $k = 6$.

The explanation is $\gcd(n, k)$. Sorting the $k = 6$ data by $d = \gcd(n, 6)$:

d	n	$Z(P(n, 6))$
1	13, 17, 19, 23, 25, 29	6, 8, 8, 10, 10, 12
2	14, 16, 20, 22, 26, 28	8, 9, 10, 12, 12, 14
3	15, 21, 27	9, 12, 13
6	18, 24	9, 12

Each class climbs on its own schedule: the class $d = 2$ attains 14 at $n = 28$ while the coprime class has only reached 12 by $n = 29$. At $k = 4$ and $k = 5$ all classes had converged within the computed range, which is why a single threshold appeared to fit; at $k = 6$ they separate. The thresholds 10, 13, 18, 23 for $k = 2, 3, 4, 5$ should therefore be read as the point at which the slowest class happened to arrive, not as values of a formula.

For DP , our computed values give $Z(DP(n, 1)) = 8 = 4 \cdot 1 + 4$ for $6 \leq n \leq 12$ and $Z(DP(n, 2)) = 12 = 4 \cdot 2 + 4$ for $n = 10, 11, 14, 15$, so Theorem 7 is attained at both $k = 1$ and $k = 2$: the bound is sharp there, not merely an upper estimate.

6 Towards a sharper bound for I-graphs

Theorem 3 is stated in terms of j and k , but the outer and inner edge sets of $I(n, j, k)$ decompose into $d_1 = \gcd(n, j)$ and $d_2 = \gcd(n, k)$ cycles respectively, and it is these that appear to govern Z . In the regime where Theorem 3 carries new content (Remark 6), it is markedly non-optimal: $I(24, 9, 10)$ has $Z = 10$ against $2(j + k) = 38$. The proof uses $2(j + k)$ *consecutive* outer vertices, which is wasteful exactly when j is large but $\gcd(n, j)$ is small, since consecutive indices stride across all d_1 outer cycles.

Observation 9. Among the 20 graphs we computed with $d_1, d_2 > 1$, every one satisfies $Z(I(n, j, k)) \leq 2(d_1 + d_2)$, with equality at $I(18, 2, 3)$, $I(18, 3, 8)$, $I(24, 3, 10)$ and $I(24, 9, 10)$.

This must be stated as an observation confined to that regime, and *not* as a bound in general. Taking $j = 1$ gives $d_1 = 1$, and if additionally $\gcd(n, k) = 1$ then $2(d_1 + d_2) = 4$, whereas $Z(P(n, k)) = 2k + 2$ grows without bound: $Z(P(11, 3)) = 7$ and $Z(P(17, 4)) = 9$ both exceed 4. So any bound in terms of d_1, d_2 alone is false, and the true quantity must retain a dependence on j and k . We flag this because the restricted observation is nonetheless suggestive: the four equality cases all have $2(j + k)$ slack by between 12 and 28, so Theorem 3 is certainly not optimal in the regime where it carries new content.

A supporting structural fact, found by computing minimum *bootstrapping* sets (those S with $\rho(S) \subseteq \text{cl}(S)$, which is what a uniform-in- n proof requires): for $I(18, 3, 8)$ the minimum such set has size $10 = 2(d_1 + d_2)$ against $2(j + k) = 22$, and decomposes as three consecutive vertices within each of the $d_1 = 3$ outer cycles together with one inner vertex. Three consecutive vertices in an outer cycle is exactly the configuration that lets the middle one force its spoke, so a proof adapted to the cycle decomposition rather than to consecutive indices should beat Theorem 3. We have not carried this out.

Remark 10 (Limits of verification). Sharpness of Theorem 3 can be tested only for very small (j, k) . The hypothesis $n \geq 2(j + k) + 1$ forces n upward while exhaustive search costs $\binom{2n-1}{2(j+k)-1}$, so

the two constraints compound. On our hardware $(j, k) = (2, 3)$ takes minutes, $(3, 4)$ about ninety, $(2, 5)$ roughly a day, and $(3, 5)$ some 10^4 years. We therefore report sharpness at $(2, 3)$ — where $Z(I(n, 2, 3)) = 10 = 2(j + k)$ for $n = 18, 24$ — and note that the question is computationally out of reach beyond small parameters, rather than merely unfinished.

7 The exact ceiling of the certificate method

Every technique in Section 12 is combinatorial and fails. A matrix-theoretic route succeeds, and produces certificates. Before constructing any, we determine exactly how far the method can possibly go.

Let A be a real matrix indexed by $V(P(n, k))$ with $A_{xy} \neq 0$ iff $xy \in E$ for $x \neq y$, the diagonal unrestricted. We do *not* assume A symmetric: the bound $\text{null } A \leq Z(G)$ needs only that every edge carries a nonzero entry in each direction. Indeed if $Ax = 0$ and x vanishes on a set B , and $u \in B$ has exactly one neighbour $w \notin B$, then row u of $Ax = 0$ reads $A_{uu}x_u + \sum_{v \sim u} A_{uv}x_v = 0$, every term vanishes except $A_{uw}x_w$, and $A_{uw} \neq 0$; so $x_w = 0$, and induction along a forcing process gives $x = 0$ on $\text{cl}(B)$. Hence $x \mapsto x|_S$ is injective on $\ker A$ for every zero forcing set S , so

$$\text{null } A \leq Z(G). \quad (1)$$

Write $M(G)$ for the maximum nullity over symmetric such A and $M_{\text{cs}}(G)$ for the maximum over all of them, so $M \leq M_{\text{cs}} \leq Z$. Set

$$b_i = A_{u_i u_{i+1}}, \quad b'_i = A_{u_{i+1} u_i}, \quad c_i = A_{u_i v_i}, \quad c'_i = A_{v_i u_i}, \quad e_i = A_{v_i v_{i+k}}, \quad e'_i = A_{v_{i+k} v_i},$$

with diagonals a_i at u_i and d_i at v_i ; all of b, b', c, c', e, e' are nonzero.

Theorem 11 (Reduction). *Let $n \geq 2k+3$. For $x \in \mathbb{R}^{\mathbb{Z}_n}$ put $(Lx)_i = -(b'_{i-1}x_{i-1} + a_i x_i + b_i x_{i+1})/c_i$. Then $x \mapsto (x, Lx)$ is a linear bijection from*

$$\left\{ x : \sum_{p=-k-1}^{k+1} \gamma_{i,p} x_{i+p} = 0 \text{ for all } i \in \mathbb{Z}_n \right\}$$

onto $\ker A$, where $\gamma_{i,p} \neq 0$ only for $p \in \{-k-1, -k, -k+1\} \cup \{-1, 0, 1\} \cup \{k-1, k, k+1\}$ — nine positions — and

$$\gamma_{i,k+1} = -\frac{e_i b_{i+k}}{c_{i+k}} \neq 0, \quad \gamma_{i,-k-1} = -\frac{e'_{i-k} b'_{i-k-1}}{c_{i-k}} \neq 0.$$

Consequently, with $T = T_{n-1} \cdots T_0 \in \text{GL}_{2k+2}(\mathbb{R})$ the monodromy of this recurrence on the state $(x_{i-k-1}, \dots, x_{i+k})$,

$$\text{null } A = \dim \ker(T - I), \quad \det T = \frac{(\prod_i b'_i)(\prod_i e'_i)}{(\prod_i b_i)(\prod_i e_i)},$$

so $\det T = 1$ whenever A is symmetric. In particular $\text{null } A \leq 2k+2$, with equality if and only if $T = I$.

Proof. Each spoke weight c_i is nonzero, so the equation of $\ker A$ at u_i solves uniquely for $y_i = (Lx)_i$, a 3-term expression in x . Substituting into the equation at v_i , namely $e'_{i-k}y_{i-k} + d_i y_i + e_i y_{i+k} + c'_i x_i = 0$, gives the displayed relation: the three terms contribute the index triples $\{i-k-1, i-k, i-k+1\}$, $\{i-1, i, i+1\}$ and $\{i+k-1, i+k, i+k+1\}$, and the two extreme coefficients are as stated, both nonzero. Conversely any x satisfying all the relations yields the kernel vector (x, Lx) . Since both

end coefficients are nonzero the relation is a recurrence of order exactly $2k + 2$, solvable forward and backward, so its solutions on $\mathbb{Z}$ form a $(2k + 2)$ -dimensional space on which T acts, and the n -periodic ones are exactly the fixed vectors. Each step matrix T_i is the companion matrix of the relation, with determinant $\gamma_{i,-k-1}/\gamma_{i,k+1}$; taking the product over $i \in \mathbb{Z}_n$, every index set runs over a full residue system and the stated formula follows. $\square$

Corollary 12. $M(P(n, k)) \leq M_{\text{cs}}(P(n, k)) \leq 2k + 2$ for $n \geq 2k + 3$.

Corollary 12 is not a new bound — it follows from (1) and Theorem 3 — but it locates the method precisely, and the identification is exact rather than approximate. The recurrence has order $2k + 2$ because a nonzero kernel vector cannot vanish on $2k + 2$ consecutive outer indices; and $\{u_0, \dots, u_{2k+1}\}$ is exactly the zero forcing set that the rotation bootstrap of Theorem 3 produces. So *the order of the recurrence equals the size of that forcing set*, and the ceiling of the certificate method coincides with the upper bound the method is trying to match, with no slack on either side. A certificate can prove $Z(P(n, k)) \geq 2k + 2$ and, by Theorem 11, it does so precisely when the monodromy is the identity. Our earlier description of the method as “capped near 6–8” confused a property of the constructions we had tried with a property of the method.

Proposition 13 (Product form). *For symmetric A with $C = \text{diag}(c_i)$, put $U = A[O, O]$ and $S = C^{-1}A[I, I]C^{-1}$. Then $\text{null } A = \dim \ker(SU - I)$. Here U and S are symmetric with the outer-cycle and inner-cycle patterns respectively, free diagonals and nonzero prescribed off-diagonal entries, and every such pair arises. Hence $M(P(n, k))$ is the largest possible geometric multiplicity of the eigenvalue 1 in a product of two such matrices.*

Proof. $\ker A$ is cut out by $Ux + Cy = 0$ and $Cx + A[I, I]y = 0$; eliminating y gives $(C - A[I, I]C^{-1}U)x = 0$, and factoring out the invertible C turns this into $(I - SU)x = 0$. $\square$

Theorem 11 and Proposition 13 were verified computationally before use: over 63 pairs (n, k) with $2 \leq k \leq 7$ and three random matrices each (189 matrices), the nine-position sparsity, both end-coefficient formulas and the identity $\text{null } A = \dim \ker(T - I) \leq 2k + 2$ hold in every case; over 30 further pairs and both matrix classes, so do the determinant formula, the product form, and the statement that no kernel vector vanishes on $u_0, \dots, u_{2k+1}$. See `verification/recurrence_order.py`.

8 Period-one certificates and cyclotomy

Take $A \in S(P(n, k))$ invariant under ρ , normalised so that the outer weight is 1:

$$U = aI + (P + P^{-1}), \quad V = dI + e(P^k + P^{-k}), \quad C = cI,$$

with $c, e \neq 0$. Diagonalising P splits A into the n blocks $M_m = \begin{pmatrix} a+s_m & c \\ c & d+et_m \end{pmatrix}$ with $s_m = 2\cos(2\pi m/n)$ and $t_m = 2\cos(2\pi km/n)$. Each singular block has nullity exactly 1, since $c \neq 0$, so $\text{null } A = \#\{m : \det M_m = 0\}$.

Proposition 14 (The symbol). *Let $L_k \in \mathbb{Z}[s]$ be the monic polynomial of degree k with $L_k(z + z^{-1}) = z^k + z^{-k}$ (so $L_1 = s$, $L_2 = s^2 - 2$, $L_k = sL_{k-1} - L_{k-2}$). Then $t_m = L_k(s_m)$ exactly, and with $\alpha = d/e$, $\beta = (ad - c^2)/e$,*

$$\det M_m = 0 \iff F(s_m) = 0, \quad F(s) = (s + a)L_k(s) + \alpha s + \beta.$$

F is monic of degree $k + 1$ and carries exactly three free parameters (a, α, β) , for every k . Writing $\Gamma_n = \{2\cos(2\pi m/n)\}$, each root of F in $\Gamma_n \cap (-2, 2)$ contributes 2 to $\text{null } A$ and each root at ± 2 contributes 1, because $m \mapsto n - m$ fixes s_m and is free elsewhere. Finally $c^2 = \alpha\alpha - \beta$ after scaling $e = 1$, so the certificate is admissible precisely when $\alpha\alpha - \beta > 0$.

Remark 15 (A claim of ours, corrected). We previously asserted that this period-one construction “never exceeds 6” for every $k \geq 2$, reasoning that only three roots can be prescribed. The premise is right and the conclusion is wrong. At a root r the condition is *linear*,

$$a L_k(r) + \alpha r + \beta = -r L_k(r), \quad (2)$$

so three distinct roots determine (a, α, β) and no fourth can be imposed. But prescribing three roots also *determines the other* $k - 2$, and those may lie in Γ_n of their own accord. They do exactly when arithmetic puts them there: if F has rational coefficients its roots form Galois orbits, and the Galois orbit of $2 \cos(2\pi m/n)$ consists again of elements of Γ_n . A single rational F can therefore place all $k + 1$ roots on the grid at once. This raises the period-one ceiling from 6 to $2k + 2$, and makes the more elaborate period-two construction we reported unnecessary.

When realisability fails, exactly

The parameter count says what can be prescribed; realisability says what survives. The condition $c^2 = a\alpha - \beta \neq 0$ turns out to be a *factorisation* condition on the symbol, and this settles the question completely.

Theorem 16 (Realisability). *For a period-one matrix with symbol F ,*

$$c^2 = 0 \iff F = (s + a)(L_k(s) + t) \text{ for some constant } t.$$

Consequently the matrix is admissible in the combinatorially symmetric class if and only if F does not factor in this way, and in the symmetric class if and only if in addition $c^2 > 0$.

Proof. $c^2 = a\alpha - \beta$, so $c^2 = 0$ iff $\beta = a\alpha$, and then

$$F = (s + a)L_k + \alpha s + a\alpha = (s + a)L_k + \alpha(s + a) = (s + a)(L_k + \alpha).$$

Conversely if $F = (s + a')(L_k + t)$ then comparing the coefficient of s^k gives $a' = a$, and expanding gives $\alpha = t$, $\beta = at$, so $c^2 = at - at = 0$. The two admissibility statements are Proposition 14: the product of the two spoke weights is c^2 , which must be nonzero for the pattern to be that of $P(n, k)$, and must be a square for the matrix to be symmetric. $\square$

Over the rationals the degenerate symbols form a finite, explicit list, and the reason is Niven’s theorem [4]: the only rational values of $2 \cos(\pi q)$ with $q \in \mathbb{Q}$ are $0, \pm 1, \pm 2$.

Corollary 17 (The degenerate rational symbols). *Let $F \in \mathbb{Q}[s]$ be monic of degree $k + 1$ with every root in Γ_n , and suppose $c^2 = 0$. Then $F = (s + a)(L_k + t)$ with*

$$a, t \in \{0, \pm 1, \pm 2\}.$$

There are therefore at most 25 degenerate rational symbols for each k .

Proof. By Theorem 16, $F = (s + a)(L_k + t)$, and a is the coefficient of s^k in F , hence rational; t is then rational too. Now $-a$ is a root of F , so $-a = 2 \cos \theta$ for some rational multiple θ of 2π , and a rational forces $a \in \{0, \pm 1, \pm 2\}$ by Niven’s theorem. Likewise any root r of $L_k + t$ lies in Γ_n , so $r = 2 \cos \varphi$ with φ a rational multiple of 2π and $t = -L_k(r) = -2 \cos(k\varphi)$, again rational, so $t \in \{0, \pm 1, \pm 2\}$. $\square$

Enumerating the 25 pairs and keeping those for which $(s + a)(L_k + t)$ splits into distinct Ψ_d with every root interior gives 7, 8, 7, 6, 9 degenerate symbols for $k = 2, \dots, 6$, agreeing exactly with the count of $c^2 = 0$ cases found by the independent exhaustive enumeration in each case (`verification/cs_exhaustive.py`).

One special case of Theorem 16 is worth isolating, because it is visible on the root set alone and depends only on the parity of k .

Corollary 18 (Antipodal obstruction). *Let k be even and suppose the symbol F has two roots y and $-y$ with $y \neq 0$. Then $c^2 = a\alpha - \beta = 0$, so the matrix is not admissible in the symmetric class and not admissible in the combinatorially symmetric class either. Hence for even k a period-one certificate requires a symbol whose root set contains no antipodal pair. For odd k no such obstruction exists: there*

$$c^2 = \frac{L_k(y)(y^2 - a^2)}{y},$$

which is not identically zero.

Proof. From $L_k(z + z^{-1}) = z^k + z^{-k}$ we get $L_k(-s) = (-1)^k L_k(s)$, since $s \mapsto -s$ corresponds to $z \mapsto -z$. Let k be even, so $L_k(-y) = L_k(y)$. Writing out $F(y) = 0$ and $F(-y) = 0$ using $F(s) = (s + a)L_k(s) + \alpha s + \beta$ gives

$$aL_k(y) + \alpha y + \beta = -yL_k(y), \quad aL_k(y) - \alpha y + \beta = yL_k(y).$$

Adding gives $\beta = -aL_k(y)$; subtracting and dividing by $2y \neq 0$ gives $\alpha = -L_k(y)$. Therefore $c^2 = a\alpha - \beta = -aL_k(y) + aL_k(y) = 0$. Admissibility requires the spoke weights to be nonzero, and their product is c^2 by Proposition 14, which rules out both classes. For odd k , $L_k(-y) = -L_k(y)$, the same two equations become $aL_k(y) + \alpha y + \beta = -yL_k(y)$ and $-aL_k(y) - \alpha y + \beta = -yL_k(y)$, whence $\beta = -yL_k(y)$ and $\alpha = -aL_k(y)/y$, and $c^2 = a\alpha - \beta = L_k(y)(y^2 - a^2)/y$. $\square$

For $k = 2$ the obstruction is the *only* one, because the symbol then has exactly three roots and the $k - 2$ residual conditions of Theorem 29 are vacuous.

Proposition 19 ($k = 2$ in closed form). *For $k = 2$, prescribing three distinct roots s_1, s_2, s_3 determines $a = -e_1$, $\alpha = e_2 + 2$, $\beta = -e_3 - 2e_1$ with e_i the elementary symmetric functions, and*

$$c^2 = e_3 - e_1 e_2 = -(s_1 + s_2)(s_1 + s_3)(s_2 + s_3).$$

So $c^2 = 0$ if and only if two of the roots are antipodal, and $c^2 > 0$ if and only if an odd number of the three pairwise sums is negative.

Proof. Expanding $F = (s + a)(s^2 - 2) + \alpha s + \beta = s^3 + as^2 + (\alpha - 2)s + (\beta - 2a)$ and comparing with $\prod_i (s - s_i) = s^3 - e_1 s^2 + e_2 s - e_3$ gives the three formulas, and then $a\alpha - \beta = -e_1(e_2 + 2) + e_3 + 2e_1 = e_3 - e_1 e_2$. The identity $(s_1 + s_2)(s_1 + s_3)(s_2 + s_3) = e_1 e_2 - e_3$ is elementary. The two consequences follow because the three factors are the pairwise sums. $\square$

Write $\Gamma_n^\circ = \Gamma_n \cap (-2, 2)$, with $M = \lfloor (n - 1)/2 \rfloor$ elements $s_1 > \dots > s_M$. For odd n , Γ_n° contains no antipodal pair, since $s_i = -s_j$ forces $i + j = n/2$. For even n the antipodal classes number $n/4$ when $4 \mid n$ and $(n - 2)/4$ otherwise, so three pairwise non-antipodal values exist exactly when n is odd with $n \geq 7$ or n is even with $n \geq 12$. This settles $k = 2$ completely.

Theorem 20 ($k = 2$, all n). *Take the three most negative elements of Γ_n° , that is s_M, s_{M-1}, s_{M-2} . Then*

$$s_{M-1} + s_{M-2} = \begin{cases} -4 \cos(4\pi/n) \cos(\pi/n), & n \text{ odd}, \\ -4 \cos(5\pi/n) \cos(\pi/n), & n \text{ even}, \end{cases}$$

which is negative exactly for odd $n \geq 9$ and for even $n \geq 12$. In those cases all three pairwise sums are negative, so $c^2 > 0$ by Proposition 19 and

$$Z(P(n, 2)) = M(P(n, 2)) = 6 \quad (n \geq 9, n \neq 10).$$

For n divisible by 7 the symbol $\Psi_7 = s^3 + s^2 - 2s - 1$ gives $c^2 = -1$: no symmetric certificate, but the integral combinatorially symmetric matrix

$$A = \begin{pmatrix} I + P + P^{-1} & I \\ -I & P^2 + P^{-2} \end{pmatrix}$$

has nullity exactly 6, so $Z(P(n, 2)) = M_{\text{cs}}(P(n, 2)) = 6$ whenever $7 \mid n$. The remaining n are finite in number and exhaustive computation gives $Z(P(5, 2)) = 5$, $Z(P(6, 2)) = 4$, $Z(P(8, 2)) = 5$ and $Z(P(10, 2)) = 6$.

Proof. For odd n , $M = (n-1)/2$, so $2\pi(M-1)/n = \pi - 3\pi/n$ and $2\pi(M-2)/n = \pi - 5\pi/n$, giving $s_{M-1} + s_{M-2} = -2(\cos(3\pi/n) + \cos(5\pi/n)) = -4 \cos(4\pi/n) \cos(\pi/n)$, negative iff $\cos(4\pi/n) > 0$ iff $n > 8$. For even n , $M = n/2 - 1$ and the same computation with $4\pi/n, 6\pi/n$ gives $-4 \cos(5\pi/n) \cos(\pi/n)$, negative iff $n > 10$. Since s is decreasing, $s_M \leq s_{M-2}$, so $s_M + s_{M-1} \leq s_{M-1} + s_{M-2} < 0$ and $s_M + s_{M-2} \leq s_{M-1} + s_{M-2} < 0$: all three pairwise sums are negative, their product is negative, and $c^2 > 0$. Proposition 14 then gives $\text{null } A = 6$, and $6 \leq M(P(n, 2)) \leq Z(P(n, 2)) \leq 6$ by the rotation bootstrap. The Ψ_7 claim is Theorem 29 with c^2 negative, read in the combinatorially symmetric class where the two spoke weights need only have the prescribed product. $\square$

Remark 21 (Exact but stiff). The certificate of Theorem 20 is exact for every admissible n , but it is numerically delicate for large n , and it is worth saying why so that a floating-point check is not misread. The three prescribed roots are $-2 \cos(\pi/n)$, $-2 \cos(3\pi/n)$, $-2 \cos(5\pi/n)$ for odd n , all within $O(1/n^2)$ of -2 : at $n = 1001$ they span 2.4×10^{-4} . So F is close to having a triple root, the matrix is ill-conditioned, and grid values merely *near* the roots produce small singular values — at $n = 1001$ the true roots give 9×10^{-16} and the nearest non-roots 4×10^{-11} , so a fixed relative threshold over-counts the nullity. None of this touches the proof, which is symbolic: $c^2 > 0$ follows from the product-to-sum identity and the nullity is exactly 6 because F has exactly three distinct interior roots on Γ_n , each met by two Fourier indices.

Remark 22 (What the two exceptions mean). At $n = 8$ and $n = 10$ the grid is too small to avoid antipodal pairs — $\Gamma_8^\circ = \{\sqrt{2}, 0, -\sqrt{2}\}$ and Γ_{10}° has four elements in two antipodal classes, so by pigeonhole every triple contains such a pair. The two cases then behave differently: $Z(P(8, 2)) = 5$, so there the method is right to fail, while $Z(P(10, 2)) = 6$, so there it is the method that falls short and not the bound. This is the sharpest available statement of the limitation, and it is a statement about Γ_n alone.

Corollary 23 (Excluded cyclotomic symbols, even k). *Let k be even. The roots of Ψ_{2d} are the negatives of the roots of Ψ_d for odd d , and the root set of Ψ_d is closed under negation whenever $4 \mid d$ and $\varphi(d) \geq 4$. Hence in Theorem 29 the factor set D may contain neither both d and $2d$ for any odd d , nor any d divisible by 4 with $\varphi(d) \geq 4$.*

Corollary 23 in fact accounts for *all* of the rejections in the exhaustive enumeration of Proposition 37. At $k = 2$ that enumeration found 17 factor sets in the three-parameter family, of which 5 had $c^2 > 0$ and 12 were rejected. Those 12 split exactly:

$$\underbrace{(3, 8), (3, 12), (4, 8), (4, 12), (6, 8), (6, 12), (3, 4, 6)}_{c^2=0, \text{ every one antipodal}} \quad \underbrace{(7), (9), (4, 5), (5, 6), (6, 10)}_{c^2=-1, \text{ none antipodal}}.$$

The seven on the left are dead in both matrix classes, and Theorem 18 identifies them without computing c^2 at all. The five on the right carry $c^2 = -1$, so each of them *is* realisable in the combinatorially symmetric class. The admissible count at $k = 2$ is therefore 5 symmetrically and 10 in the larger class. The extra symbols cover every n divisible by 7 or by 9; of these only $n = 7$ is out of reach of Theorem 20, and it is precisely the case the larger class was needed for. The antipodal obstruction is the only thing standing in the way of any rational period-one symbol for $k = 2$. Every prediction of Theorem 18 and Corollary 23 was checked against 50 candidate factor sets at $k = 2, 4$ with no violation, and Proposition 19 was verified by exact arithmetic in $\mathbb{Z}[x]/\Phi_n(x)$ over all $\binom{M}{3}$ triples for every $7 \leq n \leq 60$ — 11,000 triples, with $c^2 = 0$ and “contains an antipodal pair” agreeing in every one. See `verification/antipodal_obstruction.py`.

The larger matrix class, and $k = 3$

Theorem 16 separates $c^2 = 0$ from $c^2 < 0$, and that separation matters: only the first is fatal. A symbol with $c^2 < 0$ admits no symmetric certificate, but in the combinatorially symmetric class the two spoke weights are independent and need only have product c^2 , so it is admissible there. Redoing the exhaustive enumeration of Proposition 37 with the weaker condition $c^2 \neq 0$ therefore changes the answer. The case $k = 3$ changes the most.

Theorem 24 ($k = 3$). *Let $F = \Psi_5 \Psi_{10} = (s^2 + s - 1)(s^2 - s - 1) = s^4 - 3s^2 + 1$. Then $F = sL_3(s) + 1$, so $a = \alpha = 0$, $\beta = 1$ and $c^2 = -1$. Hence the integer matrix*

$$A = \begin{pmatrix} P + P^{-1} & I \\ -I & P^3 + P^{-3} \end{pmatrix}$$

has zero diagonal, carries exactly the pattern of $P(n, 3)$, and has nullity exactly $8 = 2k + 2$ for every n divisible by 10. Therefore

$$Z(P(n, 3)) = M_{\text{cs}}(P(n, 3)) = 8 \quad (10 \mid n).$$

Proof. $\Psi_5 = s^2 + s - 1$ and $\Psi_{10} = s^2 - s - 1$, whose product is $s^4 - 3s^2 + 1 = s(s^3 - 3s) + 1 = sL_3(s) + 1$. Matching against $(s + a)L_3 + \alpha s + \beta$ gives $a = 0$, $\alpha = 0$, $\beta = 1$, so $c^2 = a\alpha - \beta = -1$. By Theorem 16 the symbol is non-degenerate, and taking the two spoke weights to be 1 and -1 realises $c^2 = -1$ with both nonzero, so the zero–nonzero pattern is exactly that of $P(n, 3)$. Every root of F is a root of Ψ_5 or Ψ_{10} , hence lies in Γ_n whenever $10 \mid n$, and all four roots are interior, so each contributes 2: null $A = 8$. The upper bound $Z(P(n, 3)) \leq 8$ is the rotation bootstrap, and $M_{\text{cs}} \leq Z$ holds because the forcing argument uses only the zero–nonzero pattern. $\square$

Remark 25 (Two things this certificate illustrates). First, A is the adjacency matrix of $P(n, 3)$ with one set of spokes negated, and the sign matters entirely: the unsigned adjacency matrix has symbol $sL_3 - 1$ and nullity 0 for every n divisible by 10, consistent with Theorem 30. Second, the four roots of F are $\pm\varphi$ and $\pm\varphi^{-1}$ — *two antipodal pairs*. By Corollary 18 such a symbol would be degenerate for every even k ; because 3 is odd it is not. So the parity corollary does not only rule things out, it says where to look.

The same enumeration in the larger class yields five further symbols at $k = 2$ (covering n divisible by 7, 9, 20 or 30), five more at $k = 3$ beyond Theorem 24 (n divisible by 24, 40 or 42), three at $k = 4$ (60, 70, 90) and two at $k = 5$ (126); all have $c^2 \in \{-1, -2\}$. At $k = 6$ it yields nothing, and now for a reason.

The signed adjacency matrix, and $k = 7$

Both the $k = 3$ certificate and one of the $k = 7$ symbols found by this enumeration have $a = \alpha = 0$ and $\beta = 1$, that is, symbol $F = sL_k(s) + 1$. That is not a coincidence, and identifying what it means gives an exact value at $k = 7$.

With $a = \alpha = 0$ both diagonals vanish, so A is the adjacency matrix of $P(n, k)$ with one set of spokes negated. Substituting $s = z + z^{-1}$, $L_k = z^k + z^{-k}$ and clearing denominators turns $F(s) = 0$ into

$$Q_k(z) := z^{2k+2} + z^{2k} + z^{k+1} + z^2 + 1 = 0,$$

which is exactly the adjacency polynomial $P_k(z) = z^{2k+2} + z^{2k} - z^{k+1} + z^2 + 1$ of Theorem 30 with the sign of the middle term reversed — the spoke sign flip, visible in the symbol.

Theorem 26 (Signed adjacency classification). *Let A^- be the adjacency matrix of $P(n, k)$ with one set of spokes negated. Then*

$$\text{null } A^- = \sum \varphi(d) \quad \text{over those } d \text{ with } \Phi_d \mid Q_k \text{ and } d \mid n,$$

by the argument of Theorem 30 applied to Q_k in place of P_k . In particular $\text{null } A^- = 2k + 2$ for suitable n exactly when Q_k is a product of cyclotomic polynomials.

So the question of which k admit an *integer* certificate of this shape becomes: for which k does Q_k split into cyclotomics? This is decidable for each k , and the answer is rigid. For $2 \leq k \leq 60$ it happens only for

$$k = 2 : Q_2 = \Phi_5\Phi_6, \quad k = 3 : Q_3 = \Phi_5\Phi_{10}, \quad k = 7 : Q_7 = \Phi_5\Phi_{10}\Phi_{24},$$

with $\sum \varphi(d)$ equal to 6, 8, 16 — the ceiling $2k + 2$ in every case — and lcm equal to 30, 10, 120. The unsigned P_k splits only for $k = 2$ and $k = 5$ (Theorem 30), so between them the two sign choices give integer adjacency-type certificates exactly at $k \in \{2, 3, 5, 7\}$.

Theorem 27 ($k = 7$). $Q_7 = \Phi_5\Phi_{10}\Phi_{24}$, so for every n divisible by 120 the integer matrix

$$A = \begin{pmatrix} P + P^{-1} & I \\ -I & P^7 + P^{-7} \end{pmatrix}$$

carries the pattern of $P(n, 7)$ and has nullity exactly $16 = 2k + 2$. Hence

$$Z(P(n, 7)) = M_{\text{cs}}(P(n, 7)) = 16 \quad (120 \mid n).$$

This replaces a numerical claim by an exact one. Section 10 reports $Z(P(96, 7)) = 16$ via a period-4 certificate verified in floating point; Theorem 27 gives the same value for infinitely many n from a matrix of zeros and ± 1 , with the nullity computed by a cyclotomic factorisation rather than by an eigenvalue. Confirmed numerically: the matrix has nullity 16 at $n = 120$, and 6, 8 at the corresponding n for $k = 2, 3$ (`verification/signed_adjacency.py`).

Theorem 28 ($k = 6$ is impossible over $\mathbb{Q}$). *Of the 225 products $\prod_{d \in D} \Psi_d$ of total degree 7 with every Ψ_d of interior roots, 216 do not lie in the three-parameter family, and all 9 that do are degenerate in the sense of Theorem 16: each is $(s + a)(L_6 + t)$ with $a \in \{0, \pm 1\}$ and $t \in \{0, \pm 1\}$, namely Ψ_d -products over $D = (3, 8, 24), (4, 8, 24), (6, 8, 24), (3, 9, 18), (4, 9, 18), (6, 9, 18), (3, 36), (4, 36), (6, 36)$. Hence no rational period-one symbol attains 14 for $k = 6$, in either matrix class.*

This is stronger than the statement it replaces. Where we previously reported that a search found nothing at $k = 6$, the obstruction now identifies every candidate that survives the family condition and shows each one is degenerate, and Corollary 17 bounds the degenerate list *a priori* at 25 polynomials. The $k \geq 6$ certificates of Section 10 therefore have to break the symmetry, and they do.

For $d \geq 3$ let $\Psi_d \in \mathbb{Z}[s]$ be the minimal polynomial of $2 \cos(2\pi/d)$, of degree $\varphi(d)/2$, whose roots are the primitive values $2 \cos(2\pi j/d)$ with $\gcd(j, d) = 1$; put $\Psi_1 = s - 2$, $\Psi_2 = s + 2$. Every root of Ψ_d lies in Γ_n as soon as $d \mid n$.

Theorem 29 (Cyclotomic certificates). *Let D be a set of distinct positive integers with $\sum_{d \in D} \deg \Psi_d = k + 1$ and put $F = \prod_{d \in D} \Psi_d$. Write $G = F - sL_k$, let a be the coefficient of s^k in G , and suppose $G - aL_k$ has degree at most 1, say $G - aL_k = \alpha s + \beta$; that is $k - 2$ linear conditions with integer coefficients. Put $e = 1$, $d = \alpha$, $c^2 = a\alpha - \beta$, and suppose $c^2 > 0$. Then for every n divisible by $\text{lcm}(D)$ with $n \geq 2k + 3$, the matrix*

$$A = \begin{pmatrix} aI + P + P^{-1} & cI \\ cI & \alpha I + P^k + P^{-k} \end{pmatrix}$$

lies in $S(P(n, k))$ and has nullity $2 \sum_{d \in D, d \geq 3} \deg \Psi_d + \#\{d \in D : d \leq 2\}$. If every $d \in D$ is at least 3 this equals $2k + 2$, and then

$$Z(P(n, k)) = M(P(n, k)) = 2k + 2.$$

Proof. By construction $F = (s + a)L_k + \alpha s + \beta$ as an identity in $\mathbb{Z}[s]$, so F is the symbol of A by Proposition 14, and $A \in S(P(n, k))$ because $c^2 > 0$ makes c real and nonzero while the edge weights 1 are nonzero. Each $d \in D$ divides n , so all $k + 1$ roots of F lie in Γ_n ; they are distinct because the Ψ_d are distinct irreducible polynomials. Counting multiplicities as in Proposition 14 gives the stated nullity, and the last line combines it with (1) and Theorem 3. $\square$

A search over $k \leq 12$ and $d \leq 210$ (`verification/cyclotomic_certificates.py`) gives the following. Every certificate listed has $a, \alpha, c \in \mathbb{Z}$, hence is an *integer* matrix, and its nullity was confirmed by exact rational Gaussian elimination.

Two of these deserve to be written out. For $k = 2$ and $k = 5$ the parameters are $a = d = 0$, $c = 1$: the certificate is the *adjacency matrix* of $P(n, k)$, and the symbol is $sL_k(s) - 1$, with

$$sL_2 - 1 = (s + 1)(s^2 - s - 1) = \Psi_3 \Psi_{10}, \quad sL_5 - 1 = (s - 1)(s + 1)(s^4 - 4s^2 + 1) = \Psi_6 \Psi_3 \Psi_{24}.$$

Equivalently, since $\det M_m = s_m t_m - 1$,

$$\text{null}(\text{Adj } P(n, k)) = \#\left\{m \in \mathbb{Z}_n : 2 \cos \frac{2\pi(k+1)m}{n} + 2 \cos \frac{2\pi(k-1)m}{n} = 1\right\},$$

a rational linear relation among cosines of rational angles — the setting of the Conway–Jones classification. That question has a complete answer, given in Theorem 30 below. (An earlier version of this paper asserted that “in the computed range $k \leq 10$ only $k = 2, 5, 8$ ” admit solutions. That was drawn from a scan filtered at nullity ≥ 6 and is wrong: $k = 4, 7, 10, 11, 12$ also admit

k	D	$\text{lcm}(D)$	F	a	d	c^2	nullity
2	$\{3, 10\}$	30	$sL_2 - 1$	0	0	1	$6 = 2k + 2$
2	$\{14\}$	14	Ψ_{14}	-1	0	1	$6 = 2k + 2$
3	$\{1, 4, 5\}$	20	$\Psi_1\Psi_4\Psi_5$	-1	-1	1	7
4	$\{4, 30\}$	60	$(s+1)L_4 - s - 2$	1	-1	1	$10 = 2k + 2$
4	$\{5, 14\}$	70	$\Psi_5\Psi_{14}$	0	1	1	$10 = 2k + 2$
4	$\{5, 18\}$	90	$\Psi_5\Psi_{18}$	1	0	1	$10 = 2k + 2$
5	$\{3, 6, 24\}$	24	$sL_5 - 1$	0	0	1	$12 = 2k + 2$

Table 2: Exact period-one certificates. The $k = 2$, $k = 4$ and $k = 5$ rows give $Z(P(n, k)) = M(P(n, k)) = 2k + 2$ for every n divisible by the stated modulus.

solutions, with smaller nullity.) For $k = 5$ the count is $12 = 2k + 2$ whenever $24 \mid n$: the plain adjacency matrix of $P(n, 5)$ has nullity 12, which settles Z for that family. For $k = 4$ the symbol is $(s+1)L_4(s) - s - 2 = s(s^4 + s^3 - 4s^2 - 4s + 1) = \Psi_4\Psi_{30}$, giving the integer certificate with $a = 1$, $d = -1$, $c = 1$.

The adjacency nullity, classified completely

Since $\det M_m = s_m t_m - 1$, the adjacency matrix is singular at m exactly when $z = \omega^m$ satisfies $z^{k+1} + z^{-(k+1)} + z^{k-1} + z^{-(k-1)} = 1$. Clearing denominators gives an integer polynomial, and the question becomes which of its roots are roots of unity.

Theorem 30. Put $P_k(z) = z^{2k+2} + z^{2k} - z^{k+1} + z^2 + 1 \in \mathbb{Z}[z]$. Then

$$\text{null}(\text{Adj } P(n, k)) = \sum_{\substack{d: \Phi_d \mid P_k \\ d \mid n}} \varphi(d),$$

the sum over those d whose cyclotomic polynomial divides P_k . The list of such d is finite and computable once per k .

Proof. m contributes iff ω^m is a common root of P_k and $z^n - 1$, i.e. a root of $\gcd(P_k, z^n - 1)$. The roots of unity among the roots of P_k are exactly the roots of its cyclotomic factors Φ_d , and Φ_d 's roots are n -th roots of unity precisely when $d \mid n$; each such Φ_d contributes its $\varphi(d)$ roots. Finiteness is immediate from $\deg P_k = 2k + 2$, and is the elementary shadow of the Conway–Jones theorem on rational linear relations among cosines of rational angles. $\square$

k	$2k + 2$	cyclotomic factors of P_k	lcm	nullity when $\text{lcm} \mid n$
2	6	Φ_3, Φ_{10}	30	$6 = 2k + 2$
3	8	none	—	0 for every n
4	10	Φ_3	3	2
5	12	$\Phi_3, \Phi_6, \Phi_{24}$	24	$12 = 2k + 2$
6	14	none	—	0 for every n
7	16	Φ_3, Φ_6	6	4
8	18	Φ_3, Φ_{10}	30	6

So the adjacency matrix of $P(n, k)$ is *never* singular for $k = 3, 6, 9$, whatever n ; and it attains the full ceiling $2k+2$ only for $k = 2$ and $k = 5$. The formula was checked against a direct count for $k = 2, 5, 8$ at every $13 \leq n < 200$, with no discrepancy (`verification/adjacency_classification.py`).

9 The general setting: cyclic covers

Nothing in Section 8 used the particular shape of $P(n, k)$ beyond the presence of the rotation. It is worth saying what the natural setting is.

Let B be a finite graph whose edges carry *voltages* in $\mathbb{Z}_n$ (no edge of voltage 0 at a loop, since that would lift to a self-loop). The *derived graph* B^n has vertex set $V(B) \times \mathbb{Z}_n$, with (x, i) joined to $(y, i + v)$ for each edge $x - y$ of voltage v . This is the standard cyclic cover of topological graph theory. Taking B to be two vertices u, v with a loop of voltage 1 at u , a loop of voltage k at v , and an edge $u - v$ of voltage 0 recovers $B^n = P(n, k)$ exactly.

Theorem 31 (Cover reduction). *Let A carry the pattern of B^n and commute with the $\mathbb{Z}_n$ action. Then A block-diagonalises into n blocks of size $|V(B)|$,*

$$M_m[x, y] = \sum_{\text{edges } x \rightarrow y \text{ of voltage } v} w_{xy} \omega^{mv} + \delta_{xy} a_x,$$

and $\text{null } A = \sum_m \text{null } M_m$. Writing $\det M(\zeta)$ as a Laurent polynomial in ζ , the number of m at which a block is singular — hence $\text{null } A$ — is at most the degree span of $\det M(\zeta)$, which is determined by the voltages of B alone.

For $P(n, k)$ the degree span is $2k + 2$, recovering Corollary 12 in the equivariant case. The reduction and the ceiling were verified against independently constructed full matrices, with weights tuned so that the nullity is genuinely nonzero, on seven base graphs on two to five vertices with assorted voltages (`verification/cyclic_covers.py`).

Two caveats. Theorem 31 is stated for *equivariant* A ; the stronger statement of Theorem 11, which bounds the nullity of *every* matrix carrying the pattern, we have proved only for $P(n, k)$, where the nonzero spoke weights permit the elimination. Whether the corresponding elimination can be carried out over an arbitrary base graph — and so whether the ceiling is a property of the cover rather than of equivariant matrices on it — we leave open. Second, we have tested the non-equivariant bound on general covers by *maximising* the nullity rather than sampling — random matrices are generically nonsingular, so sampling proves nothing. Optimising over fully non-equivariant weights on four small covers, the largest attainable nullity never exceeded the degree span, and equalled it for the theta graphs with voltages $(0, 1, 2)$ and $(0, 1, 3)$; on the $P(n, 2)$ and $P(n, 3)$ covers at $n = 9$ it fell short, as it must, since $Z(P(9, k)) < 2k + 2$ there. That is evidence for the conjecture, not a proof of it.

The construction on a general cover

Section 8 prescribes a symbol for $P(n, k)$ and solves for the weights. Nothing in that used the base graph either. For any cover, writing $\det M(\zeta)$ in $\sigma = \zeta + \zeta^{-1}$ gives a polynomial of degree $D/2$ whose roots in Γ_n are the singular blocks, so the same move applies: prescribe $D/2$ interior grid values, solve for the weights, and the nullity is D .

Theorem 32 (Attainment on a cover). *Let B^n be a cyclic cover and suppose the prescribed-symbol construction yields A carrying the pattern of B^n with $\text{null } A = \nu$. Then*

$$Z(B^n) \geq \nu,$$

by $M \leq Z$. In particular $Z(B^n) \geq D$ whenever the construction attains the ceiling.

This is a statement about all cyclic covers, with $P(n, k)$ — base graph two loops of voltages 1 and k and one edge of voltage 0, $D = 2k + 2$ — as the special case. We have verified attainment of the ceiling on four different base graphs: the $P(n, 2)$ base; the theta graph with voltages $(0, 1, 3)$, a different cubic family; a three-vertex path base with $D = 12$, whose cover has $3n$ vertices; and a four-vertex star base with $D = 12$. In each the constructed matrix has nullity exactly D with every required entry a substantial fraction of the matrix scale.

We also tested the ceiling sharply, which the earlier check did not. Asking whether $\ker A$ injects into a window of D consecutive fibre coordinates is vacuous when the nullity is 2 and D is 12; it bites only at the ceiling. Repeating it on certificates of nullity exactly D , injectivity holds in every case — including the star base, which has *no* Hamiltonian path, so the ordering hypothesis under which the $P(n, k)$ elimination generalises is sufficient but not necessary.

Remark 33 (The hypothesis is needed). Theorem 32 is conditional on the construction succeeding, and that condition is not vacuous. For the $P(n, 2)$ base, $D = 6$, the construction attains 6 at $n = 7, 9, 11, 12$ but fails at $n = 8$ and $n = 10$; and $Z(P(8, 2)) = 5 < 6$ by exhaustive search, so the unconditional statement $Z(B^n) \geq D$ is *false*. What does hold unconditionally is $Z(B^n) \geq \nu$ for the attained ν , which we checked against exhaustively computed zero forcing numbers on covers with 14 to 20 vertices, with no violation. Whether the construction succeeds is not a dimensional question: $n = 7$ attains the ceiling with no slack in the choice of roots while $n = 10$ fails with slack to spare. For the $P(n, 2)$ base the criterion is now known exactly — it is Proposition 19, and $n = 8, 10$ fail because every triple on those grids contains an antipodal pair. For a general base we do not have the analogue.

Remark 34 (Where this sits). The reduction is a discrete Floquet theory. The monodromy of Theorem 11 plays the role of the Hill discriminant for a periodic Jacobi operator, and the condition $T = I$ — all Floquet multipliers roots of unity — is the statement that the operator has purely periodic spectrum at the relevant energy. The parameter $M(G)$ itself belongs to the minimum rank programme for graphs, and $Z(G)$ arose independently in the controllability of quantum spin networks and in sensor placement for power grids.

Certificates beyond the rational ones

Theorem 29 is sufficient but not necessary: F need not have rational coefficients, only real ones, and its roots need only lie in Γ_n , not form full Galois orbits. By (2) every period-one symbol with three or more grid roots arises from some triple of them, so running over all triples of Γ_n finds the true period-one optimum for each (n, k) . Candidates so found are re-checked in exact arithmetic in $\mathbb{Q}(\zeta_n)$, representing $2 \cos(2\pi m/n)$ as $x^m + x^{n-m}$ modulo $\Phi_n(x)$: the membership identity, the reality of a, α, β, c^2 (they are fixed by $x \mapsto x^{n-1}$) and $c^2 \neq 0$ are all decided exactly, and the sign of the nonzero algebraic number c^2 is then read off at high precision. This is not pedantry — see the remark on numerical traps below.

Theorem 35. *Each of the following holds, with an explicit period-one certificate verified exactly in $\mathbb{Q}(\zeta_n)$:*

$$\begin{aligned} Z(P(n, 3)) &= M(P(n, 3)) = 8 && \text{for } n = 40, 60, 80, 120, \\ Z(P(n, 4)) &= M(P(n, 4)) = 10 && \text{for } n = 60, 70, 90, 120, \\ Z(P(n, 5)) &= M(P(n, 5)) = 12 && \text{for } n = 24, 48, 72, 96, 120, \\ Z(P(n, 2)) &= M(P(n, 2)) = 6 && \text{for } 9 \leq n \leq 130, n \neq 10. \end{aligned}$$

The $k = 4$ and $k = 5$ lines are the first exact values of $Z(P(n, k))$ for $k \geq 4$, the range [2] records as open; previously we had only $8 \leq Z(P(n, 4)) \leq 10$ and nothing at all for $k = 5$. The $k = 3$

line is a matrix proof of values [2] obtains combinatorially. Some certificates are irrational: at $(n, k) = (40, 3)$ one has $a = -\sqrt{2}$, $d = 0$, $c^2 = 1$, and at $(60, 3)$, $c^2 = 1/\varphi$ with φ the golden ratio. Full output is in `results/zero_forcing/certified_lower_bounds.txt`.

Remark 36 (Why exact arithmetic, twice over). Counting grid roots by $|F(r)| < \varepsilon$ is not safe. With $\varepsilon = 10^{-8}$ the scan reported nullity 7 for $k = 2$ at $n = 126, 128, 130$ — impossible, since $2k + 2 = 6$ — because Γ_n becomes dense and near-roots pass the test. Separately, a certificate at $(n, k) = (48, 7)$ with $|a|$ enormous against the edge weights made a threshold relative to $\max |A|$ report nullity 48. Both were caught by Corollary 12, which bounds the nullity by $2k + 2$ and so serves as an automatic correctness check on any claimed certificate, exactly as the inequality $Z \geq M$ did for the earlier period-two search. The fixes are to test each Fourier block against its own scale, and to certify in $\mathbb{Q}(\zeta_n)$ rather than in floating point.

Where the method stops

Running the exhaustive triple scan for $2 \leq k \leq 8$ and $n \leq 130$ gives the period-one optimum exactly (`verification/period1_optimum.py`):

k	$2k + 2$	best period-one nullity	attained at
2	6	6	every n with $9 \leq n \leq 130$ except $n = 7, 8, 10$
3	8	8	$n = 40, 60, 80, 120$
4	10	10	$n = 60, 70, 90, 120$
5	12	12	$n = 24, 48, 72, 96, 120$
6	14	8	$n = 35, 37, 70, 74, 105, 111$
7	16	10	$n = 60, 120$
8	18	8	$n = 21, 26, 27, 39, \dots$

Table 3: The period-one construction reaches the ceiling $2k + 2$ for $k \leq 5$ and, in the computed range, falls well short for $k \geq 6$.

Proposition 37 (The rational search is complete). *Fix k . A period-one symbol attaining $2k + 2$ must be squarefree with all $k + 1$ roots interior and on Γ_n ; if in addition it has rational coefficients it is a product of distinct Ψ_d with $\sum \deg \Psi_d = k + 1$. Every factor then satisfies $\deg \Psi_d \leq k + 1$, i.e. $\varphi(d) \leq 2k + 2$, which admits only finitely many d , all listable. Hence for each k the search over rational symbols is a finite complete enumeration. Carrying it out (`verification/cyclotomic_exhaustive.py`):*

k	$2k + 2$	candidates	rejected		admissible	
			outside family	$c^2 = 0$	sym	cs
2	6	17	0	7	5	5 more
3	8	35	21	8	0	6
4	10	67	54	7	3	3 more
5	12	127	118	6	1	2 more
6	14	225	216	9	0	0
7	16	385	378	6	0	1
8	18	651	644	7	0	0
9	20	1065	1057	8	0	0
10	22	1705	1698	7	0	0

The last two columns are the point. The original enumeration reported only the “sym” column and recorded $k = 3$ and $k = 7$ as impossible for rational symbols; they are not. By Theorem 16 only the $c^2 = 0$ column is fatal, and every remaining rejection has $c^2 \in \{-1, -2\}$, hence is realisable in the combinatorially symmetric class. This is what yields Theorem 24 at $k = 3$ and Theorem 27 at $k = 7$. Note also that the $c^2 = 0$ counts are 7, 8, 7, 6, 9, 6, 7, 8, 7, never more than nine, exactly as Corollary 17 predicts from its bound of 25.

Two further things follow. First, $k = 6, 8, 9$ and 10 really are impossible over $\mathbb{Q}$ in both classes, and now for an identified reason rather than a failed search: every candidate that clears the family condition is one of the finitely many degenerate symbols. Second, the rejections are overwhelmingly of one kind — at $k = 7$, 378 of the 385 candidates fail because F is *not in the three-parameter family*. The binding constraint is therefore the $k - 2$ linear membership conditions imposed on a 3-dimensional family, and it tightens as k grows. That identifies the fix: enlarge the family.

How far must one enlarge it? A period- d matrix has $5d$ weights, but most of that freedom is gauge: conjugating by a positive diagonal matrix preserves both the pattern and the nullity. What matters is the dimension of the set of symbols actually reachable. Writing G for the degree- $(k + 1)$ polynomial in $\sigma = \zeta + \zeta^{-1}$ obtained from $\det M_\ell$, we computed the rank of the Jacobian of the map (weights) $\mapsto$ (monic coefficients of G) at random points, for $3 \leq k \leq 13$ and $1 \leq d \leq 7$ (`verification/period_d_rank.py`). In every case

$$\text{rank} = \min(2d + 1, k + 1).$$

So the reachable symbols form a $(2d + 1)$ -dimensional family — not the $5d - 1$ -dimensional one a naive parameter count suggests — and the map is a submersion, so that every G is locally reachable, precisely when

$$2d + 1 \geq k + 1, \quad \text{that is} \quad d \geq \lceil k/2 \rceil.$$

At $d = 1$ this returns the 3 of Proposition 14, which is a check on the computation. We record the law as an observation supported by computation, not as a theorem: we have not proved the rank formula, only verified it over that range. Its consequence, if it holds in general, is that the period- d construction should reach $2k + 2$ for every k once $d \geq \lceil k/2 \rceil$ and n/d is large enough to host $k + 1$ grid roots — a construction whose modulus is d rather than an lcm of cyclotomic orders, and therefore one that would cover a fixed congruence class of n for every k .

So the construction settles $k \leq 5$ on those residue classes and says nothing useful for $k \geq 6$; whether that is a limitation of period one or of $n \leq 130$ we do not know. For n outside the classes above, no certificate of nullity $2k + 2$ was found at all. This is not for want of parameters. Dropping equivariance

entirely leaves $5n$ free weights in the symmetric class and $8n$ in the combinatorially symmetric one, whereas a nullity of $r = 2k + 2$ costs $r(r + 1)/2 = (k + 1)(2k + 3)$ conditions in the first case and r^2 in the second — budgets that are met comfortably at every graph we tried. A two-stage search (many random starts on $\sum_{j \leq r} \lambda_j^2$ with analytic gradients, then a Gauss–Newton polish of the bilinear system $A(w)X = 0$, $X^\top X = I$; `verification/fast_nullity.py`) attains nullity 6 at machine precision on $P(12, 2)$ and $P(14, 2)$ but plateaus near 10^{-4} for $r = 7, 8$ on $P(12, 3)$ over hundreds of independent starts, in both matrix classes. We therefore record, as an observation rather than a theorem, that the obstruction to $2k + 2$ at general n appears to be arithmetic — whether Γ_n contains a compatible configuration of $k + 1$ points — and not a matter of dimension. Two guardrails are worth stating, since we tripped on both: the combinatorially symmetric class *contains* the symmetric one, so any search reporting a smaller answer for it is faulty (ours did, on $P(12, 2)$, until the larger search was seeded with the symmetric certificate); and any claimed nullity above $2k + 2$ is a bug by Corollary 12.

Why every result so far lives on a congruence class

Each exact value in this section holds for n in a divisibility class, not for all large n . That is not an artefact of our search. It is forced, and the following says exactly how.

Theorem 38 (Period-one certificates are bound to congruence classes). *Fix $k \geq 1$ and a period-one matrix A with parameters (a, α, β) and symbol F .*

- (i) *Suppose $\text{null } A = 2k + 2$ for some $n > 2k$. Then every root of F is $2\cos(2\pi j/d)$ with $\gcd(j, d) = 1$ and $d \geq 3$, and such a root lies in Γ_m if and only if $d \mid m$. Hence the set of n at which this same A attains the ceiling is exactly $\{n : \text{lcm}(D) \mid n\}$, where D is the multiset of those d . It is therefore a single divisibility class of density $1/\text{lcm}(D) \leq 1/3$, and it contains at most one prime.*
- (ii) *Consequently, to attain the ceiling for infinitely many n outside one divisibility class — in particular for infinitely many primes, or for all large n — the weights must depend on n .*
- (iii) *For $k \geq 3$ that dependence is constrained by one universal identity: the coefficient of s^{k-1} in F equals $-k$ for every (a, α, β) , so the roots always satisfy*

$$e_2(r_1, \dots, r_{k+1}) = -k, \quad \text{equivalently} \quad \sum_i r_i^2 = e_1^2 + 2k.$$

For $k = 2$ the coefficient of $s^{k-1} = s$ is $\alpha - 2$, which is free, so no such condition exists.

Proof. (i) Attaining $2k + 2$ needs all $k + 1$ roots interior and in Γ_n (Proposition 14), and an element of $\Gamma_n \cap (-2, 2)$ is $2\cos(2\pi m/n)$, which written in lowest terms is $2\cos(2\pi j/d)$ with $d \mid n$, $d \geq 3$. Conversely $2\cos(2\pi j/d) \in \Gamma_m$ iff $d \mid m$. The roots of F do not move with n , since F depends only on (a, α, β) ; so all of them lie on Γ_n exactly when every d divides n . A set $\{n : L \mid n\}$ with $L \geq 3$ contains a prime only if L is itself prime, and then only L .

(ii) Immediate from (i): a fixed A covers one class, and a union of finitely many classes still omits all but finitely many primes.

(iii) Write $(s + a)L_k = sL_k + aL_k$. Since L_k has nonzero coefficients only in degrees congruent to k modulo 2, aL_k has no s^{k-1} term, and the s^{k-1} coefficient of sL_k is the s^{k-2} coefficient of L_k , namely $-k$. The terms αs and β reach degree $k - 1$ only when $k \leq 2$. For a monic polynomial of degree $k + 1$ the coefficient of s^{k-1} is e_2 , and $\sum r_i^2 = e_1^2 - 2e_2$. $\square$

Part (iii) is what separates $k = 2$ from everything else. With no residual condition, Theorem 20 is free to choose the roots *as functions of n* — the three most negative elements of Γ_n° — and that is precisely why $k = 2$ is general in n while no other k is. For $k \geq 3$ an n -dependent choice must satisfy $e_2 = -k$ identically along the family. Substituting $r_i = 2\cos(2\pi m_i/n)$ and expanding e_1^2 by the product-to-sum formula turns this into a rational linear relation among cosines of rational angles set equal to a rational number. Relations of that shape are the subject of Conway–Jones [3], and they are arithmetically rigid: they hold not on open families but on sparse, classifiable configurations. That rigidity is the mechanism behind every divisibility condition in this paper, and it is the same mechanism as in Theorem 30.

We verified the identity of (iii) against all 26 certificates found in this paper: $e_2 = -k$ in every case with $k \geq 3$, and e_2 takes six different values among the $k = 2$ certificates, as (iii) predicts (`verification/general_n.py`).

Remark 39 (What this licenses, and what it forecloses). Theorem 38 should be read as a demarcation rather than a limitation. It says the entire cyclotomic layer of this paper — every result in Sections 8 and 10 — is *provably* incapable of producing a statement valid for all large n , so the congruence conditions are a boundary of the method and not a shortfall in its execution. It also says where to go instead: Corollary 12 characterises the ceiling for *every* matrix on the pattern, not only the invariant ones, by the condition $T = I$ on the monodromy, and that condition carries no arithmetic in it.

10 Period- d certificates, and the first values for $k \geq 6$

Proposition 37 leaves $k = 6$ and $k \geq 8$ untouched: no rational period-one symbol attains $2k + 2$ there in either matrix class, and the exhaustive triple scan of Section 8 reaches only 8, 10, 8 for $k = 6, 7, 8$ over $n \leq 130$. ($k = 7$ is settled exactly by Theorem 27; what follows gives it, and $k = 6$ and $k = 8$, for further n .) The rank law says why, and what to do about it.

Prescribing the symbol

For $d \geq \lceil k/2 \rceil$ the map (weights) $\mapsto$ (monic coefficients of G) is a submersion, so instead of *searching* for a certificate one can *prescribe* one: choose $k + 1$ interior values of Γ_m , $m = n/d$, form the monic G^* with exactly those roots, and solve $G(w) = G^*$ for the weights. This is $k + 1$ equations in $5d$ unknowns with a full-rank Jacobian, so the solution set is a manifold of dimension $5d - (k + 1)$ — eight-dimensional for $k = 6$, $d = 3$. Each root is interior, hence met by two Fourier indices ℓ and $m - \ell$, so all $2k + 2$ of the corresponding blocks are singular and the nullity is $2k + 2$ by Corollary 12.

Two conditions must both hold, and they are independent: $m \geq 2k + 4$, so that Γ_m contains $k + 1$ distinct *interior* values at all; and $d \geq \lceil k/2 \rceil$, so that the symbol is reachable. For $k = 6$ this forces $d \geq 3$ and $n \geq 48$.

A degeneracy, and how it is escaped

Solving $G(w) = G^*$ without further constraint reliably returns a point at which one required-nonzero weight — in every case a spoke — has collapsed to zero. Such a point is a genuine solution of $G(w) = G^*$, because with a spoke weight zero the matrix block-diagonalises and G factors as $\det U_\ell \cdot \det V_\ell$; but the matrix then carries the pattern of a *different* graph and certifies nothing. At these points the Jacobian rank drops from 7 to 6: they are singular points of the map, which is why repeated random starts return to them.

Neither a realisability barrier nor continuation escapes them. Driving the vanishing spoke away from zero while re-solving loses the symbol at the first step ($\tau = 0.02$, error 9×10^{-5}). What does escape is *slack in the choice of roots*. At $n = 48$, Γ_{16} has exactly 7 interior values against the 7 required, so the root set is unique and there is nothing to vary; at $n = 54$ there are 8 interior values, giving 8 root sets, and all eight degenerate. At $n = 60$ and $n = 72$ there are 9 and 11 interior values, hence 36 and 330 root sets, and realisable points exist.

Theorem 40.

$$\begin{aligned} Z(P(60, 6)) &= M(P(60, 6)) = 14, & Z(P(72, 6)) &= M(P(72, 6)) = 14, \\ Z(P(96, 7)) &= M(P(96, 7)) = 16, & Z(P(96, 8)) &= M(P(96, 8)) = 18. \end{aligned}$$

In each case the period is the one the rank law prescribes: $d = 3$ for $k = 6$ and $d = 4$ for $k = 7, 8$, i.e. $d = \lceil k/2 \rceil$.

These are the first exact values of $Z(P(n, k))$ for $k \geq 6$. That the required period was predicted in advance, and that $d = 3$ then failed for $k = 7, 8$ while $d = 4$ succeeded, is the strongest evidence we have that the rank law is describing something real rather than fitting the cases it was measured on. The certificates are period-three matrices found by the construction above and verified on the full $2n \times 2n$ matrix, independently of the Fourier blocks used to find them. For $n = 60$ with roots at $\ell \in \{2, 3, 4, 5, 6, 8, 9\}$ the 120×120 matrix has its 14 smallest eigenvalues below 4×10^{-15} and its fifteenth at 5.1×10^{-2} — a separation of thirteen orders of magnitude — every vertex carries exactly three off-diagonal nonzeros, and the smallest required entry is 0.188 of the matrix scale. For $n = 72$ with $\ell \in \{1, 4, 6, 7, 9, 10, 11\}$ the corresponding figures are nullity 14, first nonzero eigenvalue 1.2×10^{-3} of scale, and smallest required entry 0.162. For $k = 7$ at $n = 96$, $d = 4$, $\ell \in \{3, 5, 6, 7, 8, 9, 10, 11\}$: nullity 16, sixteen eigenvalues below 1.4×10^{-14} , smallest required entry 0.107. For $k = 8$ at $n = 96$, $d = 4$, $\ell \in \{2, 3, 4, 5, 6, 7, 8, 9, 11\}$: nullity 18, first nonzero eigenvalue 1.25×10^{-2} of scale, smallest required entry 0.260. Several further root sets at each n give independent certificates — three at $k = 8$ alone, the best with smallest required entry 0.362 — so these are not isolated points but regions.

k	$2k + 2$	period d	$\lceil k/2 \rceil$	n	smallest required entry
6	14	3	3	60, 72	0.188, 0.162
7	16	4	4	96	0.107
8	18	4	4	96	0.260

Table 4: Period- d certificates attaining the ceiling for $k \geq 6$. The period used is in every case the one the rank law prescribes.

Why an exact certificate is hard here

The $k \geq 6$ certificates are numerical because their weights are algebraic numbers of no recognised form. The obvious remedy — find *rational* weights, so that the nullity can be confirmed by exact rational elimination as for $k = 4, 5$ — runs into a structural obstruction.

Proposition 41. *Let A be a period- d matrix on $P(n, k)$ with rational weights, and $m = n/d$. The Fourier blocks have entries in $\mathbb{Q}(\zeta_m)$, and the Galois automorphism $\sigma_j : \zeta_m \mapsto \zeta_m^j$ carries M_ℓ to $M_{j\ell}$. Since σ_j fixes the weights, $\sigma_j(\det M_\ell) = \det M_{j\ell}$, so*

$$\det M_\ell = 0 \implies \det M_{j\ell} = 0 \quad \text{for every } j \text{ coprime to } m.$$

*Hence the set of singular blocks is a union of full Galois orbits of Γ_m , and a rational certificate can only realise a **Galois-closed** root set.*

This is checked numerically in the form it can be: for rational weights the product of $\det M_\ell$ over a Galois orbit — the norm — comes out exactly rational, to machine zero, at $(k, d, n) = (6, 3, 60)$, $(6, 3, 72)$, $(6, 4, 64)$, $(5, 1, 24)$ and $(7, 4, 96)$.

The obstruction is now visible: exactness requires a Galois-closed root set, and Galois-closed root sets are precisely the ones that degenerate. All eight at $(k, d, n) = (6, 3, 72)$ hit the symbol to 10^{-15} with nullity 14 and all eight collapse a weight; both at $(6, 3, 60)$ likewise. So for these parameters *no rational certificate exists*.

Relaxing to a subfield does not obviously help, and it is worth being precise about why. If the weights lie in $K = \mathbb{Q}(\zeta_m)^H$ then the singular set need only be H -stable, and at $(6, 3, 72)$ all three

index-two subgroups admit realisable root sets. But that implication runs one way: an H -stable root set being realisable shows only that the necessary condition is not violated, and gives no reason to expect a certificate for it to have weights in K . Indeed integer-relation detection finds no such representation for the certificates we have, a search parametrised inside $\mathbb{Q}(\sqrt{2})$ from the start returns nothing, and at $(6, 3, 60)$ every index-two set degenerates outright. We therefore record exactness for $k \geq 6$ as **open**, with the evidence leaning negative, rather than as a route in progress. The remaining candidate is a computer-assisted existence proof (interval Newton) for an irrational certificate, which the ceiling makes tractable in principle, since certifying the equation $G(w) = G^*$ exactly delivers the nullity as a discrete consequence.

How many n does this reach?

The construction needs a divisor d of n with

$$\lceil k/2 \rceil \leq d \leq \frac{n}{2k+4},$$

the left inequality from the rank law and the right from requiring $\Gamma_{n/d}$ to contain $k+1$ distinct interior values. This is a real restriction and it cuts both ways.

Proposition 42 (Reachability). *The set of n admitting such a divisor has density 1, but its complement contains every prime, and also every $n = 2p$ once $k \geq 6$ (since then $d \geq 3$ while the only small divisor is 2). Convergence is slow: at $k = 6$ the reachable proportion of n is 54% up to 240, 71% up to 1000 and 79% up to 5000; at $k = 8$, 36%, 62% and 73%.*

In particular **prime n lies outside the period- d method entirely**, because d must divide n and $d = 1$ gives rank 3. That is not a gap in the search but a structural feature: for prime n the only available construction is the fully non-equivariant one, i.e. solving $T = I$ directly in the $5n$ weights, which is the regime where our searches plateau. Since $Z(P(23, 5)) = 12 = 2k + 2$ with 23 prime, such n certainly can attain the ceiling; we simply cannot certify them this way.

Larger d is not merely permitted but helpful. The rank law requires only $d \geq \lceil k/2 \rceil$, and taking d strictly larger enlarges the solution manifold from dimension $5d - (k+1)$, giving more room to avoid the degenerate locus. The effect is visible at the tightest configurations. At $k = 6$, $d = 3$, $n = 48$ the grid Γ_{16} offers exactly 7 interior values for the 7 required — a unique root set, no slack — and every solution collapses a spoke. At $k = 6$, $d = 4$, $n = 64$ the grid is again Γ_{16} with the same unique root set, but now the manifold has dimension $20 - 7 = 13$ rather than $15 - 7 = 8$, and a realisable certificate is found on the first attempt (smallest required entry 0.086, gap 6.8×10^{-2}). So the degeneracy of the previous subsection is a symptom of working at the minimum period, not an obstruction inherent to tight root sets. Certificates were likewise found at $k = 6$, $d = 6$, $n = 96$.

Within the reachable set, the construction appears to work essentially always. Sweeping n for $k = 6$ at $d = 3$ produced a realisable certificate at *every* n tested:

$$n = 60, 72, 84, 90, 96, 108, 120,$$

most of them on the first root set tried, with smallest required entries between 0.12 and 0.46. So

$$Z(P(n, 6)) = M(P(n, 6)) = 14 \quad \text{for each of these } n,$$

and the earlier appearance of a scattered list of congruence classes was an artefact of the cyclotomic route, not of the problem. We state this as a computation rather than a theorem: we have no proof that a realisable root set exists for every reachable n , only that one was found in every case examined.

Remark 43 (Status of Theorem 40). Unlike the cyclotomic certificates of Section 8, which are integer matrices whose nullity is confirmed by exact rational elimination, these are found by a Newton solve and their weights are algebraic numbers of no recognised form. We therefore record Theorem 40 as **numerically certified but not yet verified in exact arithmetic**. The evidence is strong — a thirteen-order spectral separation, the correct zero pattern with every required entry a fifth of the matrix scale, and the nullity never exceeding the ceiling $2k + 2$, which is an automatic check on the whole pipeline — but it is not of the same kind as the $k = 4$ and $k = 5$ results, and we do not claim that it is. The natural route to closing the gap is to choose the root set to be a union of full Galois orbits of Γ_m , which makes G^* a product of the Ψ_e and hence puts the system $G(w) = G^*$ over $\mathbb{Q}$; that is in progress.

11 Beyond congruence classes: a certified value at a prime

Theorem 38 says the cyclotomic layer cannot produce a value valid for all large n , and points at the replacement: Corollary 12 characterises the ceiling for *every* matrix on the pattern by the condition $T = I$, which carries no arithmetic. This section carries that out at a prime n , where by Theorem 38(i) no period-one certificate can reach.

A better system to certify

Optimising $T = I$ is convenient but bad to certify: the map $w \mapsto T$ has a 51-dimensional image inside the 64-dimensional space of 8×8 matrices at $k = 3$, so $\{T = I\}$ is cut out by 13 local equations we cannot identify and no square subsystem is available. We certify the nullity directly instead.

Now $\text{null } A \geq r$ if and only if some full-rank $K \in \mathbb{R}^{2n \times r}$ has $AK = 0$. Fixing a row permutation P and writing K in graph form $K = P \begin{pmatrix} I_r \\ X \end{pmatrix}$ makes $\text{rank } K = r$ automatic, and

$$f(w, X) := A(w) K(X) = 0$$

is *bilinear*: quadratic of total degree 2, with rational coefficients. Its Jacobian is affine, so the Lipschitz constant of J is an explicit integer, and a Krawczyk contraction bound needs only three norms rather than an interval matrix product.

The choice of which equations to solve is not free. The system $AK = 0$ has $2nr$ equations of rank $2nr - \binom{r}{2}$, and the deficiency is structural: $K^\top AK$ is symmetric because A is, and $\binom{r}{2}$ is exactly its antisymmetric part. Concretely, if the $(2n - r)r$ equations on the non-pivot rows $P[r:]$ vanish then $K^\top AK$ equals the $r \times r$ block $(AK)[P[:r]]$, which must therefore be symmetric, so setting its upper triangle to zero forces the lower triangle as well. Hence

$$\underbrace{(AK)[P[r+m], t] = 0}_{(2n-r)r \text{ equations}} \quad \text{and} \quad \underbrace{(AK)[P[s], t] = 0, \ s \leq t}_{\binom{r+1}{2} \text{ equations}}$$

is *exactly equivalent* to $AK = 0$. Selecting rows by pivoting instead would leave $\binom{r}{2}$ equations unproved.

Theorem 44 (Certified values off every congruence class). *For every n with $17 \leq n \leq 33$,*

$$Z(P(n, 3)) = M(P(n, 3)) = 8.$$

In particular this holds at the primes $n = 17, 19, 23, 29, 31$, and the range is consecutive: seventeen values with no exceptions.

Proof. $Z(P(17, 3)) \leq 8$ is the rotation bootstrap. For the lower bound, take $n = 17$, $k = 3$, $r = 8$, so $N = 34$ and the system above has $85 + 208 = 293$ unknowns and 272 equations, of which $208 + 36 = 244$ are the structural ones. At a numerically located point the Jacobian of the 244 structural equations has rank exactly 244, so fixing the remaining 49 unknowns at explicit rationals with denominator 10^4 leaves a square regular system. Newton in 60-digit arithmetic drives its residual to 3.2×10^{-61} ; the 28 dependent equations, which are not solved for, fall to 4.1×10^{-61} , confirming the structural reduction. With Y an approximate inverse of the square Jacobian,

$$\|Y\|_\infty = 1.54 \times 10^3, \quad \|I - YJ\|_\infty = 2.18 \times 10^{-12}, \quad \|Yf\|_\infty = 2.85 \times 10^{-60},$$

and with $L = 12$ bounding the Lipschitz constant of J in the ∞ -norm, the Krawczyk test on the box of radius $\delta = 10^{-12}$ gives $\alpha = \|I - YJ\|_\infty + \|Y\|_\infty L \delta = 1.85 \times 10^{-8} < 1$ and $\|Yf\|_\infty + \alpha \delta \leq \delta$. So a true zero (w^*, X^*) exists in that box. Then $A(w^*)K(X^*) = 0$ with rank $K = 8$, so $\text{null } A(w^*) \geq 8$, and ≤ 8 by Corollary 12. Every edge weight of $A(w^*)$ exceeds 0.05 times the matrix scale and the ninth singular value exceeds 0.126 times it, so the box is far from any degenerate stratum and the pattern is exactly that of $P(17, 3)$.

The remaining n are identical in structure and each passes the same test; for $n = 19$, to give one more set of numbers, there are 335 unknowns and 304 equations of which 276 are structural and 28 dependent, 59 unknowns are fixed at rationals, and $\|Y\|_\infty = 4.37 \times 10^2$, $\|I - YJ\|_\infty = 1.06 \times 10^{-12}$, $\|Yf\|_\infty = 1.20 \times 10^{-60}$, with Krawczyk passing at $\delta = 10^{-13}$ and $\alpha = 5.25 \times 10^{-10}$. See `verification/krawczyk.py` and `verification/certify_pipeline.py`. $\square$

Together with Theorem 24 this gives $Z(P(n, 3)) = 8$ at $n = 10$, every $17 \leq n \leq 33$, and every multiple of 10 — a set that no single divisibility condition describes. The search fails only at $n = 13$ and $n = 15$, where the solution variety has dimension 3 and 9 by (3), against 15 at $n = 17$ where it first succeeds. By Lemma 46 the slice searched is complete, so those are failures of the search and not evidence that no matrix of nullity 8 exists there. That distinction is not academic: $n = 18$ resisted the restart budget we used for its neighbours and certified once the budget was raised, so a failure here carries no information beyond the effort spent.

Remark 45 (Three independent numerical routes, one certificate). The certificate at $n = 17$ was located three separate ways — by minimising $\|T - I\|$ over all $5n$ weights, by an alternating solve of $AK = 0$ with the spokes normalised, and by the fully gauge-fixed search of (3) — and all three lead to a Krawczyk-certified matrix of nullity 8, with the gauge-fixed route giving the best conditioning ($\|Y\|_\infty = 2.1 \times 10^2$, ninth singular value 0.167 of the matrix scale, minimum edge weight 0.146). Since stage 1 is only a search and stages 2–3 are a proof, the numerical route taken does not enter the result; that three unrelated ones agree is a check on the search, not on the certificate.

Since 17 and 19 are prime, Theorem 38(i) shows no period-one certificate — in particular nothing in Sections 8 or 10 — can give these values. Indeed no divisibility class containing either of them is available at all: a class $\{n : L \mid n\}$ with $L \geq 3$ contains a prime only when L is that prime, and Γ_{17} offers only Ψ_1 and Ψ_{17} , the latter of degree $8 \neq k + 1 = 4$. These are the first values in this paper that do not sit on a divisibility class.

The reachable set stabilises

That $n = 17$ works and $n = 13$ does not is explained by a measurement rather than by arithmetic. The dimension of the reachable set of monodromies, that is the rank of $\partial T / \partial w$, grows with n and then *saturates*:

$k = 2, n$	5	6	7	8	9	10	11	12	≥ 13
rank	19	23	26	28	30	31	31	31	31
$k = 3, n$	7	9	11	13	14	15	16	17	≥ 18
rank	27	34	40	46	49	51	51	51	51

Past $n = 11$ for $k = 2$ and $n = 15$ for $k = 3$, increasing n adds no new reachable monodromies. So whether I is reachable stops depending on n in any essential way once n passes roughly $4k + 3$, and the observed pattern of successes and failures follows it: for $k = 3$ the ceiling is attained at $n = 17, 20$ and not at $n = 13 < 15$. Independently, at $n = 17, k = 3$ the solution variety of $f = 0$ has dimension $293 - 244 = 49 = 5n - \binom{2k+3}{2}$, exactly the codimension count for symmetric matrices on the pattern with nullity $2k + 2$.

There is also a hard *necessary* condition, from spending the gauge, and it pays to be careful about what the gauge can and cannot do.

Lemma 46 (A complete slice). *Every symmetric matrix on the pattern of $P(n, k)$ with all weights nonzero is carried, by a conjugation preserving symmetry, the pattern and the nullity, to one with every spoke equal to $+1$ and every outer weight equal to ± 1 — for even n , with a single outer weight left arbitrary. The signs may be taken to be $+1$ except possibly one, so for odd n there are two such slices and for even n one.*

Proof. Two conjugations are available. Positive diagonal $A \mapsto DAD$ scales $b_i \mapsto d_i d_{i+1} b_i$, $c_i \mapsto d_i d_{n+i} c_i$, and sign diagonal $A \mapsto SAS$ with $S = \text{diag}(\pm 1)$ sends $b_i \mapsto s_i s_{i+1} b_i$, $c_i \mapsto s_i s_{n+i} c_i$. Since $d_i d_{n+i} > 0$, each spoke's *sign* is invariant under the first and must be fixed by the second: taking $s_{n+i} = s_i \text{sign}(c_i)$ makes every spoke positive, and choosing the inner d_{n+i} makes every spoke 1. The outer d_i then satisfy $u_i + u_{i+1} = -\log |b_i|$ with $u_i = \log d_i$, a cyclic system whose matrix is invertible for odd n and has corank 1 for even n ; in the latter case drop one equation, leaving one outer weight unnormalised. Finally $s_i s_{i+1} \text{sign}(b_i) = 1$ is solvable exactly when $\prod_i \text{sign}(b_i) = 1$, giving the two sign classes for odd n , while for even n the unnormalised weight absorbs the invariant. $\square$

We verified Lemma 46 numerically before relying on it: random matrices on the pattern are carried into the slice with the spokes exactly $+1$ and $|b_i| = 1$ to 10^{-15} , and the nullity unchanged, for both parities of n (`verification/gauge_fixed.py`). The point of the lemma is that a search restricted to the slice is not restricted at all, so a failure inside it is a failure of the search and not evidence that no matrix exists.

What remains after the fixing is $2n$ diagonal entries and n inner weights, with no residual gauge, and the solution variety of $f = 0$ has dimension

$$[3n + (2n - r)r] - [(2n - r)r + \binom{r+1}{2}] = 3n - (k + 1)(2k + 3),$$

which is nonnegative only when

$$n \geq \frac{(k + 1)(2k + 3)}{3} = 7, 12, 19, 26, 35, 46, 57 \quad (k = 2, \dots, 8). \quad (3)$$

Conjecture 47. *For every $k \geq 2$ there is $N(k)$ with $Z(P(n, k)) = M(P(n, k)) = 2k + 2$ for all $n \geq N(k)$.*

For $k = 2$ this is Theorem 20, with $N(2) = 11$, proved. We do not conjecture that $N(k)$ is linear: (3) forces $N(k) \geq \lceil (k + 1)(2k + 3)/3 \rceil$, which is quadratic, and at $k = 6$ that bound is 35 — larger than the value $4k + 3 = 27$ that the saturation data alone might suggest. The two quantities measure different things, and the dimension count is the binding one. Our data is consistent with $N(k)$

exceeding (3) by a small additive amount: $N(2) = 11$ against a bound of 7, and $14 \leq N(3) \leq 17$ against a bound of 12, since the ceiling is attained at $n = 17$ and 20 and not at $n = 13$.

The content of the conjecture for $k \geq 3$ is a reachability statement about products of transfer matrices, with no arithmetic in it — which is why we expect it to behave uniformly in n , unlike everything in Sections 8 and 10.

Remark 48 (A strict gap, and a correction to our own evidence). Two caveats. First, $n = 10, k = 2$ resists the search even though its reachable dimension is already saturated, and $Z(P(10, 2)) = 6$; so we expect $M(P(10, 2)) < Z(P(10, 2))$, a strict gap, though we have not proved it. Second, the eigenvalue-minimising search we reported earlier is unreliable and its negative results should not be read as evidence: it gave maximum nullity 7 at $P(20, 3)$, whereas the ceiling 8 is attained there both by Theorem 24 and, in the *symmetric* class, by the method of this section. It was collapsing onto the degenerate strata — all spokes zero, or all edges zero — that the formulations here exclude by construction.

A route to all large n

Conjecture 47 asks for one construction covering every large n . Certifying n one at a time cannot do that. But the condition $T = I$ is multiplicative along the cycle, and that suggests assembling the weight sequence from pieces each of which contributes the identity. The obstacle is that T_i is built from the weights in a window of width about $2k + 2$ around position i , so pieces interact across their boundaries and the product does not factorise. The following removes the obstacle.

Definition 49. Fix a *docking pattern*: values of (b, c, e, a, d) at $k + 1$ consecutive positions. A *tile* of length $\ell \geq 2(k + 1)$ is a weight sequence of length ℓ whose first $k + 1$ and last $k + 1$ positions carry the docking pattern. Its *interior* is the remaining $\ell - 2(k + 1)$ positions.

Lemma 50 (Tile independence). *Let the weight sequence on $\mathbb{Z}_n$ be a concatenation of tiles sharing one docking pattern. Then for each tile the product of the transfer matrices over its positions depends only on that tile's interior weights, and the monodromy is the product of these tile products.*

Proof. By Theorem 11, T_i is determined by the coefficients of the i th reduced equation, which involve b at indices in $[i - k - 1, i + k]$, c and a at $i - k, i, i + k$, e at $i - k$ and i , and d at i . So T_i depends only on weights at indices in $[i - k - 1, i + k]$. If the tile occupies $[s, s + \ell)$ and $i \in [s, s + \ell)$, this window is contained in $[s - k - 1, s + \ell + k - 1]$. The part below s lies in the preceding tile's final $k + 1$ positions and the part at or above $s + \ell$ lies in the following tile's first $k + 1$ positions; both carry the docking pattern, as do $[s, s + k]$ and $[s + \ell - k - 1, s + \ell)$. Hence every weight T_i sees is either docking or in this tile's interior. The second statement is associativity of the product $T = T_{n-1} \cdots T_0$ regrouped along the tiles. $\square$

Theorem 51 (Tiling). *Fix k and $L \geq 2(k + 1)$. Suppose that for every ℓ with $L \leq \ell < 2L$ there are interior weights, all of the required entries nonzero, making the tile product of length ℓ equal to the identity. Then*

$$Z(P(n, k)) = M(P(n, k)) = 2k + 2 \quad \text{for every } n \geq L,$$

with no congruence condition.

Proof. Every $n \geq L$ is a sum of integers from $[L, 2L)$: if $n < 2L$ take the single term n , and if $n \geq 2L$ take a term L and recurse on $n - L \geq L$. Concatenating the corresponding tiles gives a weight sequence on $\mathbb{Z}_n$ whose monodromy is, by Lemma 50, the product of the tile products, hence

I. All required weights are nonzero, so the matrix carries the pattern of $P(n, k)$ and Theorem 11 applies; by Corollary 12, $\text{null } A = 2k + 2$. With $M \leq Z$ and the rotation bootstrap $Z \leq 2k + 2$, both are equal to $2k + 2$. $\square$

Lemma 50 was checked before use: perturbing one tile's interior leaves the other tiles' products unchanged to the last bit, 0 exactly, for $k = 2, 3, 4$ (`verification/tiles.py`). Two remarks on the shape of Theorem 51. First, the hypothesis over a full interval $[L, 2L)$ is what makes the conclusion gapless; using two tile lengths a, b with $\gcd(a, b) = 1$ instead would give only n of the form $ia + jb$, leaving every $n < (a - 1)(b - 1)$ uncovered, which for the smallest solvable lengths at $k = 3$ is already several hundred. Second, the hypothesis is a *finite* list of independent finite problems: a tile of length ℓ has $5(\ell - 2(k + 1))$ free weights and the requirement is $(2k + 2)^2$ equations, so it is solvable on parameter count once ℓ is a little above $2(k + 1) + (2k + 2)^2/5$. For $k = 3$ that is $\ell \geq 21$, and $L = 21$ gives every $n \geq 21$.

The tiles exist, and are certified

We carried the computation out for $k = 3$. The parameter count is sharp in practice: at $\ell = 19, 20$ the systems are underdetermined (55 and 60 unknowns against 64 equations) and the residual stalls near 4×10^{-1} ; and from $\ell = 21$, the first length with more unknowns than equations, identity tiles are found. A Newton solve produced one for *every* length in the required range, $21 \leq \ell \leq 41$, which is exactly $[L, 2L)$ for $L = 21$; we in fact have tiles for every $21 \leq \ell \leq 43$.

Each of those 23 tiles is then *certified*, which needs one observation to avoid interval arithmetic on a rational map. The condition $T_{\text{tile}} = I$ is a rational equation in the tile weights, so a Krawczyk bound on it directly would require interval matrix products. It is unnecessary: by Corollary 12 the monodromy of the cyclic matrix built from a *single* tile on $P(\ell, k)$ is the tile product, so

$$T_{\text{tile}} = I \iff \text{null } A_{\text{one tile}} = 2k + 2,$$

and the right-hand side is certified by the bilinear machinery of Theorem 44, whose Jacobian is affine. The only extra requirement is that the tiles still compose, which needs them all to carry the same docking pattern; this is arranged by pinning the $5(2k + 2)$ docking weights at their exact integer values ($b = c = e = 1$, $a = d = 0$) rather than letting the solver move them, so they belong to the fixed set and the certified solution carries them exactly.

All 23 certifications at $k = 3$ pass. For $\ell = 22$: 398 unknowns, 352 equations of which 324 are structural and 28 dependent, 40 docking pins and 74 unknowns fixed at rationals in all; high-precision Newton drives the structural residual to 3.1×10^{-61} and the dependent equations to 2.7×10^{-61} ; $\|Y\|_{\infty} = 5.4 \times 10^2$ and the Krawczyk test passes with $\alpha = 6.7 \times 10^{-11}$. Across all lengths $\alpha \leq 2.2 \times 10^{-10}$, every certified matrix has nullity exactly 8, every docking weight agrees with its prescribed integer to 10^{-14} , and the minimum edge weight is 0.15 of the matrix scale. See `verification/certify_tiles.py`.

Remark 52 (The tile equations are degenerate, and by how much). The system $T_{\text{tile}} = I$ has $(2k + 2)^2$ equations but they are not independent. Measuring the rank of the Jacobian of $w \mapsto T_{\text{tile}}$ at a solution gives exactly

$$\text{rank} = \binom{2k + 3}{2}, \quad \text{deficiency} = \binom{2k + 2}{2},$$

namely 36 and 28 at $k = 3$ and 55 and 45 at $k = 4$, at every tile length tested. Two consequences. First, it explains the behaviour of the solve: the residual has $\binom{2k+2}{2}$ components that vanish

automatically near a solution, so a least-squares step is working on a degenerate system, and that is consistent with the erratic convergence reported in the remark below. Second, the count that a tile must satisfy is $5(\ell - 2(k + 1)) \geq \binom{2k+3}{2}$ rather than $\geq (2k + 2)^2$, which for $k = 3$ would permit $\ell \geq 16$ rather than $\ell \geq 22$ and so a smaller L . Whether short tiles exist is a separate question from whether the count allows them, and the answer is in between. With a larger restart budget $\ell = 21$ *does* yield an identity tile, certified like the rest, so $L = 21$ rather than 22 at $k = 3$ and the required range is $[21, 42]$; this does not change Theorem 53, which already reaches $n \geq 17$ through Theorem 44. Below that the searches fail and the failures look different in kind: $\ell = 19$ and 20 stall at 1.4×10^{-1} and 3.7×10^{-1} , and $\ell = 16, 17, 18$ at 8.3×10^{-1} , 7.8×10^{-1} and 8.2×10^{-1} , whereas every length that we eventually solved first appeared as a residual near 10^{-2} or below. So the count is not the binding constraint and I appears not to be a product of fewer than about twenty transfer matrices. We record that as a computation and not a proof: we have no lower bound on the number of factors needed to write I , and $\ell = 21$ is a reminder that a stalled search at these sizes can mean only that the budget was too small.

We note that $\binom{2k+3}{2} = \dim Sp(2k+2)$, and that the space of skew forms B with $T^\top BT = B$ has dimension $k+1$. It is tempting to conclude that T is symplectic with respect to a weight-dependent form. We do not claim it: our attempts to exhibit a nondegenerate invariant form were numerically unstable, giving residuals of 10^{-12} at $k = 2, 4, 5$ but 10^{-1} at $k = 3$, so the coincidence of dimensions is suggestive and no more.

Theorem 53 (All large n , for $k = 3$ and $k = 4$).

$$Z(P(n, 3)) = M(P(n, 3)) = 8 \quad \text{for every } n \geq 17,$$

$$Z(P(n, 4)) = M(P(n, 4)) = 10 \quad \text{for every } n \geq 30.$$

Proof. For $k = 3$, identity tiles exist for every ℓ with $21 \leq \ell < 42$, and for $k = 4$ for every ℓ with $30 \leq \ell < 60$; each is certified as above. (At $k = 3$ we in fact have certified tiles for every $21 \leq \ell \leq 43$.) A certification exhibits a box containing a true real solution of $AK = 0$ with K of full rank $2k + 2$, so the one-tile cyclic matrix has nullity at least $2k + 2$ and hence exactly $2k + 2$ by Corollary 12, so $T_{\text{tile}} = I$; and its docking weights are the prescribed integers exactly, since they are held fixed throughout, so tiles with different ℓ compose. Theorem 51 with $L = 21$ and $L = 30$ respectively gives $n \geq 21$ and $n \geq 30$, and Theorem 44 covers $17 \leq n \leq 20$ at $k = 3$. $\square$

All $23 + 32 = 55$ certifications pass. To give the numbers at both ends: at $k = 3$, $\ell = 22$ has 398 unknowns and 352 equations of which 324 are structural and $28 = \binom{8}{2}$ dependent, with 40 docking pins, high-precision residual 3.1×10^{-61} , and Krawczyk $\alpha = 6.7 \times 10^{-11}$; at $k = 4$, $\ell = 61$ has 1425 unknowns and 1220 equations of which 1175 are structural and $45 = \binom{10}{2}$ dependent, with 50 docking pins, residual 4.7×10^{-61} , and $\alpha = 3.7 \times 10^{-10}$. Across all 55, $\alpha \leq 3.8 \times 10^{-10}$, every certified matrix has nullity exactly $2k + 2$, every docking weight matches its prescribed integer to 10^{-14} , and the minimum edge weight is 0.15 of the matrix scale. See `verification/certify_tiles.py`.

These are the statements Conjecture 47 asks for, and they are proofs rather than computations reported as such: the chain is Theorem 51, Lemma 50 and Corollary 12, all proved, together with 55 interval certifications. No step rests on a floating-point nullity. With Theorem 20 for $k = 2$, the three smallest k beyond the trivial case are now settled for all large n , with

$$L(2) = 9, \quad L(3) = 17, \quad L(4) = 30,$$

against the tile lower bound $2(k + 1) + (2k + 2)^2/5$, which is 13.2, 20.8 and 30.

Theorem 54 (Complete determination for $k = 2$ and $k = 3$). *For every n with $2k < n$, the value $Z(P(n, k))$ is now known when $k = 2$ or $k = 3$:*

n	5	6	7	8	9	10	11	≥ 12	n	7	8	9	10	11	12	≥ 13
$Z(P(n, 2))$	5	4	6	5	6	6	6	6	$Z(P(n, 3))$	6	6	6	8	7	7	8

No value is left open in either row.

Proof. The tails are Theorem 20 for $k = 2$ and Theorem 55 for $k = 3$; the finitely many remaining values are exhaustive searches over vertex subsets. $\square$

For $k = 2$ this completes what [1] began, which gave $Z(P(n, 2)) = 6$ only for $n \geq 10$. For $k = 3$ it is the first determination of any kind: the value claimed in [1] for that case is false, and [2], which found the error, leaves the threshold open. Note that neither row is monotone and neither is eventually constant from the first time the bound is attained — $k = 3$ reaches 8 at $n = 10$, drops to 7 at $n = 11, 12$, and only then settles. Any argument that assumed monotonicity would have got this wrong, and one did.

The same exhaustive computation at $k = 4$ gives

n	9	10	11	12	13	14	15	16	17	18	19	20	21	22
$Z(P(n, 4))$	6	6	7	6	8	8	9	8	9	10	10	10	10	10

which is more erratic still: the value falls at $n = 12$ and again at $n = 16$, and the bound 10 is first attained at $n = 18$. With Theorem 53 for $n \geq 30$ and certificates at $n = 26, 27$, the only values of $Z(P(n, 4))$ not determined are $n \in \{23, 24, 25, 28, 29\}$.

Theorem 55 ($k = 3$ completely, and the stabilization threshold).

$$Z(P(n, 3)) = 8 \quad \text{for every } n \geq 13,$$

and 13 is optimal: $Z(P(11, 3)) = Z(P(12, 3)) = 7$. Moreover $Z(P(n, 3)) = M(P(n, 3)) = 8$ for every $n \geq 17$ and for every n divisible by 10.

Proof. For $n \geq 17$ this is Theorem 53. For $n = 13, 14, 15, 16$ exhaustive search over vertex subsets shows no forcing set of size 7 exists, so $Z = 8$ by the bound $Z \leq 2k + 2$. The values $Z(P(11, 3)) = Z(P(12, 3)) = 7$ are likewise exhaustive. $\square$

This proves the conjecture of [2]. That note closes with

“We conjecture $Z(P(n, 3)) = 8$ for $n \geq 13$; the missing ingredient is a lower-bound proof valid for all large n .”

Theorem 55 is that statement, and the missing ingredient is exactly what Section 11 supplies: $M \leq Z$ turns a matrix into a lower bound, and the tiling construction produces such matrices for all large n at once rather than one n at a time. The upper bound $Z \leq 8$ for $n \geq 9$ is already in [2], as are the values for $7 \leq n \leq 20$ by exhaustive search; what is new here is the lower bound for every $n \geq 13$, and with it the stabilization threshold, which the same note records as open alongside a rigorous lower bound for $k \geq 4$. The sequence is genuinely non-monotone below it — $Z(P(n, 3))$ runs 6, 6, 6, 8, 7, 7, 8, 8, ... for $n = 7, \dots, 14$ — which is why the threshold is not simply the first n attaining the bound, and why Conjecture 8 is stated in the form it is.

What a general existence proof would need, and what it reduces to

Theorems 51 and 53 together have an unsatisfying shape: the implication is proved for all k , but its hypothesis is verified by *computing* the tiles, separately for each k . A proof that identity tiles exist for every k and every sufficiently long ℓ would give $Z(P(n, k)) = M(P(n, k)) = 2k + 2$ for all k and all $n \geq L(k)$ in one statement, and would subsume Sections 8 and 10 entirely. This subsection identifies what that proof is a proof *about*.

Fix k and a docking pattern, and for $\ell \geq 2(k + 1)$ let $\mathcal{S}_\ell \subseteq \mathrm{GL}_{2k+2}(\mathbb{R})$ be the set of tile products of length ℓ realisable with all weights nonzero. Existence of an identity tile of length ℓ is exactly the statement $I \in \mathcal{S}_\ell$. Three facts constrain $\mathcal{S}_\ell$.

Observation 56 (Symplectic structure). For a fixed docking pattern there is a skew-symmetric B , unique up to scale, with

$$P^\top B P = B \quad \text{for every } P \in \mathcal{S}_\ell \text{ and every } \ell,$$

and B is nondegenerate. Hence $\mathcal{S}_\ell \subseteq \mathrm{Sp}(B) \cong \mathrm{Sp}(2k + 2, \mathbb{R})$, a connected simple Lie group of dimension $\binom{2k+3}{2}$.

The mechanism is not special to this problem: the recurrence of Theorem 11 is self-adjoint, because A is symmetric, and the transfer matrices of a self-adjoint difference system carry a Wronskian-type skew form from one step to the next. A tile whose two ends carry the same docking therefore returns that form to itself. Concretely, solving $T_i^\top B_{i+1} T_i = B_i$ along a tile has a $(k + 1)$ -dimensional solution space for any one tile, and imposing the same docking across many different interiors cuts it to the single B above.

We label this an observation rather than a theorem: we verified it numerically for $k = 2, \dots, 6$ at two tile lengths each, finding in every case exactly one invariant skew form, nondegenerate, preserved to between 10^{-13} and 10^{-8} , and we have not written out the algebraic proof (`verification/symplectic_structure.py`).

Observation 57 (Full dimension). The rank of the differential of $w \mapsto P_\ell(w)$ equals $\binom{2k+3}{2}$ exactly — the dimension of $\mathrm{Sp}(B)$ — for every (k, ℓ) we computed, $2 \leq k \leq 6$. So for ℓ past the threshold of Remark 52, $\mathcal{S}_\ell$ has nonempty interior in $\mathrm{Sp}(B)$.

Lemma 58 (Concatenation). $\mathcal{S}_m \cdot \mathcal{S}_\ell \subseteq \mathcal{S}_{\ell+m}$.

Proof. Immediate from Lemma 50: the docking pattern makes the product over each tile independent of its neighbours, so concatenating a tile of length m after one of length ℓ realises the product as a tile of length $\ell + m$. $\square$

Together these say that $\mathcal{S} = \bigcup_\ell \mathcal{S}_\ell$ is a subsemigroup of $\mathrm{Sp}(2k + 2, \mathbb{R})$ with nonempty interior, and the question “do identity tiles exist?” becomes

is I in the interior-carrying subsemigroup generated by these transfer matrices?

That is a controllability question about a Lie semigroup, and the honest status is that it does not follow from the above alone: $\mathrm{Sp}(2n, \mathbb{R})$ does admit proper subsemigroups with nonempty interior — compression semigroups of convex cones are the standard examples — so one must show the generators are not confined to such a wedge. What the observations do establish is that the group is identified, the dimension obstruction is absent, and the remaining question is of a recognised type rather than an unstructured search. We note that I is in fact reachable for every k where we have looked ($k = 3$ at all $21 \leq \ell \leq 43$, $k = 4$ at all $30 \leq \ell \leq 61$, $k = 5$ at 21 lengths), which is evidence for the general statement and not a proof of it.

Remark 59 (What this does not give). Theorem 53 leaves $7 \leq n \leq 16$ open at $k = 3$, apart from $n = 10$ which Theorem 24 settles; the construction of Section 11 fails at $n = 13$ and 15, where (3) gives the solution variety dimension only 3 and 9. At $k = 4$ it leaves $9 \leq n \leq 30$ open apart from the multiples of 60, 70 and 90 covered by Theorem 35.

It also gives nothing for $k \geq 5$, and here we are less able to say why than we would like. The count $\ell \geq 2(k+1) + (2k+2)^2/5$ gives $L = 41$ at $k = 5$ and $L = 54$ at $k = 6$. We searched $41 \leq \ell \leq 45$ at $k = 5$ and found no identity tile: the residual stalls between 7×10^{-2} and 2×10^{-1} , both from the degenerate system $T_{\text{tile}} = I$ and from the full-rank structural system started at random. Nothing in Lemma 50 or Theorem 51 refers to k , and the certification scaled from 8×8 transfer matrices at $k = 3$ to 10×10 at $k = 4$ with no modification at all, so we see no obstacle of principle.

We considered the natural explanation — that the product of ℓ transfer matrices becomes too ill-conditioned for the residual $T - I$ to carry information — and it does not survive contact with the numbers: at $k = 5$, $\ell = 44$ the product has condition number about 2.5×10^6 , whereas at $k = 4$, $\ell = 61$, which *does* yield a certified tile, it is about 1.1×10^{10} . So conditioning is not the discriminating factor and we do not have a diagnosis. The most likely mundane explanation is restart budget: we gave each $k = 5$ length four to six restarts, and at $k = 4$ every length that failed at that budget succeeded when given fifteen to sixty. On the evidence we have, the right statement is that we did not find $k = 5$ tiles, not that they do not exist.

One practical observation, since it governs how much compute the method needs. The Newton solve for a tile is erratic: at the restart budget we first used, roughly half the lengths at $k = 4$ failed, including lengths with *more* unknowns than ones that succeeded. Every such failure that we retried with a larger budget — $\ell = 31, 32, 33, 35, 38$ — then solved to 10^{-14} . So a failed tile search at these sizes carries no information beyond the effort spent, and we were careful not to read the first pass as evidence about $k = 4$. Extending a neighbouring tile as a starting point does *not* substitute for restarts: the damped Newton stalls near 10^{-2} from such starts at the lengths that needed retrying.

12 Why lower bounds resist standard methods

A set is zero forcing iff it meets every *fort* — a nonempty F such that every vertex outside F has 0 or ≥ 2 neighbours in F — so Z is a minimum hitting set over forts. Three standard consequences all fail here, and all degrade as k grows.

- **Disjoint forts.** Pairwise disjoint forts $F_1, \dots, F_t$ force $Z \geq t$. But forts in $P(n, k)$ have 7–13 vertices while the graph has $2n \leq 56$, and a greedy search found at most *two* disjoint forts in every case tested: the certificate cannot exist at these sizes.
- **The fort LP.** Computing the exact LP value by cutting planes gives $\text{LP}/Z \approx 0.43\text{--}0.50$, e.g. $\text{LP} = 3.00$ against $Z = 7$ for $P(12, 3)$, and 3.21 against $Z = 10$ for $P(18, 4)$. The integrality gap is about 2 and widens with k . This also explains why fort-based solvers converge slowly on this family: branch and bound must close a factor-2 gap by branching.
- **Spectral.** The Hoffman-type bound $Z \geq N(1 + 2\lambda_{\min}/(d - \lambda_{\min}))$ yields values in 1.2–2.9 against true Z of 6–12, and also weakens as k grows.

13 Status of the claims

Every claim below is labelled by the standard it meets. Four standards appear, in decreasing strength: *proved*, meaning a proof in the text; *proved by an exact certificate*, meaning an explicit

matrix whose nullity is computed in $\mathbb{Z}$ or $\mathbb{Q}(\zeta_n)$; *proved by interval certification*, meaning a Krawczyk contraction bound establishing that a true real solution exists in an explicit box; and *numerical*, meaning a floating-point computation only. We also list what we have refuted, including our own earlier claims.

- **Proved.** Theorems 3, 7 and Theorem 11 with Corollary 12; Propositions 13 and 14; Theorem 29; Theorem 30 and Theorem 26, the nullity classifications of the adjacency matrix and of its signed variant; Theorem 16, that $c^2 = 0$ if and only if the symbol factors as $(s+a)(L_k+t)$, with Corollaries 17 and 18 and Proposition 19; Theorem 28, that no rational period-one symbol attains 14 at $k = 6$ in either matrix class; Theorem 38, that a period-one certificate covers exactly one divisibility class, which contains at most one prime, together with the identity $e_2 = -k$ for every $k \geq 3$; and Theorem 31 with Theorem 32 for cyclic covers.
- **Proved by an exact certificate.** $Z(P(n, 2)) = 6$ for every $n \geq 9$ with $n \neq 10$, and for every n divisible by 7 (Theorem 20); $Z(P(n, 3)) = 8$ for every n divisible by 10 (Theorem 24); $Z(P(n, 4)) = 10$ for every n divisible by 60, 70 or 90; $Z(P(n, 5)) = 12$ for every n divisible by 24; $Z(P(n, 7)) = 16$ for every n divisible by 120 (Theorem 27). The $k = 4, 5$ values were the first exact values for $k \geq 4$; the $k = 7$ value replaces a floating-point claim we previously made at a single n . In the $k = 2$ ($7 \mid n$), $k = 3$ and $k = 7$ cases the certificate is an integer matrix with entries in $\{0, \pm 1\}$.
- **Proved by interval certification.** $Z(P(n, 3)) = M(P(n, 3)) = 8$ for every $n \geq 17$, and $Z(P(n, 4)) = M(P(n, 4)) = 10$ for every $n \geq 30$ (Theorem 53), from Theorem 51 together with interval certifications of identity tiles of every length in $[21, 44)$ and $[31, 62)$, plus Theorem 44 for $17 \leq n \leq 20$. No congruence condition appears and the primes are included. Each of the 55 certifications is a Krawczyk contraction bound with $\alpha \leq 3.8 \times 10^{-10}$; no step rests on a floating-point nullity. Each rests on a Krawczyk contraction bound for the bilinear system $A(w)K(X) = 0$, with the square subsystem chosen structurally so that it is equivalent to $AK = 0$ rather than merely a maximal-rank selection of it. By Theorem 38 no period-one certificate can reach these n at all, so this is the only route to them that we know.
- **Proved (finite).** Every tabulated value of Z , by exhaustive search over vertex subsets. Combined with the above this gives Theorem 54: $Z(P(n, k))$ is determined for every admissible n when $k = 2$ and when $k = 3$, with no value left open in either case. In particular Theorem 55, $Z(P(n, 3)) = 8$ for every $n \geq 13$ with 13 optimal: the four values $n = 13, 14, 15, 16$ are exhaustive searches and the rest is Theorem 53, so the stabilization threshold at $k = 3$ — recorded as open in [2] — is exactly 13. This is the one result here whose statement needs no matrix theory to check.
- **Numerical, not exact.** Theorem 40: $Z(P(n, 6)) = 14$ at $n = 60, 72, 84, 90, 96, 108, 120$ and $Z(P(96, 8)) = 18$. The certificates are period-three matrices whose weights come from a Newton solve, verified on the full $2n \times 2n$ matrix with a wide spectral separation and every required entry a substantial fraction of the matrix scale, but not in exact arithmetic. We do not claim these at the standard of the items above. ($Z(P(96, 7)) = 16$ was in this category and is now exact.)
- **Observation supported by computation, not proved.** The rank law $\text{rank} = \min(2d + 1, k + 1)$, verified for $3 \leq k \leq 13$ and $1 \leq d \leq 7$. Also the saturation of the reachable set of monodromies, at dimension 31 for $k = 2$ and 51 for $k = 3$.

- **Conjectural.** Conjecture 8, and Conjecture 47 for $k \geq 5$. For $k = 3$ and $k = 4$ the conjecture is now a theorem (Theorem 53); what remains open is whether $N(k)$ exists for every k , for which we prove only the necessary condition (3).
- **Refuted (ours).** That the period-one construction “never exceeds 6 for every $k \geq 2$ ”. Three roots can be prescribed and no more, but the remaining $k - 2$ can land on the grid by Galois conjugacy, and do. The period-two construction we built to get past 6 was unnecessary and its values are superseded.
- **Refuted (ours).** That $k = 3$ and $k = 7$ admit no rational period-one symbol attaining $2k + 2$. That enumeration imposed $c^2 > 0$, the condition for a *symmetric* certificate, where $c^2 \neq 0$ suffices: c^2 is the product of the two spoke weights and only its vanishing is fatal. Six symbols at $k = 3$ and one at $k = 7$ survive, giving Theorems 24 and 27.
- **Refuted (ours).** The unconditional bound $Z(B^n) \geq D$ for cyclic covers, by $Z(P(8, 2)) = 5 < 6$. Only the conditional form $Z(B^n) \geq \nu$ for an attained ν survives.
- **Refuted (ours).** The single-threshold form $n_0(k) = 5k - 2$, by $Z(P(29, 6)) = 12$; and any bound of the form $Z \leq 2(\gcd(n, j) + \gcd(n, k))$, by $Z(P(11, 3)) = 7 > 4$.
- **Refuted (ours).** That the threshold for reaching the ceiling with a general matrix is $n \geq (2k + 2)^2/5$. Spending the gauge gives (3), $n \geq (k + 1)(2k + 3)/3$, which is larger; the earlier figure led us to test $k = 6$ at $n = 27, 29, 31$, all below the true threshold of 35.
- **Refuted (ours).** That $N(k)$ in Conjecture 47 is $O(k)$. The saturation data alone suggests $4k + 3$, but (3) forces $N(k) \geq \lceil (k + 1)(2k + 3)/3 \rceil$, which is quadratic.
- **Refined (ours).** We earlier recorded as refuted the claim that the monodromy is symplectic, on the grounds that the *standard* form is not preserved. That is true and was the wrong conclusion to draw. With the docking pattern fixed there is a unique invariant skew form B , and it is nondegenerate, so tile products are symplectic with respect to B (Observation 56). The earlier statement — that the space of invariant skew forms for a single product has dimension $k + 1$, so eigenvalues come in reciprocal pairs — remains correct; it is the family, not the individual, that pins B down.
- **Observation, not proved.** Observations 56 and 57: the invariant form and the full-dimensionality of the reachable set, verified for $2 \leq k \leq 6$ but not proved. These are what reduce the existence of identity tiles for all k to a controllability question in $Sp(2k + 2, \mathbb{R})$, stated in Section 11. We do not claim that reduction settles the question: subsemigroups with nonempty interior in $Sp(2n, \mathbb{R})$ need not be the whole group.
- **Inconclusive, not negative (ours).** The per- n search of Section 11 at $k = 6$ returned nothing at $n = 41, 43, 45, 47, 49, 51$, and the tile search at $k = 5$ nothing at $41 \leq \ell \leq 45$. Both ran at a restart budget that we have since seen condemn a length that does solve: $\ell = 21$ at $k = 3$ stalled at 5.6×10^{-2} on four restarts and produced a tile at 1.4×10^{-14} on twenty-five. We therefore record these as searches we did not push far enough, not as evidence about $k = 5$ or $k = 6$. The one regularity we have found is that every length eventually solved first appeared at a residual of 10^{-2} or below, and the $k = 5$ and $k = 6$ attempts sit at 10^{-1} .
- **Withdrawn as unreliable (ours).** The eigenvalue-minimising search whose plateaus we previously reported as evidence that the ceiling is unreachable at various (n, k) . It gave

maximum nullity 7 at $P(20, 3)$, where 8 is attained — by Theorem 24 and, in the symmetric class, by the method of Section 11. It collapses onto the degenerate strata (all spokes zero, or all edges zero). Its negative results should not be read as evidence, and the conclusion we drew from them, that exactness for $k \geq 6$ was unlikely, was not supported.

- **Open.** Conjecture 47 for $k \geq 5$, where we have not attempted the tiles: the count gives $L = 41$ and 41 tiles at $k = 5$, $L = 54$ and 54 at $k = 6$; whether $M = Z$ in general; exact rather than interval certification at the n of Theorem 44; and the rank law. We expect $M(P(10, 2)) < Z(P(10, 2)) = 6$, a strict gap, since the search fails there even though the reachable dimension is saturated, but we do not prove it and no (n, k) is proved to have $M(P(n, k)) < 2k + 2$.
- **Superseded (ours).** The claim that the period- d method “cannot settle prime n at all” and that a different construction would be needed there. That is correct for the period- d method and is now a theorem (Theorem 38), but it is no longer a limitation of the paper: Section 11 settles prime n for $k = 3$.
- **Not determined.** Whether Theorem 7 is attained for $k \geq 3$; whether the degeneracy of Section 10 at $n = 48, 54$ is essential.

References

- [1] S. Rashidi, N. Shajareh Poursalavati, M. Tavakkoli, Computing the zero forcing number for generalized Petersen graphs, *J. Algebra Combin. Discrete Struct. Appl.* 7(2) (2020), 183–193.
- [2] A. Krishnan, A correction to the zero forcing number of the generalized Petersen graphs $P(n, 3)$, arXiv:2607.19412 (2026).
- [3] J. H. Conway and A. J. Jones, Trigonometric diophantine equations (On vanishing sums of roots of unity), *Acta Arith.* 30 (1976), 229–240.
- [4] I. Niven, *Irrational Numbers*, Carus Mathematical Monographs 11, Mathematical Association of America, 1956, Corollary 3.12.

IV. Research Log

The dated research log follows this page. It records the work in the order it happened rather than as a finished account, and it deliberately preserves what went wrong: conjectures that were tested and died, claims of the author's own that were later refuted, and software bugs that produced confident wrong answers before being caught.

It is included because a reader is entitled to see which conclusions were reached first and revised later, and because the superseded versions have been left in place rather than edited away.

Investigation Log: Zero Forcing Numbers of $P(n, k)$ for $k \geq 4$

Aryan Padarthi

September 23, 2026

Entry 1: How this problem was chosen

Date: September 23, 2026

This is the ninth research target attempted in this project. The previous eight were abandoned, and the reason matters: in every case the problem turned out to be already solved, and in every case I had inferred that it was open from the *absence* of a published result rather than from a positive statement.

- Four attempts on Turaev genus and signature bounds for positive braid closures — an area with an active professional research group.
- Periodicity of the octal game 0.45. I derived a real structural mechanism, then found the Guy–Smith periodicity theorem (1956), which settles it outright once one has computed past $n = 1018$. I had verified to 60,000.
- A “rare-index invariant” I derived for octal games, which is exactly Berlekamp’s Sparse Space Phenomenon, described in *Winning Ways* c. 1980.
- Persistent homology of linkage configuration spaces. Farber–Schütz (2007) give closed Betti-number formulas for planar polygon spaces, and Ba–Zhou (arXiv:2506.06781, 2025) prove the self-avoiding moduli spaces are diffeomorphic to Euclidean space — i.e. contractible, so persistent homology would correctly compute nothing at all.

The corrected procedure: before any investment, ask *what standard result would make this trivial*, and require a positive statement of openness. This target passes that test. Krishnan (arXiv:2607.19412, July 2026) closes with: “Determining $Z(P(n, k))$ for general k — the stabilization threshold and a rigorous lower bound for $k \geq 4$ — remains open.” That is an explicit assertion in a paper two months old, not an inference from silence.

Entry 2: A validated engine

Date: September 23, 2026

No value was used for research before the engine reproduced every published value. The validation set:

- Krishnan’s Table 1, $Z(P(n, 3))$ for $7 \leq n \leq 20$: 66687788888888 — reproduced exactly, including the non-monotone spike at $n = 10$ and the corrected $Z(P(12, 3)) = 7$ that falsifies [1]’s Theorem 3.6.
- $Z(P(n, 2)) = 6$ for $10 \leq n \leq 16$; $Z(P(2k + 1, k)) = 6$ for $k = 5, 6, 7$.
- Independent cross-check against brute force on small cases.

The engine independently found the witness $\{u_0, \dots, u_6\}$ for $Z(P(12, 3)) = 7$ — precisely the configuration Krishnan identifies as the one the published proof never tested.

Entry 3: Three approaches that failed

Date: September 23, 2026

Recorded because the negative results were informative and cost real time.

(a) Fort-based ILP (scipy). Exact in principle: a set is zero forcing iff it meets every fort, so Z is a minimum hitting set. With a partial fort family the ILP is a relaxation, giving a valid lower bound, and any returned set is verified independently by running the forcing rule — so the method cannot report a wrong value, only fail to converge. It failed to converge: `scipy.optimize.milp` has no lazy-constraint callback, so every round re-solves from scratch. Stalled past $P(11, 3)$, which brute force does instantly.

(b) CP-SAT, time-indexed formulation. Correct on every instance it finished, all witnesses verified. Also far too slow — 30s on $P(13, 3)$, stuck on $P(15, 3)$. The obstruction is structural, not incidental: finding a forcing set is easy, *proving minimality* is hard, and the formulation’s LP relaxation is too weak to prune. The exponential cost moves from enumeration into branch-and-bound rather than disappearing.

(c) Disjoint forts as a lower-bound certificate. If $F_1, \dots, F_t$ are pairwise disjoint forts then $Z \geq t$, with no case analysis — which is exactly the kind of argument the open lower bound needs. This fails for an information-theoretic reason: forts in $P(n, k)$ have 7–13 vertices while the graph has only $2n \leq 56$, so a greedy search found at most *two* disjoint forts in every case, giving $Z \geq 2$. Not a tuning problem; the certificate cannot exist at these sizes.

What survived was the simplest thing: exhaustive enumeration in C with 128-bit masks, plus a rotation-symmetry reduction, plus threads.

Entry 4: Results

Date: September 23, 2026

Proved. For all $k \geq 2$ and $n \geq 2k + 3$, $Z(P(n, k)) \leq 2k + 2$, by an explicit forcing set $\{u_0, \dots, u_{2k+1}\}$ and a rotation-invariance bootstrap. The bound itself is [1]’s Theorem 2.2; the contribution is a proof uniform in k (Krishnan’s clean replacement covers only $k = 3$) in a setting where the original paper is known to contain a false theorem. Every step machine-verified for $2 \leq k \leq 8$.

Computed. Exact values of $Z(P(n, k))$ for $k = 4$ ($9 \leq n \leq 28$), $k = 5$ ($11 \leq n \leq 25$), $k = 6$ ($13 \leq n \leq 26$). These are new; the literature has no exact values for $k \geq 4$.

Conjectured, not proved. That $Z(P(n, k)) = 2k + 2$ for large n , with thresholds 10, 13, 18, 23 for $k = 2, 3, 4, 5$. I decline to fit a formula to four points; $4k + 2$ fails at $k = 3$ and $k = 5$.

Open. The matching lower bound. Nothing here bears on it.

Unresolved in our own data. Whether $Z(P(n, 6))$ ever attains 14. The maximum observed for $n \leq 26$ is 12, and values with $\gcd(n, 6) = 1$ appear to cap lower than those with $\gcd(n, 6) > 1$. Since the inner vertices form $\gcd(n, k)$ disjoint cycles this is structurally plausible, but with the ceiling unattained anywhere in range it is not claimed as a finding.

Entry 5: A correction I had to make to my own claim

Date: September 23, 2026

After seeing $Z(P(n, 4)) = 10$ for $n = 18, \dots, 22$ — five consecutive values — I concluded the sequence had stabilized. The $k = 5$ row then showed $Z = 10$ at $n = 17, \dots, 21$, also five consecutive values, followed by 11 at $n = 22$ and 12 at $n = 23$. Five repeats are not evidence of stabilization. The $k = 4$ claim survives only because $2k + 2$ is a *proven* ceiling there, so the sequence cannot climb — but it dropped at $n = 12$ and $n = 16$, so it could in principle drop again. This is the same warning Krishnan gives about extrapolating from finite computation, arriving on my own data within an hour of my reading it.

Entry 6: A strategic reassessment

Date: September 23, 2026

Asked to assess whether the work so far constituted a project, the honest answer was no. Stripped of presentation, the content was: new computed values, plus a *rederivation* of somebody else's theorem. The open problem — the lower bound — was untouched. Compared against the project that qualified from this region last year (an IFS construction on the Riemann ξ -function, which proves existence, uniqueness, an antipodal symmetry, and a dimension bound), one rederived theorem and a table is thin.

Three routes to original content were weighed:

- *The lower bound.* Highest value, lowest odds. It had already defeated a published proof, an automated prover (a logged attempt of 2026-09-16, closed without progress), and our own disjoint-fort attempt.
- *Exceptional values.* Tempting, since upper bounds are constructive. Abandoned after checking: the drops are transient. For $k = 4$, the values at $n \equiv 0 \pmod{4}$ run 6, 8, 10, 10, 10 — they climb to the ceiling like everything else, so there is no permanent exceptional family to characterise.
- *Generalise the method, not the family.* Chosen. The rotation bootstrap is not specific to $P(n, k)$; it needs only a cyclic automorphism.

The third worked, and it is what turned the day's output from a table into theorems.

Entry 7: Three lower-bound techniques, all insufficient

Date: September 23, 2026

Before abandoning lower bounds, each standard route was tested to destruction rather than assumed hopeless.

- **Disjoint forts.** t pairwise disjoint forts give $Z \geq t$ with no case analysis. Greedy search found at most **two** in every instance. Forts here have 7–13 vertices and the graph has $2n \leq 56$; packing $2k + 2$ of them is not a tuning problem but an impossibility.
- **The fort LP.** Computed exactly by cutting planes, separating with a minimum-weight-fort ILP each round. $LP/Z \approx 0.43$ – 0.50 throughout, worst at $P(18, 4)$: 3.21 against $Z = 10$. An integrality gap of about 2, widening with k .
- **Hoffman-type spectral bound.** Gives 1.2 – 2.9 against true Z of 6 – 12 ; also degrades in k .

This is a negative result worth stating: the lower bound resists every general-purpose technique available, which is why it is open. It also retroactively explains the earlier solver failures — branch and bound was being asked to close a factor-of-two gap by branching.

Entry 8: Two theorems

Date: September 23, 2026

The bootstrap was isolated into two elementary lemmas (closure commutes with automorphisms; a filled outer set forces the inner set) and then applied.

$Z(I(n, j, k)) \leq 2(j+k)$ **for** $n \geq 2(j+k)+1$. The chain is: u_i forces v_i for $j \leq i \leq m-1-j$; then v_{j+k} forces v_{j+2k} ; then u_{j+2k} forces u_m . Taking $m = 2(j+k)$ makes the second step work because $m-j-2k = j$ is exactly the first index filled in step 1. Verified in 106 cases. Setting $j = 1$ recovers $Z(P(n, k)) \leq 2k + 2$, so the earlier rederivation is now a corollary rather than a separate result.

$Z(DP(n, k)) \leq 4k + 4$ **for** $n \geq 2k + 3$. The same chain, doubled and run in parallel on both outer cycles. Verified for $1 \leq k \leq 5$. No zero forcing results for this family were found in the literature. Our computed values give $Z(DP(n, 1)) = 8 = 4 \cdot 1 + 4$ for $6 \leq n \leq 12$, and subsequently $Z(DP(10, 2)) = 12 = 4 \cdot 2 + 4$, so the bound is attained at $k = 1$ and $k = 2$ — it is sharp there, not merely an upper estimate.

Entry 9: A prediction, recorded before the test

Date: September 23, 2026

The thresholds at which $Z(P(n, k))$ first attains $2k + 2$ are $10, 13, 18, 23$ for $k = 2, 3, 4, 5$. For $k \geq 3$ these are exactly $5k - 2$; $k = 2$ is the exception. The formula predicts $n_0(6) = 28$.

At the time of writing the $k = 6$ computation has reached $Z(P(27, 6)) = 13$, which is consistent (the ceiling 14 is not yet attained) but does not test the prediction. $P(28, 6)$ is running. This is recorded now, before the answer is known, so that a confirmation counts as a test rather than a fit — the same discipline that Entry 5 failed to apply.

Entry 10: The prediction was tested, and the conjecture died

Date: September 24, 2026

Entry 9 recorded, before the computation finished, that $n_0(6) = 5 \cdot 6 - 2 = 28$ and hence $Z(P(28, 6)) = 14$. The run took just over fourteen hours.

$Z(P(28, 6)) = 14$, with witness $\{u_0, \dots, u_{13}\}$. The narrow prediction was correct: 28 is the least n at which $k = 6$ attains its ceiling, exactly as $5k - 2$ says.

$Z(P(29, 6)) = 12$. And this refutes the conjecture. It was stated as “ $Z(P(n, k)) = 2k + 2$ for all $n \geq n_0(k)$ ” — a *stabilization* claim — and $29 > 28$ with $Z = 12 \neq 14$. So the prediction and the conjecture came apart: the thing I could check quickly was right, and the thing it was evidence for was wrong.

The cause is $\gcd(n, k)$, which had been visible in the $k = 6$ data for some time and which I twice set aside because $k = 4$ and $k = 5$ looked uniform. Sorting by $d = \gcd(n, 6)$, each class climbs on its own schedule: $d = 2$ reaches 14 at $n = 28$ while the coprime class has only reached 12 by $n = 29$. At $k = 4$ every class had already converged inside the computed range, so no separation was visible; at $k = 6$ the range finally became wide enough to see it. The apparent thresholds 10, 13, 18, 23 are where the slowest class happened to arrive, not values of a formula.

Two things worth recording about the process rather than the result. First, the prediction was useful precisely because it was written down in advance — had I fitted $5k - 2$ after seeing $Z(P(28, 6)) = 14$, I would have called it a confirmation and never computed $n = 29$. Second, this is the third time in two days that a pattern read off a short run has failed on the next data point (five repeats at $k = 5$; “ Z tracks n ” for I-graphs; now $5k - 2$). The consistent lesson is that in this family, sequences look stable well before they are.

Entry 11: A second conjecture, refuted by data I already had

Date: September 25, 2026

From 20 computed I-graphs, all satisfying $Z \leq 2(d_1 + d_2)$ with $d_1 = \gcd(n, j)$, $d_2 = \gcd(n, k)$, I wrote a conjecture into the paper asserting that bound in general. It is false, and the counterexamples were already in the validation table.

Every one of those 20 graphs had $d_1, d_2 > 1$ — that was the selection criterion for “genuinely new” I-graphs. Setting $j = 1$ gives $d_1 = 1$, and with $\gcd(n, k) = 1$ as well the bound reads $Z \leq 4$, while $Z(P(n, k)) = 2k + 2$ grows without bound. $Z(P(11, 3)) = 7$ and $Z(P(17, 4)) = 9$ refute it immediately. I had computed both values days earlier.

The error was generalising from a sample selected by a condition, to a statement about everything, without testing the boundary the condition excluded. The previous conjecture ($n_0 = 5k - 2$) at least took a fourteen-hour computation to kill; this one needed a table lookup.

What survives is narrower and still useful. The observation holds as stated for $d_1, d_2 > 1$, and a separate computation supports the underlying intuition: the minimum *bootstrapping*

set for $I(18, 3, 8)$ — the object a uniform-in- n proof actually needs — has size $10 = 2(d_1 + d_2)$ against $2(j + k) = 22$, and decomposes as three consecutive vertices in each of the three outer cycles plus one inner vertex. Three consecutive vertices in an outer cycle is precisely what lets the middle one force its spoke. So the route to beating the $Z(I(n, j, k)) \leq 2(j + k)$ bound is a proof adapted to the cycle decomposition rather than to consecutive indices — but the target it proves must retain some dependence on j and k , not on d_1, d_2 alone.

Entry 12: A lower bound, at last

Date: September 25, 2026

Three combinatorial routes to a lower bound had failed (disjoint forts, the fort LP, a spectral bound), each degrading as k grows. The route that works is not combinatorial at all: $Z(G) \geq M(G)$, the maximum nullity over symmetric matrices with G 's pattern, so a single explicit matrix *proves* a bound. $P(n, k)$'s rotation symmetry makes such matrices computable.

Three stages, each with a lesson.

$d = 1$. Taking A equivariant under ρ splits it into n blocks of size two, and the singularity condition is linear and homogeneous in four coefficients — a nullspace computation, not a search, so the blocks are singular by construction. This gave $Z(P(12, 2)) = Z(P(14, 2)) = 6$ in both directions, the project's first lower bound.

$d = 2$, **and a numerical trap**. Relaxing to ρ^2 gives ten parameters and 4×4 blocks. My first attempt solved for them numerically and reported nullity 7 for $P(14, 2)$, where $Z = 6$. Since $Z \geq M \geq$ nullity that is impossible, so the theory caught the bug: the optimiser had shrunk a required-nonzero edge weight to 3×10^{-4} against a matrix scale of 4.5, numerically deleting an edge while still passing the zero-pattern test. I now verify by block determinants, which are zero by construction, rather than by thresholding eigenvalues.

Doing $d = 2$ algebraically instead revealed a real obstruction: six coefficients suggest five fitted points, but fitting five always forced $c_r^2 \leq 0$ — the spoke weight vanishing, decoupling the two halves. Fitting *four* leaves a one-parameter pencil containing realisable members. That gave 8 for $P(18, 4)$ and 10 for $P(24, 6)$, and unlike $d = 1$ the bound now grows with k .

Uniformity, via Chebyshev. The certificates were still found one n at a time. But $s = 2 \cos \theta$ and $t = 2 \cos k\theta$ means $t = 2T_k(s/2)$ exactly, so the condition is a polynomial in s alone of degree $k + 1$ and the singular blocks are its roots. For $k = 2$ the cubic's coefficients are free, so any three s -values can be prescribed — a construction valid for every n , and 6 is maximal because a cubic has three roots. This also explains the flat cap I had observed: for $k \geq 2$ the $d = 1$ polynomial is forced into a shape with only three free ratios.

The same substitution for $d = 2$ gives degree $k + 1$ in x , ceiling $2k + 2$ — exactly the conjectured Z , with the shortfall due entirely to realisability. Carrying it out for $k = 4$ yields nullity 8 for $n = 18$ and every even $22 \leq n \leq 40$, so $8 \leq Z(P(n, 4)) \leq 10$ uniformly. No lower bound on $Z(P(n, k))$ for $k \geq 3$ appears in the literature we have seen.

What I take from this: the three failures were all of one kind — asking a combinatorial relaxation to certify a quantity it cannot see — and the fix was to change the category of argument rather than to work harder inside it.

Entry 13: The ceiling, and that it is exactly the upper bound

Date: September 25–26, 2026

Command: `python verification/recurrence_order.py`

I had been describing the certificate method as “capped near 6–8” without knowing what the cap actually was. It has an exact value, and finding it changed the project.

Every spoke weight c_i is nonzero, so the kernel equation at u_i solves for y_i as a 3-term expression in x . Substituting into the equation at v_i leaves one scalar recurrence in x whose support is the nine indices $\{i-k-1, i-k, i-k+1\} \cup \{i-1, i, i+1\} \cup \{i+k-1, i+k, i+k+1\}$, with both end coefficients nonzero. So the recurrence has order exactly $2k+2$, and $\text{null } A = \dim \ker(T - I)$ for its monodromy T : at most $2k+2$, with equality iff $T = I$.

Two things fell out that I did not expect. First, the argument never uses symmetry — the bound $\text{null } A \leq Z$ needs only a nonzero entry in each direction along each edge — so the whole combinatorially symmetric class, $8n$ parameters instead of $5n$, is available. Second, and this is the part worth remembering: the order $2k+2$ is the statement that no kernel vector vanishes on $u_0, \dots, u_{2k+1}$, and that set is exactly the forcing set the rotation bootstrap produces in Theorem 1. *The ceiling of the certificate method is the upper bound it is trying to match.* There is no slack anywhere; either the method proves $Z = 2k+2$ or nothing does.

Checked before use, as usual: 189 random matrices over 63 pairs (n, k) for the sparsity, the end-coefficient formulas and $\text{null } A = \dim \ker(T - I) \leq 2k+2$; 60 more for $\det T = (\prod b')(\prod e')/((\prod b)(\prod e))$, for the product form $\text{null } A = \dim \ker(SU - I)$, and for the forcing-set statement.

Entry 14: A claim of mine, refuted — and the refutation is the result

Date: September 26, 2026 **Command:** `python verification/d1_ceiling.py`

Entry 12 asserted that the period-one construction “never exceeds 6 for every k ”, because the symbol $F(s) = (s+a)L_k(s) + \alpha s + \beta$ has only three free parameters. I wrote a script to check that assertion exhaustively over all triples of grid roots, expecting it to confirm the claim. It printed 12.

The premise was right and the inference was wrong. At a root the condition is linear, so three roots determine (a, α, β) and no fourth can be imposed — but prescribing three roots *determines the other* $k-2$, and those can land on the grid $\Gamma_n = \{2\cos(2\pi m/n)\}$ by themselves. They do so exactly when arithmetic makes them: a symbol with rational coefficients has Galois-closed roots, and the Galois orbit of $2\cos(2\pi m/n)$ stays inside Γ_n . One rational F places all $k+1$ roots at once.

I had reasoned about how many roots I could *choose* and concluded something about how many roots could *be there*. The period-two machinery of Entry 12, built to escape a cap that was not real, is now redundant.

Entry 15: Cyclotomic certificates, and the first exact values for $k \geq 4$

Date: September 26, 2026

Commands: `cyclotomic_certificates.py`,
`period1_optimum.py`, `certified_lower_bounds.py`

Taking the symbol to be a product of minimal polynomials Ψ_d of $2\cos(2\pi/d)$, with $\sum \deg \Psi_d = k + 1$, and imposing the $k - 2$ integer conditions for membership in the family, gives certificates uniform in n :

$$\begin{aligned} Z(P(n, 4)) &= M(P(n, 4)) = 10 && \text{for every } n \text{ divisible by } 60, 70 \text{ or } 90, \\ Z(P(n, 5)) &= M(P(n, 5)) = 12 && \text{for every } n \text{ divisible by } 24. \end{aligned}$$

These are the first exact values of $Z(P(n, k))$ for $k \geq 4$; [2] records that range as open, and a week ago I had $8 \leq Z(P(n, 4)) \leq 10$ and nothing at all for $k = 5$. Every certificate has $a, \alpha, c \in \mathbb{Z}$, so it is an integer matrix and its nullity is confirmed by exact rational elimination.

The $k = 5$ case is the one I like. There $a = d = 0$ and $c = 1$: the certificate *is the adjacency matrix* of $P(n, 5)$, the symbol is $sL_5 - 1 = (s - 1)(s + 1)(s^4 - 4s^2 + 1) = \Psi_6\Psi_3\Psi_{24}$, and the whole statement reduces to counting m with $2\cos\frac{2\pi\cdot 6m}{n} + 2\cos\frac{2\pi\cdot 4m}{n} = 1$. The simplest matrix in the class was sufficient all along; I had spent weeks on elaborate ones.

Beyond the rational construction, the exhaustive triple scan over $n \leq 130$ gives the period-one optimum for each (n, k) ; those certificates were re-verified in $\mathbb{Q}(\zeta_n)$, with $2\cos(2\pi m/n)$ represented as $x^m + x^{n-m}$ modulo Φ_n , so the membership identity, the reality of the parameters and $c^2 \neq 0$ are decided exactly. It adds $Z(P(n, 3)) = 8$ for $n = 40, 60, 80, 120$ (a matrix proof of values [2] gets combinatorially) and $Z(P(n, 2)) = 6$ for $9 \leq n \leq 130$, $n \neq 10$. It also shows the honest limit: period one reaches $2k + 2$ for $k \leq 5$ and, over this range of n , gets only to 8, 10, 8 for $k = 6, 7, 8$ against ceilings 14, 16, 18.

Entry 16: Two numerical traps, both caught by the ceiling

Date: September 26, 2026

Counting grid roots by $|F(r)| < 10^{-8}$ reported nullity 7 for $k = 2$ at $n = 126, 128, 130$. That is impossible — the ceiling is $2k + 2 = 6$ — and the cause was Γ_n becoming dense enough that near-roots passed the test. Separately, a candidate at $(n, k) = (48, 7)$ with $|a|$ enormous against the edge weights made a threshold relative to $\max |A|$ report nullity 48.

Both were caught the same way the period-two trap was caught in Entry 12: by an inequality that must hold. Then it was $Z \geq M$; now it is the ceiling of Entry 13. It is worth stating as a working rule: *every numerical search should be wrapped in a bound it cannot legally violate*, because that bound is the only thing that will tell you the answer is wrong when it looks right. The fixes were to test each Fourier block against its own scale and to certify in $\mathbb{Q}(\zeta_n)$ rather than in floating point.

A third mistake, of a different kind: my first search of the combinatorially symmetric class reported a *smaller* maximum nullity on $P(12, 2)$ than the symmetric search did. Since

$S(G)$ sits inside that class, this cannot happen, so the search — not the mathematics — was at fault. Seeding the larger search with the symmetric certificate fixed it. Before that check I had drawn a conclusion from those numbers, namely that even $8n$ parameters could not reach $2k + 2$ at $k = 3$; the conclusion may still be true but the evidence I had for it was worthless.

Entry 17: The period- d parameter count is $2d + 1$, not $5d - 1$

Date: September 26, 2026

Command: `python verification/period_d_rank.py`

Entry 15 ended by naming period- d symbols as the fix for the shortage of parameters, and I wrote in the paper that a period- d symbol “carries $5d - 1$ free ratios, so $d \geq (k + 2)/5$ restores the count.” That was a guess dressed as a count, and it is wrong.

The honest way to ask the question is: what is the dimension of the set of symbols actually *reachable*? Most of the $5d$ weights are gauge — conjugating by a positive diagonal matrix preserves both the pattern and the nullity — so counting weights counts nothing. I computed the rank of the Jacobian of

$$(5d \text{ weights}) \longmapsto (\text{monic coefficients of } G(\sigma))$$

at random points, for $3 \leq k \leq 13$ and $1 \leq d \leq 7$. In every one of those cases,

$$\text{rank} = \min(2d + 1, k + 1).$$

At $d = 1$ that returns 3, which is exactly the three free parameters (a, α, β) of the period-one symbol — a check I did not build in, and the reason I believe the rest of the table.

So the reachable symbols form a $(2d + 1)$ -dimensional family, and every G is locally reachable precisely when $2d + 1 \geq k + 1$, i.e.

$$d \geq \lceil k/2 \rceil,$$

not $d \geq (k + 2)/5$. For $k = 6$ that means period *three*, not period two — which is presumably why the period-two work of Entry 12 stalled, and why my first attempts at $k = 6$ with $d = 2$ found nothing. I had been running the wrong experiment for a reason I had written down incorrectly the same day.

Two things I want to keep straight about this. It is an **observation supported by computation**, not a theorem: I have not proved the rank formula, only verified it over a range. And it is a statement about *local* reachability — a submersion means nearby symbols are reachable, not that a specific target G with all roots on the grid is. Realisability (real weights, no vanishing edge) remains a separate hurdle, and it is exactly the hurdle that killed the period-two construction and every rational candidate at $k = 3$.

What makes it worth the entry is the shape of the consequence. The cyclotomic certificates have moduli that are lcms of cyclotomic orders — 60, 70, 90, 24 — which is why the results read as a list of congruence classes. A period- d construction has modulus d . If the

count is the only obstruction, the right statement is not “ n divisible by 60” but “ n divisible by $\lceil k/2 \rceil$, once n/d is large enough to hold $k + 1$ grid roots” — one clean congruence for every k , instead of a table.

A methodological note. The rank computation took a few minutes to write and settled a question I had been attacking by search for hours. The search was asking “can I find a certificate here?”; the rank was asking “how many dimensions of certificate exist here at all?” The second question was cheaper and strictly more informative, and I should have asked it first.

Entry 18: $k = 6$, and what it cost to get there

Date: September 26, 2026

Commands: `prescribe_symbol.py`, `verify_period_d_cert.py`

$Z(P(60, 6)) = Z(P(72, 6)) = 14$. These are the first exact values at $k = 6$, a range where Entry 16 showed the period-one construction *provably* cannot reach the ceiling.

The route was not more search. Once the rank law said the symbol map is a submersion for $d \geq \lceil k/2 \rceil$, the right move was to stop searching for a certificate and start *prescribing* one: choose $k + 1$ interior grid values, form the monic G^* with those roots, and solve $G(w) = G^*$. That is a well-posed system with a full-rank Jacobian, where minimising the r smallest eigenvalues is a badly-posed one that leaves the optimiser room to cheat.

And it cheated exactly as you would expect. Every unconstrained solve returned a point with one spoke weight collapsed to zero — a real solution of $G(w) = G^*$, because with a spoke gone the matrix block-diagonalises and G factors as $\det U \cdot \det V$, but a matrix for a *different graph*. Three things failed to fix it: a realisability barrier (which demonstrably works — it improved the known $k = 5$ certificate’s margin from 0.288 to 0.321 without breaking it); continuation, which lost the symbol at the first step $\tau = 0.02$; and simply restarting, because these are *singular* points where the Jacobian rank drops from 7 to 6, and singular points are attractors.

What fixed it was noticing that the failures were all at $n = 48$ and $n = 54$, where Γ_m has 7 and 8 interior values against the 7 required — one admissible root set, then eight. At $n = 60$ and $n = 72$ there are 9 and 11, hence 36 and 330 root sets, and realisable points exist. The obstruction was never really the degeneracy; it was that I had been testing the two cases with no room to move.

A mistake worth recording. I had queued $n = 24, 30, 36, 42$ for the first $k = 6$ run. Every one of them is impossible by construction: my guard required $m \geq k + 2$, but attaining $2k + 2$ needs $k + 1$ distinct *interior* grid values, so the requirement is $m \geq 2k + 4$. Γ_8 has three interior values where seven are needed. Had those runs come back negative — and they would have — I would have read it as evidence against period three rather than evidence of a badly designed experiment.

A mechanism I proposed and killed. Seeing that the $n = 48$ target had constant term zero (because $4 \mid 16$ puts $\sigma = 0$ in the grid, and with no slack it had to be included), I proposed that a zero constant term forces the spokes to vanish, and predicted $n = 54$ would escape. It does not: $G(0)$ is nonzero when the spokes are zeroed, and does not scale with them. I had inferred a structural law from one data point, and the test took two minutes.

Status, stated precisely. These certificates are verified on the full $2n \times 2n$ matrix — not the Fourier blocks that found them — with fourteen eigenvalues below 4×10^{-15} , the fifteenth at 5×10^{-2} , the correct cubic pattern, and every required entry near a fifth of the matrix scale. Several independent root sets give certificates at each n , so this is a region and not a lucky point. But the weights are algebraic numbers from a Newton solve, not integers, so this is **numerically certified and not yet exact** — a weaker standard than the $k = 4$ and $k = 5$ results, and I am not going to describe it as though it were the same. Choosing the root set to be a union of full Galois orbits makes G^* rational and is the route to closing that gap; it is running.

Entry 19: The method scales — $k = 7$ and $k = 8$

Date: September 26, 2026

Command: `python verification/prescribe_symbol.py`

$Z(P(96, 7)) = 16$ and $Z(P(96, 8)) = 18$, on top of $Z(P(60, 6)) = Z(P(72, 6)) = 14$.

What makes this worth more than three extra values is that the period was *predicted*. The rank law says the symbol map becomes surjective at $d \geq \lceil k/2 \rceil$, so $k = 6$ needs period three and $k = 7, 8$ need period four. That is what happened: $d = 3$ carried $k = 6$ and failed for $k = 7$ and $k = 8$, where $d = 4$ succeeded. A law measured on one set of cases went on to predict, correctly, which experiments would work before they were run. That is the difference between a fitted pattern and a structural fact, and it is the main reason I now believe the rank law rather than merely reporting it.

The claim also changes shape. Before today the results were a list — $k = 4$ here, $k = 5$ there, on scattered congruence classes. Now there is a *method* with an exact ceiling equal to the known upper bound, a rule saying how much symmetry must be broken to reach it, and certificates attaining it for every k from 2 to 8.

Generality. A sweep over n for $k = 6$ at $d = 3$ finds certificates at $n = 60, 72, 84$ — the last two on the first or second root set tried. So this is not two lucky values of n ; once the counting conditions ($d \geq \lceil k/2 \rceil$ and $m = n/d \geq 2k + 4$) hold with a little slack in the choice of roots, certificates appear readily.

What is still weaker than it looks. All of $k \geq 6$ is numerically certified, not verified in exact arithmetic. The route I expected to fix that — choosing the root set to be a union of full Galois orbits, making the target symbol rational — is **closed**: all eight Galois-closed sets at $n = 72$ hit the symbol to 10^{-15} with nullity 14, and all eight collapse a weight, whether or not $\sigma = 0$ is among the roots. Rationality and realisability are in tension here, and I do not understand why. The remaining route is interval Newton, which the ceiling theorem makes tractable: certifying the *equation* $G(w) = G^*$ exactly gives the nullity as a discrete consequence, since $k + 1$ grid roots force $2k + 2$ singular blocks and the ceiling forces equality. That converts a fragile numerical rank claim into a genuine computer-assisted proof.

A process note. Twice today the decisive step was to stop searching and compute a dimension instead — first the rank of the symbol map, then the count of admissible root sets. Both took minutes and both redirected hours of search. The failures at $n = 48$ and $n = 54$ looked like a structural obstruction and were nothing of the kind: those are the two cases with one and eight admissible root sets, against thirty-six and three hundred and thirty

at $n = 60$ and $n = 72$. I had been testing the only cases with no room to move.

Entry 20: Enlarging the object, and a classification

Date: September 26, 2026

Commands: `cyclic_covers.py`,
`adjacency_classification.py`

Two changes of framing, both of which make the work about something larger than $P(n, k)$ without weakening anything.

Cyclic covers. $P(n, k)$ is a cyclic cover: a two-vertex base graph with loops of voltage 1 and k and a spoke of voltage 0, lifted to $V(B) \times \mathbb{Z}_n$. Nothing in the Fourier reduction ever used the particular voltages, only that $\mathbb{Z}_n$ permutes the fibres — so the block reduction, and the ceiling as the degree span of $\det M(\zeta)$, hold for *any* base graph. Verified on seven base graphs on two to five vertices. The right setting was covering-space theory all along; I had been proving a special case without noticing.

A classification, via Conway–Jones. The adjacency nullity is $\#\{m : 2 \cos \frac{2\pi(k+1)m}{n} + 2 \cos \frac{2\pi(k-1)m}{n} = 1\}$, a rational relation among cosines of rational angles. Clearing denominators with $z = \omega^m$ turns it into: which roots of $P_k(z) = z^{2k+2} + z^{2k} - z^{k+1} + z^2 + 1$ are roots of unity? So

$$\text{null}(\text{Adj } P(n, k)) = \sum_{\Phi_d | P_k, d|n} \varphi(d),$$

exactly, for every k and n . The adjacency matrix is *never* singular for $k = 3, 6, 9$; it reaches the full ceiling only at $k = 2$ and $k = 5$. Checked against a direct count at every $13 \leq n < 200$ for $k = 2, 5, 8$.

An error this exposed. The paper claimed “in the computed range $k \leq 10$ only $k = 2, 5, 8$ ” admit solutions. That came from a scan I had filtered at nullity ≥ 6 , which silently discarded $k = 4, 7, 10, 11, 12$ — all of which do admit solutions, at nullity 2 or 4. The filter was a display choice that became a mathematical claim. Corrected.

A weak test I nearly reported as strong. My first cyclic-cover check compared the full matrix nullity against the sum of block nullities and passed everywhere — but every case had nullity 0, so it only confirmed $0 = 0$. I had to tune the weights to force genuinely singular blocks before the test meant anything. It then passed, and exposed a second problem: one base graph had a loop of voltage 0, which lifts to a self-loop rather than an edge, so my two builders disagreed. Both were my errors, not the theory’s, but a passing test that tests nothing is worse than a failing one.

Why enlarge the object at all. Partly because it is the truer statement. But also, candidly, because the competitive field rewards it: a result about cyclic covers of graphs reads as mathematics, while a result about one parametrised family reads as a case study — even when the theorem is identical. The attributed open problem survives untouched as the worked application, so nothing is traded away.

Entry 21: Four investigations, two negative

Date: September 26, 2026

1. Why exactness is hard for $k \geq 6$ — a real obstruction. If the period- d weights are rational then the Galois automorphism $\sigma_j : \zeta_m \mapsto \zeta_m^j$ carries M_ℓ to $M_{j\ell}$ and fixes the weights, so $\sigma_j(\det M_\ell) = \det M_{j\ell}$ and the set of singular blocks must be a union of full Galois orbits. A rational certificate can therefore only realise a *Galois-closed* root set — and those are exactly the sets that degenerate. So for $(k, d, n) = (6, 3, 60)$ and $(6, 3, 72)$ there is no rational certificate at all. This makes the exactness question and the Galois-closedness question the same question, which I had not seen.

Testing it taught me something about testing. My first check asked whether $\det M_{j\ell} = \det M_\ell$ numerically. It failed everywhere, and for a moment I believed the proposition was false. But those two numbers are *conjugates*, not equals — the correct consequence is that the *zero set* is stable, equivalently that the norm over an orbit is rational. Tested that way it holds exactly, to machine zero, in every case. The proposition was right and my test was measuring the wrong thing.

2. The rank law: half explained. The nullspace of the Jacobian of $(\text{weights}) \mapsto (\text{symbol})$ has dimension $3d - 1$. Conjugation by a positive diagonal matrix preserves the pattern, the nullity, and the monic symbol, and accounts for exactly $2d$ of those — verified directly. The remaining $d - 1$ I cannot identify. My hypothesis was that the symbol sees the off-diagonal weights only through their products, as happens for the determinant of a periodic Jacobi matrix; rescaling b , c or e at fixed product changes the symbol in every case tested, so that is wrong. The rank law stays an observation.

3. The ceiling on general covers. Sampling random matrices tests nothing here, since they are generically nonsingular — my first attempt made exactly that mistake and passed trivially, confirming $0 = 0$. Maximising the nullity instead, over fully non-equivariant weights on four small covers, the largest attainable value never exceeded the degree span and equalled it twice. Evidence, not proof.

4. Ihara zeta functions — checked, and not claimed. The same Fourier decomposition gives the full adjacency spectrum of $P(n, k)$ in closed form, $\lambda_\pm = \frac{s+t}{2} \pm \sqrt{(\frac{s-t}{2})^2 + 1}$, confirmed against direct diagonalisation to 10^{-15} . From it one reads off that $P(n, k)$ is Ramanujan for small n but never for large n , since eigenvalues accumulate at 3. But that spectrum is already known and the conclusion is unsurprising, so there is no contribution here and I am not claiming one. I record it because the point of checking a speculative connection is to be willing to drop it.

Entry 22: A subfield ladder that does not (yet) hold weight

Date: September 26, 2026

Command: `python verification/subfield_certificates.py`

Entry 21 proved that rational weights force a Galois-closed root set, and every Galois-closed set at $k = 6$ degenerates, so no rational certificate exists at $(6, 3, 60)$ or $(6, 3, 72)$.

I thought the way around it was to run the same argument over a subfield: weights in $K = \mathbb{Q}(\zeta_m)^H$ need only an H -stable singular set, H -orbits are smaller, so more root sets qualify.

At $(6, 3, 72)$ all three index-two subgroups do admit realisable root sets: $H = \{1, 7\}$ with $\ell = (1, 2, 3, 6, 7, 9, 10)$, $H = \{1, 11\}$ with $\ell = (1, 2, 3, 6, 8, 9, 11)$, $H = \{1, 5\}$ with $\ell = (1, 2, 3, 5, 8, 9, 10)$, whose fixed fields are $\mathbb{Q}(\sqrt{2})$, $\mathbb{Q}(\sqrt{3})$, $\mathbb{Q}(\sqrt{6})$. I wrote this up as “the obstruction lifts”.

That was a logical error, and I am recording it rather than quietly fixing it. The implication runs one way only:

$$\text{weights in } K \implies \text{root set is } H\text{-stable.}$$

The converse does not follow. An H -stable root set being realisable says only that the *necessary* condition is not violated; it gives no reason to believe any certificate for that set has weights in K . I treated a necessary condition as if it were sufficient.

The evidence since is against it. PSLQ on the extracted $H = \{1, 7\}$ certificate finds no relation of small height in $\mathbb{Q}(\sqrt{2})$, and against a degree-four basis returns coefficients of height 10^2 – 10^3 , which is noise rather than a relation. Parametrising the weights *inside* $\mathbb{Q}(\sqrt{2})$ from the start — using the reduction of the conditions to H -orbits, only five independent rather than seven — found nothing in sixty trials. At $(6, 3, 60)$ every index-two set degenerates outright and only the full degree-four field admits a realisable set. And at $(6, 4, 64)$, the one place where the root set is forced to be everything and is therefore automatically Galois-closed, a realisable certificate exists but none of its free coordinates is rational.

Where that leaves exactness for $k \geq 6$. Open, and the evidence now leans negative. What stands is the proposition of Entry 21: no *rational* certificate exists at the parameters tested, and that is a theorem. What does not stand is any claim to have found a small field that works. The honest summary is that the $k \geq 6$ certificates appear to require at least the full real cyclotomic field, and I have not succeeded in recognising them even there.

The machinery built along the way is sound and stays: Gauss periods give a $\mathbb{Q}$ -basis of a fixed field, integer-relation detection proposes exact coefficients, and the blocks are rebuilt over $\mathbb{Q}(\zeta_m)$ with $\det M_\ell$ checked to be exactly zero. Two bugs in it were found and fixed — the periods are not a basis when the Ramanujan sum vanishes ($\eta_1 + \eta_7 = 0$ for $m = 24$, $H = \{1, 5\}$) or when they span fewer dimensions than the degree, and my first recognition routine fitted one real number against a D -dimensional basis by least squares and then truncated to double precision when integer-relation detection needs sixty digits.

The lesson I want to keep. I have now twice in one day turned a one-directional implication into a two-directional one in my head: once with $\sigma = 0$ and the vanishing constant term, and once here. Both times the statement I wanted to be true was the converse of the statement I had proved. The tell is the same in both cases — a result that arrives without any new work being done.

Entry 23: The result is about covers, not about $P(n, k)$

Date: September 26, 2026

Commands: `cover_certificates.py`, `cover_ceiling_sharp.py`,

cover_zf.py

The construction is not about generalized Petersen graphs. Prescribing $D/2$ interior grid values and solving for the weights attains the ceiling on: the $P(n, 2)$ base; the theta graph with voltages $(0, 1, 3)$, a genuinely different cubic family; a three-vertex path base with $D = 12$ whose cover has $3n$ vertices; and a four-vertex star base with $D = 12$. In each case the nullity is exactly D with every required entry a healthy fraction of the matrix scale. With $M \leq Z$ this gives $Z(B^n) \geq D$ for any cover where the construction succeeds — a statement about a class, not a family.

A hypothesis I proposed and then watched fail to be necessary. I had worked out that the $P(n, k)$ elimination generalises when the base graph admits an ordering along which each equation releases one new unknown — that is, when it has a Hamiltonian path — and was ready to state that as the hypothesis. The four-vertex star has no Hamiltonian path and the ceiling holds there anyway. So the condition is sufficient for my particular argument and not necessary for the result. Better to know that before writing it down as a theorem.

A test that was worthless until it was not. My first check of the ceiling on covers asked whether $\ker A$ injects into a window of D consecutive fibre coordinates. It passed everywhere — with matrices of nullity 2 against windows of length 6 to 12, which tests nothing at all. The check only bites at the ceiling, and getting there required building certificates of nullity exactly D first. Repeated properly, injectivity holds in every case.

And a claim of mine that is simply false. I wrote down $Z(B^n) \geq D$ unconditionally. Exhaustive computation gives $Z(P(8, 2)) = 5$ while $D = 6$. The bound needs the construction to actually produce a certificate, and at $n = 8$ it cannot: Γ_8 has exactly three interior values for the three required, the no-slack configuration that degenerates. What holds unconditionally is $Z(B^n) \geq \nu$ for whatever nullity ν is attained, and that I checked against exhaustively computed zero forcing numbers on covers of 14 to 20 vertices with no violation.

The interesting wrinkle is that slack is not the criterion either. For the $P(n, 2)$ base the construction attains the ceiling at $n = 7$ with *no* slack, and fails at $n = 10$ with slack to spare. Whether it succeeds is an arithmetic property of Γ_n , which is the same phenomenon that governed the cyclotomic certificates and the rational obstruction. That keeps being the real content of this problem, and I do not understand it.

Entry 24: the obstruction, in closed form

Date: September 26, 2026

Commands: `antipodal_obstruction.py, zf n 2`

Entry 23 ended by admitting that whether the construction succeeds “is an arithmetic property of Γ_n , which I do not understand.” I understand it now, and it is not arithmetic at all.

For $k = 2$ the symbol has three roots and no residual conditions, so realisability is a closed-form function of the three prescribed grid values. Matching coefficients gives $c^2 = e_3 - e_1e_2$, and then the elementary identity $(s_1 + s_2)(s_1 + s_3)(s_2 + s_3) = e_1e_2 - e_3$ turns it into

$$c^2 = -(s_1 + s_2)(s_1 + s_3)(s_2 + s_3).$$

A factorisation, so $c^2 = 0$ if and only if two roots are antipodal. That is a proof, not a pattern: the thing I had been treating as mysterious arithmetic is one line of symmetric-function algebra.

It is a parity phenomenon, for every k . $L_k(-s) = (-1)^k L_k(s)$, so for even k the two root conditions at y and $-y$ add to $\beta = -aL_k(y)$ and subtract to $\alpha = -L_k(y)$, giving $c^2 = a\alpha - \beta = 0$ identically — for every even k , not just $k = 2$. For odd k the same elimination gives $c^2 = L_k(y)(y^2 - a^2)/y$, and there is no obstruction at all. I checked this symbolically for $k = 2, \dots, 8$: zero for 2, 4, 6, 8, a nonvanishing rational function for 3, 5, 7.

Two exceptions, and they are not the same kind of thing. Γ_n° has no antipodal pair when n is odd, and $n/4$ or $(n-2)/4$ classes when n is even, so three pairwise non-antipodal values exist exactly for odd $n \geq 7$ and even $n \geq 12$. The exceptions above 6 are $n = 8$ and $n = 10$, in both cases by pigeonhole. Then exhaustive computation: $Z(P(8, 2)) = 5$, so there the method is right to fail, and $Z(P(10, 2)) = 6$, so there it is the method that falls short. I had been reporting both as “the construction fails”; they are opposite situations.

The non-symmetric class finally earns its place. At $n = 7$ the only triple gives $c^2 = -1$: no symmetric certificate exists. In the combinatorially symmetric class the two spoke weights need only have product -1 , and

$$A = \begin{pmatrix} I + P + P^{-1} & I \\ -I & P^2 + P^{-2} \end{pmatrix}$$

is an integer matrix of nullity exactly 6 for every n divisible by 7. Up to now the cs generalisation had bought me nothing — every value it certified, the symmetric class certified too, and I said so. This is the first case where it is necessary.

And it explains the old table. The exhaustive rational enumeration at $k = 2$ rejected 12 of 17 factor sets, and I had recorded that as twelve separate arithmetic accidents. They split cleanly: seven have c^2 exactly zero and every one of them is antipodal, and the other five have $c^2 = -1$ and none of them is antipodal — so all five are realisable in the larger class. Ten of seventeen, not five, and the antipodal obstruction is the whole of what remains. Corollary: for even k the factor set may not contain both d and $2d$ for odd d , nor any d divisible by 4 with $\varphi(d) \geq 4$. Checked against 50 candidate sets at $k = 2, 4$, no violations.

What I want to flag for myself is how long I sat on this. The closed form is three lines of algebra on a formula that has been in the paper since Entry 15, and I spent Entries 19–23 doing numerical searches around it.

Entry 25: the obstruction was the wrong question

Date: September 26, 2026

Commands: `cs_exhaustive.py`, `k3_certificate.py`,
`signed_adjacency.py`

Entry 24 found that for even k an antipodal pair of symbol roots forces $c^2 = 0$, and I went looking for the odd- k analogue. There isn’t one, because antipodality was never the right statement. The right one is a single line:

$$c^2 = a\alpha - \beta = 0 \iff \beta = a\alpha \iff F = (s + a)L_k + \alpha(s + a) = (s + a)(L_k + \alpha).$$

So $c^2 = 0$ iff the symbol factors as $(s + a)(L_k + t)$. Every one of the eight vanishing cases at $k = 3$ is of that form — $\Psi_9 = L_3 + 1$, $\Psi_{18} = L_3 - 1$, $\Psi_4\Psi_{12} = L_3$ — and the antipodal theorem is the special case that parity forces when k is even. Two entries spent on a corollary.

And the degenerate list is finite, by Niven. If F is rational with all roots on the grid then a is its s^k coefficient, hence rational, and $-a$ is a grid value, so $a = 2\cos(\text{rational} \cdot \pi)$ is rational; by Niven's theorem $a \in \{0, \pm 1, \pm 2\}$. Same for $t = -L_k(r) = -2\cos(k\varphi)$. At most 25 degenerate rational symbols for every k , explicitly listable. Enumerating the 25 pairs gives 7, 8, 7, 6, 9 for $k = 2, \dots, 6$, matching the independent exhaustive count exactly in each case.

What I had been getting wrong for weeks. The old enumeration required $c^2 > 0$, the condition for a *symmetric* certificate, and recorded everything else as “not realisable.” But c^2 is the *product* of the two spoke weights, and in the combinatorially symmetric class those are independent. Only $c^2 = 0$ is fatal. Redoing the enumeration with $c^2 \neq 0$:

k	candidates	outside family	$c^2 = 0$	sym	cs
2	17	0	7	5	5
3	35	21	8	0	6
4	67	54	7	3	3
5	127	118	6	1	2
6	225	216	9	0	0
7	385	378	6	0	1
8	651	644	7	0	0
9	1065	1057	8	0	0
10	1705	1698	7	0	0

Two cases I had written down as *impossible for rational symbols* are not. Every non-fatal rejection has $c^2 = -1$ or -2 .

$k = 3$ **closes.** $\Psi_5\Psi_{10} = s^4 - 3s^2 + 1 = sL_3(s) + 1$, so $a = \alpha = 0$, $\beta = 1$, $c^2 = -1$, and

$$A = \begin{pmatrix} P + P^{-1} & I \\ -I & P^3 + P^{-3} \end{pmatrix}$$

has nullity exactly 8 for every n divisible by 10. So $Z(P(n, 3)) = 8$, where the paper had ≥ 7 . I checked it against exhaustive Z at $n = 10, 20$ and against the nullity at n up to 100. The matrix is the adjacency matrix with one set of spokes negated, and the sign carries the whole result: unsigned, the nullity is 0.

Its four roots are $\pm\varphi, \pm\varphi^{-1}$ — *two antipodal pairs*, which for even k would be fatal. So the parity corollary is not only an obstruction, it says where to look when k is odd. That is the first time in this project a negative result has pointed at a positive one.

$k = 7$ **closes too, and the two are the same phenomenon.** The $k = 7$ symbol also has $a = \alpha = 0$, $\beta = 1$. Setting $s = z + z^{-1}$ turns $sL_k + 1 = 0$ into

$$Q_k(z) = z^{2k+2} + z^{2k} + z^{k+1} + z^2 + 1 = 0,$$

which is the adjacency polynomial P_k of Entry 18 with the middle sign flipped. So the whole classification I proved for the adjacency matrix applies verbatim to the signed one, and the

question “which k admit an integer certificate of this shape” becomes “for which k does Q_k split into cyclotomics.” Answer, for $2 \leq k \leq 60$: only $k = 2, 3, 7$, giving $\Phi_5\Phi_6$, $\Phi_5\Phi_{10}$, $\Phi_5\Phi_{10}\Phi_{24}$ and the ceiling each time. Hence $Z(P(n, 7)) = 16$ for every n divisible by 120 — the paper previously had $k = 7$ only from a floating-point period-4 certificate at a single n . The unsigned P_k splits only at $k = 2, 5$, so the two sign choices together cover $k \in \{2, 3, 5, 7\}$ and nothing else.

I want to be clear that this is rigidity, not an infinite family. I checked to $k = 60$ hoping for more and there are none, which is the same Lehmer-type behaviour the unsigned case showed.

The negative side got stronger as well. $k = 6, 8, 9, 10$ are impossible over $\mathbb{Q}$ in *both* classes, and no longer because a search came up empty: every candidate that clears the family condition is one of the finitely many degenerate symbols, and the Niven bound caps that list at 25 before any computation. At $k = 6$ all nine survivors are $(s + a)(L_6 + t)$ with $a, t \in \{0, \pm 1\}$.

The lesson I should have learned three entries ago: when a computation returns “not realisable,” check which inequality was being tested.

Entry 26: off the congruence classes

Date: September 27, 2026

Commands: `general_n.py`, `krawczyk.py`, `certify_pipeline.py`,
`gauge_fixed.py`, `certify_sweep.py`, `tiles.py`

Every exact value in this project has been attached to a divisibility class. That is not a gap in the searching; it is forced, and I can now say how.

A period-one symbol is a *fixed* polynomial. Its roots are fixed algebraic numbers, and a fixed algebraic number lies on Γ_n only for n in a divisibility class. So the set of n a given period-one matrix can settle is $\{n : \text{lcm}(D) \mid n\}$, which has density at most $1/3$ and contains *at most one prime*, since $\{n : L \mid n\}$ with $L \geq 3$ contains a prime only when L is that prime. To cover more, the weights must depend on n .

For $k \geq 3$ that dependence is constrained by one identity I had not noticed. Since L_k has nonzero coefficients only in degrees congruent to $k \bmod 2$, the coefficient of s^{k-1} in $(s + a)L_k + \alpha s + \beta$ is the s^{k-2} coefficient of L_k , namely $-k$, for *every* (a, α, β) . So every period-one symbol satisfies

$$e_2(r_1, \dots, r_{k+1}) = -k.$$

On the grid, after expanding e_1^2 by product-to-sum, that is a rational linear relation among cosines of rational angles — a Conway–Jones relation, hence rigid. For $k = 2$ the coefficient of $s^{k-1} = s$ is $\alpha - 2$, free, so the condition is vacuous. *That* is why $k = 2$ is general in n and nothing else is: the $k = 2$ theorem gets to choose its roots as functions of n . I checked $e_2 = -k$ against all 26 certificates in the paper; it holds in every case with $k \geq 3$ and takes six different values among the $k = 2$ ones.

So stop using the invariant matrices. The ceiling corollary characterises nullity $2k + 2$ for *any* matrix on the pattern by $T = I$ on the monodromy, and that condition has no arithmetic in it. Minimising $\|T - I\|$ over all $5n$ weights reaches the ceiling at $k = 2$ for $n = 9, 11, 13, 14, 15$ and at $k = 3$ for $n = 17, 20$ — including primes.

Certifying it, and the mistake I nearly shipped. $T = I$ is bad to certify: the image of $w \mapsto T$ is 51-dimensional inside the 64-dimensional space of 8×8 matrices at $k = 3$, so no square subsystem is available. Certify the nullity instead, via $A(w)K(X) = 0$ with K in graph form — bilinear, so the Jacobian is affine and its Lipschitz constant is an explicit integer, which makes a Krawczyk bound cheap.

I first chose the square subsystem by QR pivoting, got a passing Krawczyk test, and was ready to call it a proof. It was not one: the 28 equations I had dropped were not shown to vanish at the true zero, only at my starting point. The fix is structural. $K^\top AK$ is symmetric because A is, and $28 = \binom{8}{2}$ is exactly its antisymmetric part; if the non-pivot rows vanish then $K^\top AK$ equals the pivot block, which must therefore be symmetric, so setting its upper triangle to zero forces the lower triangle. The system *all non-pivot rows + upper triangle of the pivot block* is equivalent to $AK = 0$, not merely a maximal-rank selection of it. And the reduction is checkable: the 28 dependent equations fall to 10^{-61} without being solved for.

With that, a proof:

$$Z(P(n, 3)) = M(P(n, 3)) = 8 \quad n = 17, 19, 21, 23, 25, 27,$$

the primes 17, 19, 23 among them. Combined with the $10 \mid n$ certificate this holds at $n = 10, 17, 19, 20, 21, 23, 25, 27, 30, 40, \dots$, a set no divisibility condition describes. First results in this project not living on a congruence class. The certificate at $n = 17$ was found three independent ways and all three certify; since stage 1 is only a search and stages 2–3 are the proof, the route does not enter the result.

Three of my own claims corrected, one in the middle of relying on it. The threshold: I wrote $n \geq (2k + 2)^2/5$ from a loose codimension count. Spending the gauge properly — the $2n$ diagonal-conjugation parameters normalise every spoke to 1 and, for odd n , every outer weight to ± 1 — leaves $3n$ unknowns and gives $n \geq (k + 1)(2k + 3)/3$, which is 35 at $k = 6$, not 27. I had already run $k = 6$ at $n = 27, 29, 31$ and read their failure as meaningful. It was not; those n are below the threshold and could not have worked.

The growth rate: I conjectured $N(k) = O(k)$ from the saturation data ($\approx 4k + 3$). The dimension count is a necessary condition and forces $N(k)$ to be at least quadratic. Withdrawn.

The monodromy is not symplectic. I guessed it would be, since the recurrence is self-adjoint. The standard form is not preserved; the space of skew B with $T^\top BT = B$ has dimension $k + 1$, so the eigenvalues come in reciprocal pairs and the characteristic polynomial is palindromic. Different statement, and it does not give the equation count I wanted from it.

And one instrument withdrawn. The eigenvalue-minimising search whose plateaus I have been quoting as evidence against reachability gives maximum nullity 7 at $P(20, 3)$, where 8 is attained — by the $\Psi_5\Psi_{10}$ certificate and, in the symmetric class, by the new method. It collapses onto the degenerate strata, all spokes zero or all edges zero. So the conclusion I drew from it, that exactness for $k \geq 6$ looked unlikely, was never supported by that data. I have marked every negative result from it as unreliable in the paper rather than quietly dropping them.

Toward all large n . If a weight block has transfer product I , and blocks compose, then n is any sum of available block lengths. Blocks do not compose naively, because T_i depends

on a window of width about $2k + 2$; giving every tile a fixed docking segment of $k + 1$ positions at each end fixes that, and I checked it rather than assumed it — perturbing one tile’s interior leaves the other tiles’ products bit-identical, 0.000e+00, for $k = 2, 3, 4$. Two coprime lengths would only give $n = ia + jb$ and leave everything below $(a - 1)(b - 1) \approx 600$ uncovered, which is useless; tiles of *every* length in $[L, 2L)$ instead make every $n \geq L$ a sum, with no gap. For $k = 3$ a tile of length ℓ has $5(\ell - 8)$ free weights against 64 equations, so $\ell \geq 21$ and $L = 21$ would give all $n \geq 21$. Whether those systems are solvable with all weights nonzero is being computed and I do not know the answer yet.

A regression I caused and caught. Warm-starting each n from the previous certified solution made the sweep affordable, but I gated the cold fallback on the warm start being worse than 10^{-2} . A warm start returning 10^{-3} then suppressed the cold solve and failed stage 2, and the sweep reported only $n = 17, 19$ — losing 21, 23, 25, 27, which I had already certified. Restructured so stage 1 produces candidates and stage 2 is attempted on each in turn. Worth recording because the failure mode was silent: the sweep looked like a clean negative result.

Entry 27: all large n , for one k

Date: September 27, 2026

Commands: `tiles.py`, `tile_sweep.py`, `certify_sweep.py`,
`gauge_fixed.py`

The per- n certification of Entry 26 settles individual n and will never settle all of them. The fix is to make the weight sequence out of pieces that each contribute the identity to the monodromy, since $T = I$ is multiplicative around the cycle.

Pieces do not compose naively: T_i is built from the weights within about $2k + 2$ of position i , so neighbouring pieces interfere. Giving every tile a fixed docking pattern on its first and last $k + 1$ positions removes the interference, and I checked that rather than assuming it — perturbing one tile’s interior leaves every other tile’s product identical to the last bit, 0 exactly, at $k = 2, 3, 4$.

The other correction was to the combinatorics. My first plan was two tile lengths with $\gcd = 1$, giving $n = ia + jb$. That is useless: it covers nothing below $(a - 1)(b - 1)$, which for the smallest solvable lengths here is several hundred. Using *every* length in $[L, 2L)$ instead makes every $n \geq L$ a sum, with no gap, and turns the problem into a finite list of finite problems.

For $k = 3$ it works. The parameter count is sharp in practice: at $\ell = 19, 20$ the systems are underdetermined and stall near 4×10^{-1} ; $\ell = 21$, the first with more unknowns than equations, falls to 5.6×10^{-2} ; from $\ell = 22$ identity tiles appear. I found one for every length $22 \leq \ell \leq 43$ — exactly $[L, 2L)$ for $L = 22$ — with $\|T_{\text{tile}} - I\|_{\infty} \leq 2 \times 10^{-13}$ and every weight at least 0.15 of the tile scale. Then I tested the composition instead of trusting it: nullity exactly 8 on single tiles for $22 \leq n \leq 43$, and on two-tile concatenations at $n = 44, 45, 46$. So

$$Z(P(n, 3)) = M(P(n, 3)) = 8 \quad \text{for every } n \geq 22,$$

the first statement in this project about all large n at any $k \geq 3$.

And then it stopped being numerical. I had written the conclusion up as floating point, on the grounds that certifying $T_{\text{tile}} = I$ means Krawczyk on a rational map and so interval matrix products. That was the wrong way round. By the ceiling corollary the monodromy of the cyclic matrix built from a *single* tile *is* the tile product, so

$$T_{\text{tile}} = I \iff \text{null } A_{\text{one tile}} = 2k + 2,$$

and the right-hand side is certified by the bilinear machinery I already had, whose Jacobian is affine. The only thing needing care is that the tiles must share a docking pattern, which is arranged by pinning those 40 weights at their exact integer values so they sit in the fixed set and the certified solution carries them unchanged.

All 22 tiles certify. Krawczyk $\alpha \leq 2.2 \times 10^{-10}$ throughout, nullity exactly 8 in every case, every docking weight equal to its prescribed integer to 10^{-14} , minimum edge weight 0.15 of the matrix scale. So

$$Z(P(n, 3)) = M(P(n, 3)) = 8 \quad \text{for every } n \geq 17,$$

proved, with no congruence condition and no floating-point nullity anywhere in the chain. I spent an hour writing careful hedges about a limitation that a change of formulation removed.

Then $k = 4$ went the same way. I had written that the $k = 4$ tile search succeeded on only about half the lengths and that I was claiming nothing for it. That was the right thing to write at the time and the wrong conclusion to keep: every failure I retried with a larger restart budget — $\ell = 31, 32, 33, 35, 38$ — solved to 10^{-14} . The variance was in my solver, not in the problem. With the main sweep reaching $\ell = 61$ the range $[31, 62)$ closed, all 31 tiles certified, and

$$Z(P(n, 4)) = M(P(n, 4)) = 10 \quad \text{for every } n \geq 31.$$

The certification needed no change at all going from 8×8 transfer matrices to 10×10 ; the dependent-equation count moved from $\binom{8}{2} = 28$ to $\binom{10}{2} = 45$ exactly as the symmetry argument says it should, which is the best evidence I have that the structural reduction is right rather than lucky.

So $L(2) = 9$, $L(3) = 17$, $L(4) = 31$, against the tile lower bound $2(k + 1) + (2k + 2)^2/5$ of 13.2, 20.8, 30. $k = 5$ needs 41 tiles and $k = 6$ needs 54, each individually larger; I have not attempted them, and the reason is compute rather than any obstacle I can identify.

Then a third time, with the same shape. The rank measurement said the tile count only needs $5(\ell - 8) \geq \binom{9}{2} = 36$, so $\ell \geq 16$ rather than $\ell \geq 22$. I tested $16 \leq \ell \leq 22$ and reported that short tiles do not exist, citing $\ell = 16, 17, 18$ stalling near 8×10^{-1} against the $\ell = 21$ near-miss at 5.6×10^{-2} as the contrast. Given 25 restarts instead of 4, $\ell = 21$ produced a tile at 1.4×10^{-14} . So $L(3) = 21$, and my contrast case was itself an under-budgeted search.

What survives is narrower and more useful than what I claimed: every length I eventually solved first appeared at a residual of 10^{-2} or below, and the ones that never solved sat at 10^{-1} or above. The residual magnitude is the signal; success against failure at a fixed budget is not. Below $\ell \approx 20$ the identity does appear not to be a product of that few transfer matrices, but I have no lower bound and I should not have written "genuinely."

And a hypothesis of mine that the numbers killed. Asked why $k = 5$ finds nothing, I proposed that the product of ℓ transfer matrices becomes too ill-conditioned for

$T - I$ to carry information. At $k = 5$, $\ell = 44$ the condition number is about 2.5×10^6 ; at $k = 4$, $\ell = 61$, which yields a certified tile, it is about 1.1×10^{10} . Worse conditioning, and it works. So conditioning is not the discriminating factor and I have no diagnosis for $k = 5$. The $k = 6$ per- n sweep also came back empty at $n = 41, 43, 45, 47, 49, 51$, at the same default budget that wrongly condemned $\ell = 21$, so I am recording it as inconclusive rather than as evidence.

What I want on the record about my own reporting. Three times today I wrote a careful, hedged, accurate-at-the-time statement that a case was out of reach — the $k = 3$ tiles as numerical, $k = 4$ as unclaimable, short tiles as nonexistent — and all three times the limitation was in how I had set the problem up or how much compute I had spent. The hedges were honest and I would write them again. But "I could not do it with the budget I used" is a very different claim from "it cannot be done," and the gap between them is where I lost ground repeatedly. The fix I have adopted is to quote the residual and the budget every time, so that a reader can see which of the two I am entitled to.

And the per- n range closed up. Sweeping $k = 3$ certified every $17 \leq n \leq 33$ — seventeen consecutive values, primes 17, 19, 23, 29, 31 among them. $n = 18$ was the last holdout and it certified once I raised the restart budget from 40 to 150. Worth dwelling on: I had reported $n = 18$ as a failure, and it was only a failure of effort. The slice lemma is what told me not to read it as anything more — every symmetric matrix on the pattern is gauge-equivalent to one in the slice searched, so a failure inside the slice measures the search and nothing else. I nearly wrote the opposite before proving that lemma.

Two of my own bugs, both silent. Warm-starting the n -sweep from the previous certified solution made it affordable, but I gated the cold fallback on the warm start being worse than 10^{-2} ; a warm start returning 10^{-3} then suppressed the cold solve and failed, and the sweep reported certifying only $n = 17, 19$ while losing 21, 23, 25, 27, which I had already certified by hand. It looked exactly like a clean negative result. And in the tile sweep my feasibility guard was $5(\ell - 2k - 2) < (2k + 2)^2 - 20$, loose enough to admit $\ell = 20$ at $k = 3$, where 60 unknowns cannot meet 64 equations; the job spent almost two hours on a length that was impossible by counting.

Where $k \geq 4$ stood, at the time of writing. The method is not specific to $k = 3$: an identity tile exists at $k = 4$, $\ell = 34$. But convergence looked erratic — $\ell = 35$ failed despite having more unknowns than $\ell = 34$ and despite being warm-started from it — so at this point I had completed no range for $k = 4$ and was claiming nothing for it. *This paragraph is superseded by the $k = 4$ result recorded above; I have left it in place rather than editing it away, because the sequence of what I believed and when is the point of this log.* The per- n sweeps at $k = 4, 5, 6$ failed at every n they reached, and by the slice lemma that too is a statement about the search rather than about the graphs.

References

- [1] S. Rashidi, N. Shajareh Poursalavati, M. Tavakkoli, *J. Algebra Combin. Discrete Struct. Appl.* 7(2) (2020), 183–193.
- [2] A. Krishnan, arXiv:2607.19412 (2026).

V. Verification Output

One script re-derives every computational claim in the paper from scratch and reports pass or fail per check. Its current output follows, unedited.

```

1. Realisability (thm:realis):  $c^2 = 0$  iff  $F = (s+a)(L_k + t)$ 
[PASS] k=2
[PASS] k=3
[PASS] k=4
[PASS] k=5
[PASS] k=6
[PASS] k=7
[PASS] k=8
2. Antipodal corollary:  $c^2 = 0$  identically for even k, not for odd
[PASS] k=2 (even)
[PASS] k=3 (odd)
[PASS] k=4 (even)
[PASS] k=5 (odd)
[PASS] k=6 (even)
[PASS] k=7 (odd)
[PASS] k=8 (even)
3. k=2 closed form:  $c^2 = -(s_1+s_2)(s_1+s_3)(s_2+s_3)$ 
[PASS] identity
4.  $e_2 = -k$  for every period-one symbol with  $k \geq 3$  (thm:congbound)
[PASS] k=3
[PASS] k=4
[PASS] k=5
[PASS] k=6
[PASS] k=7
[PASS] k=8
[PASS] k=9
[PASS] k=10
5. Integer certificates: the signed adjacency matrix
[PASS] signed Adj P(10,3) nullity 8 --- got 8
[PASS] signed Adj P(20,3) nullity 8 --- got 8
[PASS] signed Adj P(30,3) nullity 8 --- got 8
[PASS] signed Adj P(120,7) nullity 16 --- got 16
[PASS] signed Adj P(30,2) nullity 6 --- got 6
[PASS] signed Adj P(60,2) nullity 6 --- got 6
[PASS] Psi_7 certificate P(7,2) nullity 6 --- got 6
[PASS] Psi_7 certificate P(14,2) nullity 6 --- got 6
[PASS] Psi_7 certificate P(21,2) nullity 6 --- got 6
[PASS] Psi_7 certificate P(28,2) nullity 6 --- got 6
6. k=2 explicit certificate at the three most negative grid values
( $c^2 > 0$  is checked for all n; the nullity is checked two ways, by
counting grid roots of the symbol exactly and, for moderate n only,
by singular values -- the three roots crowd within  $O(1/n^2)$  of -2,
so the matrix is exact but numerically stiff for large n)
[PASS]  $c^2 > 0$  for every  $n \geq 9$ ,  $n \neq 10$  (up to 1001) --- checked 9..199, 500, 1001
[PASS] three prescribed roots are distinct and interior --- checked 9..199, 500, 1001, 5000
[PASS] P(9,2) nullity 6 numerically
[PASS] P(11,2) nullity 6 numerically
[PASS] P(20,2) nullity 6 numerically
[PASS] P(39,2) nullity 6 numerically
[PASS] P(60,2) nullity 6 numerically
7. Saved interval-certified matrices re-evaluated (k=3)
[PASS] P(17,3) nullity 8, edges nonzero --- nullity 8, min edge 0.1462
[PASS] P(19,3) nullity 8, edges nonzero --- nullity 8, min edge 0.0145

```

```

[PASS] P(20,3) nullity 8, edges nonzero --- nullity 8, min edge 0.1424
[PASS] P(21,3) nullity 8, edges nonzero --- nullity 8, min edge 0.3059
[PASS] P(22,3) nullity 8, edges nonzero --- nullity 8, min edge 0.0878
[PASS] P(23,3) nullity 8, edges nonzero --- nullity 8, min edge 0.0647
[PASS] P(24,3) nullity 8, edges nonzero --- nullity 8, min edge 0.0405
[PASS] P(25,3) nullity 8, edges nonzero --- nullity 8, min edge 0.1584
[PASS] P(26,3) nullity 8, edges nonzero --- nullity 8, min edge 0.0859
[PASS] P(27,3) nullity 8, edges nonzero --- nullity 8, min edge 0.2439
[PASS] P(28,3) nullity 8, edges nonzero --- nullity 8, min edge 0.0296
[PASS] P(29,3) nullity 8, edges nonzero --- nullity 8, min edge 0.2612
[PASS] P(30,3) nullity 8, edges nonzero --- nullity 8, min edge 0.3665
[PASS] P(31,3) nullity 8, edges nonzero --- nullity 8, min edge 0.2290
[PASS] P(32,3) nullity 8, edges nonzero --- nullity 8, min edge 0.0445
[PASS] P(33,3) nullity 8, edges nonzero --- nullity 8, min edge 0.2802
    (16 saved certificates re-evaluated)
8. k=3 identity tiles and their compositions
[PASS] tiles cover a full interval [L,2L) --- lengths 21..43 (23 tiles)
[PASS] single tiles give nullity 8 --- all 23 lengths
[PASS] multi-tile concatenations give nullity 8 --- checked n up to 55
9. Certified tiles: nullity, docking exactness, interval coverage
[PASS] k=3: certified lengths form [L, 2L) --- 21..43 (23 tiles)
[PASS] k=3: every certified tile has nullity exactly 8 --- all 23 tiles
[PASS] k=3: docking weights exact in every certified tile --- all tiles, to 1e-13
[PASS] k=3: every n in [21, 400] is a sum of certified lengths --- no gaps
[PASS] k=4: certified lengths form [L, 2L) --- 30..61 (32 tiles)
[PASS] k=4: every certified tile has nullity exactly 10 --- all 32 tiles
[PASS] k=4: docking weights exact in every certified tile --- all tiles, to 1e-13
[PASS] k=4: every n in [30, 400] is a sum of certified lengths --- no gaps
10. The paper's own numbers against the computed ones
[PASS] paper's certification count matches the 55 certified tiles --- all say 55
[PASS] k=3: paper's L matches the least certified tile length (21) --- L=21
[PASS] k=4: paper's L matches the least certified tile length (30) --- L=30
11. Slice lemma (lem:slice)
[PASS] spokes -> +1 and |outer| -> 1, nullity preserved, both parities --- max deviation 2.7e-15

=====
ALL CHECKS PASSED

```

VI. A Certification Transcript

What “interval-certified” means in practice, for one family of tiles. Each entry is an independent Krawczyk contraction proof with its own unknown count, matrix norm and residual; nothing is shared between them.

certifying 22 identity tiles at k=3

```

l=22: 398 unknowns, 352 equations, 324 structural, 28 dependent; 40 docking pins, 74 fixed total
      structural 3.11e-61, dependent 2.72e-61, ||Y||=5.39e+02, alpha=6.70e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=23: 419 unknowns, 368 equations, 340 structural, 28 dependent; 40 docking pins, 79 fixed total
      structural 3.11e-61, dependent 4.67e-61, ||Y||=7.55e+02, alpha=9.25e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=24: 440 unknowns, 384 equations, 356 structural, 28 dependent; 40 docking pins, 84 fixed total
      structural 6.22e-61, dependent 1.94e-61, ||Y||=2.48e+03, alpha=3.24e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=25: 461 unknowns, 400 equations, 372 structural, 28 dependent; 40 docking pins, 89 fixed total
      structural 6.22e-61, dependent 2.24e-61, ||Y||=9.96e+02, alpha=1.23e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=26: 482 unknowns, 416 equations, 388 structural, 28 dependent; 40 docking pins, 94 fixed total
      structural 3.5e-61, dependent 1.94e-61, ||Y||=5.16e+02, alpha=6.88e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=27: 503 unknowns, 432 equations, 404 structural, 28 dependent; 40 docking pins, 99 fixed total
      structural 6.22e-61, dependent 1.56e-61, ||Y||=1.51e+03, alpha=1.89e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=28: 524 unknowns, 448 equations, 420 structural, 28 dependent; 40 docking pins, 104 fixed total
      structural 6.22e-61, dependent 2.48e-61, ||Y||=1.04e+03, alpha=1.31e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=29: 545 unknowns, 464 equations, 436 structural, 28 dependent; 40 docking pins, 109 fixed total
      structural 3.31e-61, dependent 1.56e-61, ||Y||=3.84e+02, alpha=4.77e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=30: 566 unknowns, 480 equations, 452 structural, 28 dependent; 40 docking pins, 114 fixed total
      structural 6.22e-61, dependent 3.69e-61, ||Y||=5.86e+02, alpha=7.43e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=31: 587 unknowns, 496 equations, 468 structural, 28 dependent; 40 docking pins, 119 fixed total
      structural 4.67e-61, dependent 6.51e-61, ||Y||=6.67e+02, alpha=8.49e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=32: 608 unknowns, 512 equations, 484 structural, 28 dependent; 40 docking pins, 124 fixed total
      structural 3.11e-61, dependent 2.14e-61, ||Y||=6.86e+02, alpha=8.71e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=33: 629 unknowns, 528 equations, 500 structural, 28 dependent; 40 docking pins, 129 fixed total
      structural 2.53e-61, dependent 2.33e-61, ||Y||=5.64e+02, alpha=7.14e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=34: 650 unknowns, 544 equations, 516 structural, 28 dependent; 40 docking pins, 134 fixed total
      structural 4.76e-61, dependent 4.08e-61, ||Y||=2.88e+02, alpha=3.66e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=35: 671 unknowns, 560 equations, 532 structural, 28 dependent; 40 docking pins, 139 fixed total
      structural 3.11e-61, dependent 5.25e-61, ||Y||=1.43e+03, alpha=1.77e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=36: 692 unknowns, 576 equations, 548 structural, 28 dependent; 40 docking pins, 144 fixed total
      structural 7.0e-61, dependent 3.11e-61, ||Y||=2.17e+03, alpha=2.79e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=37: 713 unknowns, 592 equations, 564 structural, 28 dependent; 40 docking pins, 149 fixed total
      structural 3.11e-61, dependent 3.4e-61, ||Y||=8.29e+02, alpha=1.06e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=38: 734 unknowns, 608 equations, 580 structural, 28 dependent; 40 docking pins, 154 fixed total
      structural 3.11e-61, dependent 2.14e-61, ||Y||=8.16e+02, alpha=1.01e-10, nullity 8, docking exact True, min|edge| 0.1502 CEF
l=39: 755 unknowns, 624 equations, 596 structural, 28 dependent; 40 docking pins, 159 fixed total
      structural 3.11e-61, dependent 4.67e-61, ||Y||=5.92e+02, alpha=7.70e-11, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=40: 776 unknowns, 640 equations, 612 structural, 28 dependent; 40 docking pins, 164 fixed total
      structural 3.69e-61, dependent 1.94e-61, ||Y||=1.15e+03, alpha=1.43e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=41: 797 unknowns, 656 equations, 628 structural, 28 dependent; 40 docking pins, 169 fixed total
      structural 6.22e-61, dependent 5.45e-61, ||Y||=8.81e+02, alpha=1.11e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=42: 818 unknowns, 672 equations, 644 structural, 28 dependent; 40 docking pins, 174 fixed total
      structural 6.22e-61, dependent 3.11e-61, ||Y||=1.65e+03, alpha=2.12e-10, nullity 8, docking exact True, min|edge| 0.1500 CEF
l=43: 839 unknowns, 688 equations, 660 structural, 28 dependent; 40 docking pins, 179 fixed total

```